

WORKSHOP INTRODUCTION

This workshop provides an introduction to chaos engineering using AWS tooling, with a focus on AWS Fault Injection Simulator (AWS FIS). It introduces the core elements of chaos engineering: stress, observe, and improve. You will learn how to use AWS FIS and other AWS tools to inject faults in your infrastructure to validate your system's resilience as well as verifying your alarms, observability, and monitoring practices.

Target audience

This is a technical workshop introducing chaos engineering concepts for dev, QA and Ops teams. For best results, the participants should have familiarity with the AWS console as well as some proficiency with command-line tooling.

Additionally, chaos engineering is about proving or disproving a hypothesis of how a particular fault might affect the overall system behavior (steady-state) so an understanding of the systems being disrupted is helpful but not required to do the workshop.

Duration

When run in a prepared AWS account the core sections (EC2, RDS, SSM) of the workshop will take between 1-2h. The whole workshop about 2-4h. Using the Amazon DevOps Guru section will require an additional 2-24h of wait time after the infrastructure has been configured.

When run in a customer account, deploying the workshop's core infrastructure will require an additional 45min.

Cost

When run in a private customer account, this workshop will incur costs on the order of USD1/h for the infrastructure created. Please ensure you clean up all infrastructure after finishing the workshop to prevent continuing expenses. You can find instructions in the [Cleanup](#) section.



START THE WORKSHOP

To start the workshop, follow one of the following links: depending on whether you are...

- [Running the workshop in your own account](#)
- [Running in an AWS provided account \(using AWS provided hashes\)](#)

Once you have completed one of the setup paths above, continue with [Region Selection](#)



...ON YOUR OWN

Running the workshop on your own

Warning

Only complete this section if you are running the workshop on your own. If you are at an AWS hosted event (such as re:Invent, Kubecon, Immersion Day, etc), go to [Start the workshop at an AWS event](#).

Next step:

- Create an AWS account



CREATE AN AWS ACCOUNT

Warning

Your account must have the ability to create new IAM roles and scope other IAM permissions.

1. If you don't already have an AWS account with "**Administrator**" access:

[Create an AWS account by clicking here.](#)

2. Once you have an AWS account, ensure you are following the remaining workshop steps as an IAM user with administrator access to the AWS account: [Create a new IAM user to use for the workshop](#)

3. Enter the user details:

Add user

1 2 3 4

Set user details

You can add multiple users at once with the same access type and permissions. [Learn more](#)

User name* workshop  1) Create a new user

+ Add another user

Select AWS access type

Select how these users will access AWS. Access keys and autogenerated passwords are provided in the last step. [Learn more](#)

Access type* **Programmatic access**
Enables an **access key ID** and **secret access key** for the AWS API, CLI, SDK, and other development tools.

2) Enable Console  **AWS Management Console access**
Enables a **password** that allows users to sign-in to the AWS Management Console.

Console password* Autogenerated password
 Custom password

3) Set a password 
 Show password

Require password reset **4) Uncheck reset** next sign-in
Users automatically get the [IAmUserChangePassword](#) policy to allow them to change their own password.

5) Next 

* Required

Cancel [Next: Permissions](#)

4. Attach the “**AdministratorAccess**” IAM Policy:

Add user

1 2 3 4

▼ Set permissions

Add user to group Copy permissions from existing user Attach existing policies directly  1) Attach Policy

Create policy 

	Policy name ▾	Type	Used as	Description
<input checked="" type="checkbox"/>	AdministratorAccess	Job function	Permissions policy (3)	Provides full access to AWS services and re...
<input type="checkbox"/>	AlexaForBusinessSS... 2) Check this policy	AWS managed	None	Provide device setup access to AlexaForBu...
<input type="checkbox"/>	AlexaForBusinessR...	AWS managed	None	Grants full access to AlexaForBusiness reso...
<input type="checkbox"/>	AlexaForBusinessG...	AWS managed	None	Provide gateway execution access to Alexa...
<input type="checkbox"/>	AlexaForBusinessR...	AWS managed	None	Provide read only access to AlexaForBusine...
<input type="checkbox"/>	AmazonAPIGatewa...	AWS managed	None	Provides full access to create/edit/delete A...
<input type="checkbox"/>	AmazonAPIGatewa...	AWS managed	None	Provides full access to invoke APIs in Amaz...
<input type="checkbox"/>	AmazonAPIGatewa...	AWS managed	None	Allows API Gateway to push logs to user's ...

Filter policies ▾  Search Showing 359 results

▶ Set permissions boundary  3) Next

Cancel **Previous** **Next: Review**

5. Select “**Create user**”:

Add user

1 2 3 4

Review

Review your choices. After you create the user, you can view and download the autogenerated password and access key.

User details

User name	workshop
AWS access type	AWS Management Console access - with a password
Console password type	Custom
Require password reset	No
Permissions boundary	Permissions boundary is not set

Permissions summary

The following policies will be attached to the user shown above.

Type	Name
Managed policy	AdministratorAccess

1) Create User



[Cancel](#)

[Previous](#)

[Create user](#)

6. Take note of the sign-in URL and save:

Add user

1 2 3 4

Success

You successfully created the users shown below. You can view and download user security credentials. You can also email users instructions for signing in to the AWS Management Console. This is the last time these credentials will be available to download. However, you can create new credentials at any time.

Users with AWS Management Console access can sign-in at: [https://\[REDACTED\].signin.aws.amazon.com/console](https://[REDACTED].signin.aws.amazon.com/console)

Save this URL

[Download .csv](#)

	User
▶	workshop

7. Sign out of your current AWS Console session: on the top menu, select your login and select "**Sign out**"



Oregon ▾

Support ▾

The screenshot shows the AWS Lambda console interface. On the left, there's a sidebar with a "Stay connected to your AWS go" section featuring an icon of a smartphone with a gear inside, and text about the AWS Console Mobile App. Below this is a "Explore AWS" button. On the right, a vertical navigation bar has a dropdown menu open under "My Account". The menu items are: My Account 7209 (highlighted with a white box), My Organization, My Service Quotas, My Billing Dashboard, Switch Roles, and Sign Out (highlighted with a red box). The "Sign Out" button is located at the bottom of the menu.

8. Sign in to a new AWS Console session by using the sign-in URL saved and the newly created user credentials.
9. Once you have completed the steps above, you can head straight to the **Region Selection**.



...AT AN AWS EVENT

Running the workshop at an AWS Event

Warning

Only complete this section if you are at an AWS hosted event (such as re:Invent, AWS Summit, Immersion Day, or any other event hosted by an AWS employee). If you are running the workshop on your own, go to: **Start the workshop on your own**.

Next step:

- [AWS Workshop Portal](#)



AWS WORKSHOP PORTAL

Login to AWS Workshop Portal

This workshop uses an AWS account and a AWS Cloud9 environment. You will need the **Participant Hash** provided by the event organizers and your email address to track your unique session.

Connect to the portal by following instructions sent by the organizers or by browsing to <https://dashboard.eventengine.run/>. You should see the following screen:

Who are you?

Terms & Conditions:

1. By using the Event Engine for the relevant event, you agree to the Event Terms and Conditions and the AWS Acceptable Use Policy. You acknowledge and agree that are using an AWS-owned account that you can only access for the duration of the relevant event. If you find residual resources or materials in the AWS-owned account, you will make us aware and cease use of the account. AWS reserves the right to terminate the account and delete the contents at any time.
2. You will not: (a) process or run any operation on any data other than test data sets or lab-approved materials by AWS, and (b) copy, import, export or otherwise create derivative works of materials provided by AWS, including but not limited to, data sets.
3. AWS is under no obligation to enable the transmission of your materials through Event Engine and may, in its discretion, edit, block, refuse to post, or remove your materials at any time.
4. Your use of the Event Engine will comply with these terms and all applicable laws, and your access to Event Engine will immediately and automatically terminate if you do not comply with any of these terms or conditions.

Team Hash (e.g. abcdef123456)

This is the 12 digit hash that was given to you or your team.

✓ Invalid Hash

Enter the provided hash in the text box. The button on the bottom right corner changes to **Accept Terms & Login**. Select that button to continue.

Team Dashboard



Event

[Set Team Name](#)

[AWS Console](#)

[SSH Key](#)

Event: reinvent

Select **AWS Console** on dashboard.

AWS Console Login

Remember to only use "us-west-2" as your region, unless otherwise specified.

Login Link



[Open AWS Console](#)



[Copy Login Link](#)

Credentials / CLI Snippets

Mac / Linux

Windows

Mac or Linux A blue clipboard icon.

```
export AWS_DEFAULT_REGION=us-west-2
```

```
-----  
Access Key ID: ACTASDDEDFEFGH  
Secret Access Key: FGHIJKLMNOP1234567890
```

Keep the defaults and select **Open AWS Console**. This will open AWS Console in a new browser tab.

Once you have completed the steps above, you can head straight to the **Region Selection**.

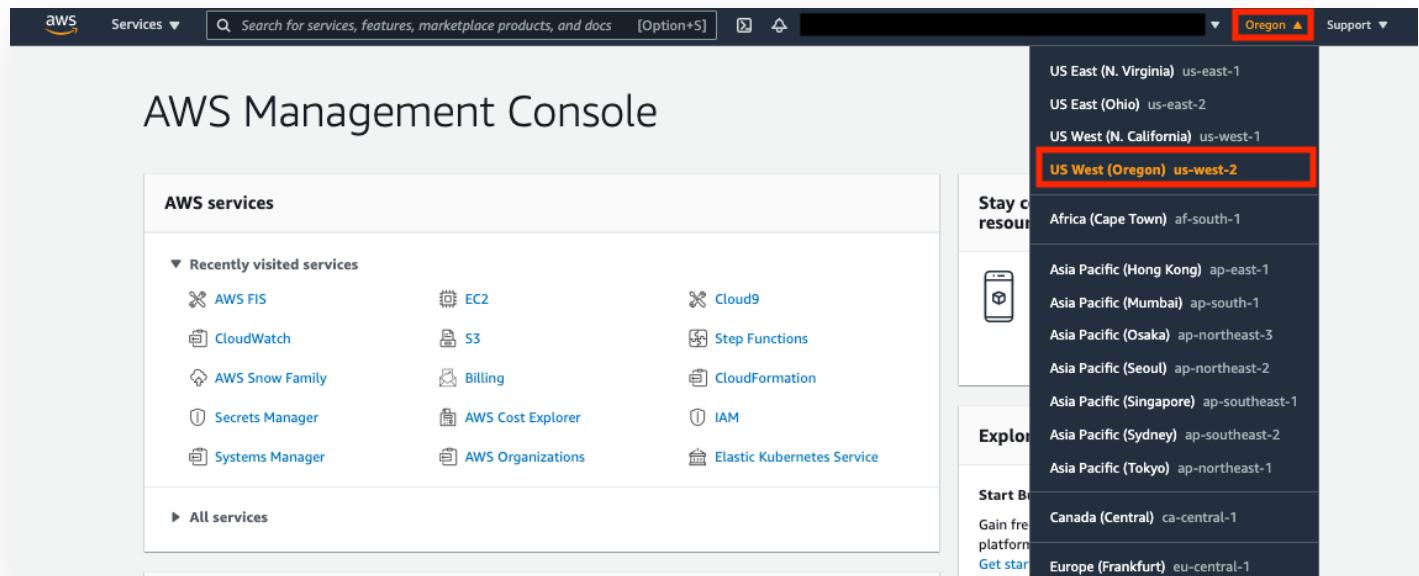
<

>

REGION SELECTION

This workshop relies heavily on AWS Fault Injection Simulator (AWS FIS) and assumes you will be running all your experiments in the same region. If not otherwise instructed, please **choose a region** in which AWS FIS is currently available.

To select this region for navigate to the **AWS Console** and select the desired region from the drop-down menu on the top right:



Note

A number of services, in particular AWS Identity and Access Management (IAM), are region independent and will show "**Global**" as the selection.



CREATE A WORKSPACE

Info

A list of supported browsers for AWS Cloud9 is found [here](#).

Tip

Ad blockers, javascript disablers, and tracking blockers should be disabled for the cloud9 domain, or connecting to the workspace might be impacted. Cloud9 requires third-party-cookies. You can whitelist specific domains by following [these instructions](#).

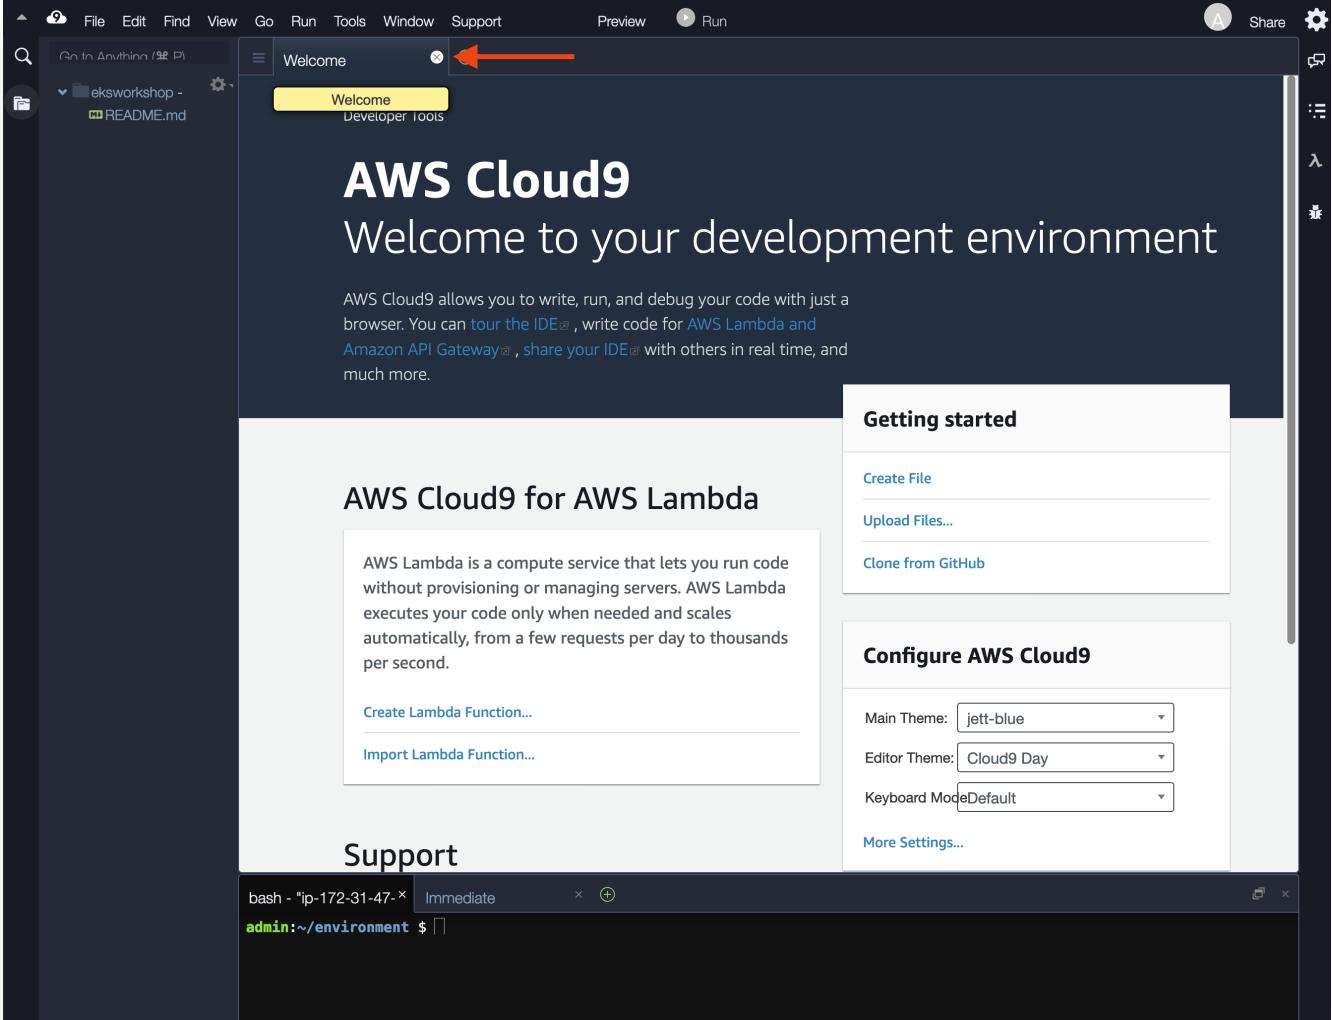
Launch Cloud9 in the region selected previously

Using the region selected in [Region Selection](#), navigate to the [Cloud9 console](#).

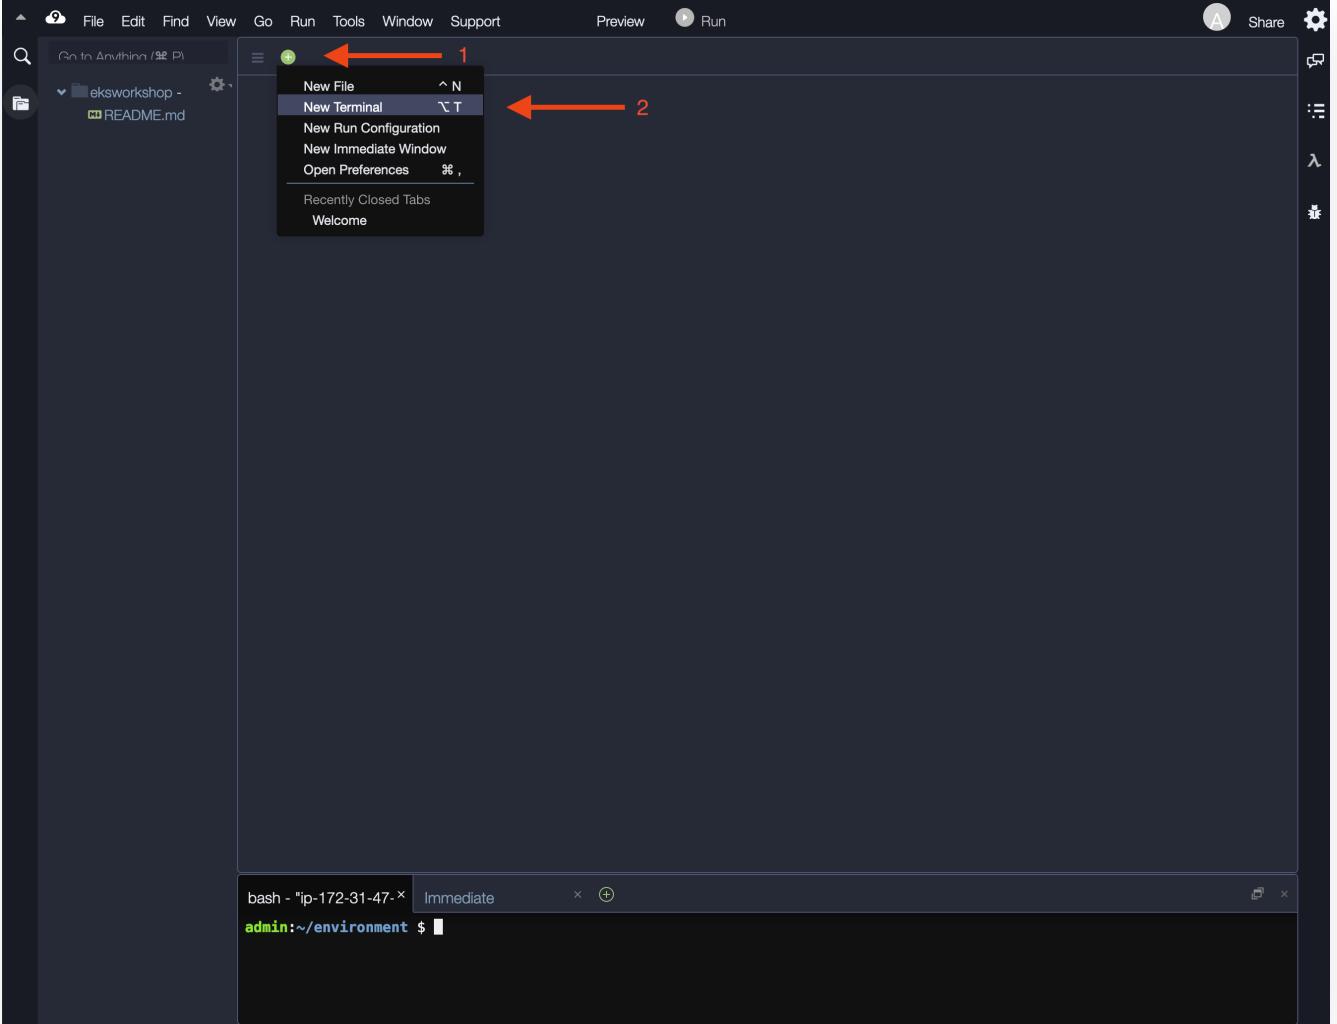
- Select **Create environment**
- Name it **fisworkshop** and select **Next step**.
- Since we only need to access our Cloud9 environment via web browser, please select the **Create a new no-ingress EC2 instance for environment (access via Systems Manager)** under the Environment Type.
- Choose **t3.small** for instance type, go through the wizard with the default values. Finally select **Create environment**

When it comes up, customize the environment by:

- Closing the **Welcome tab**



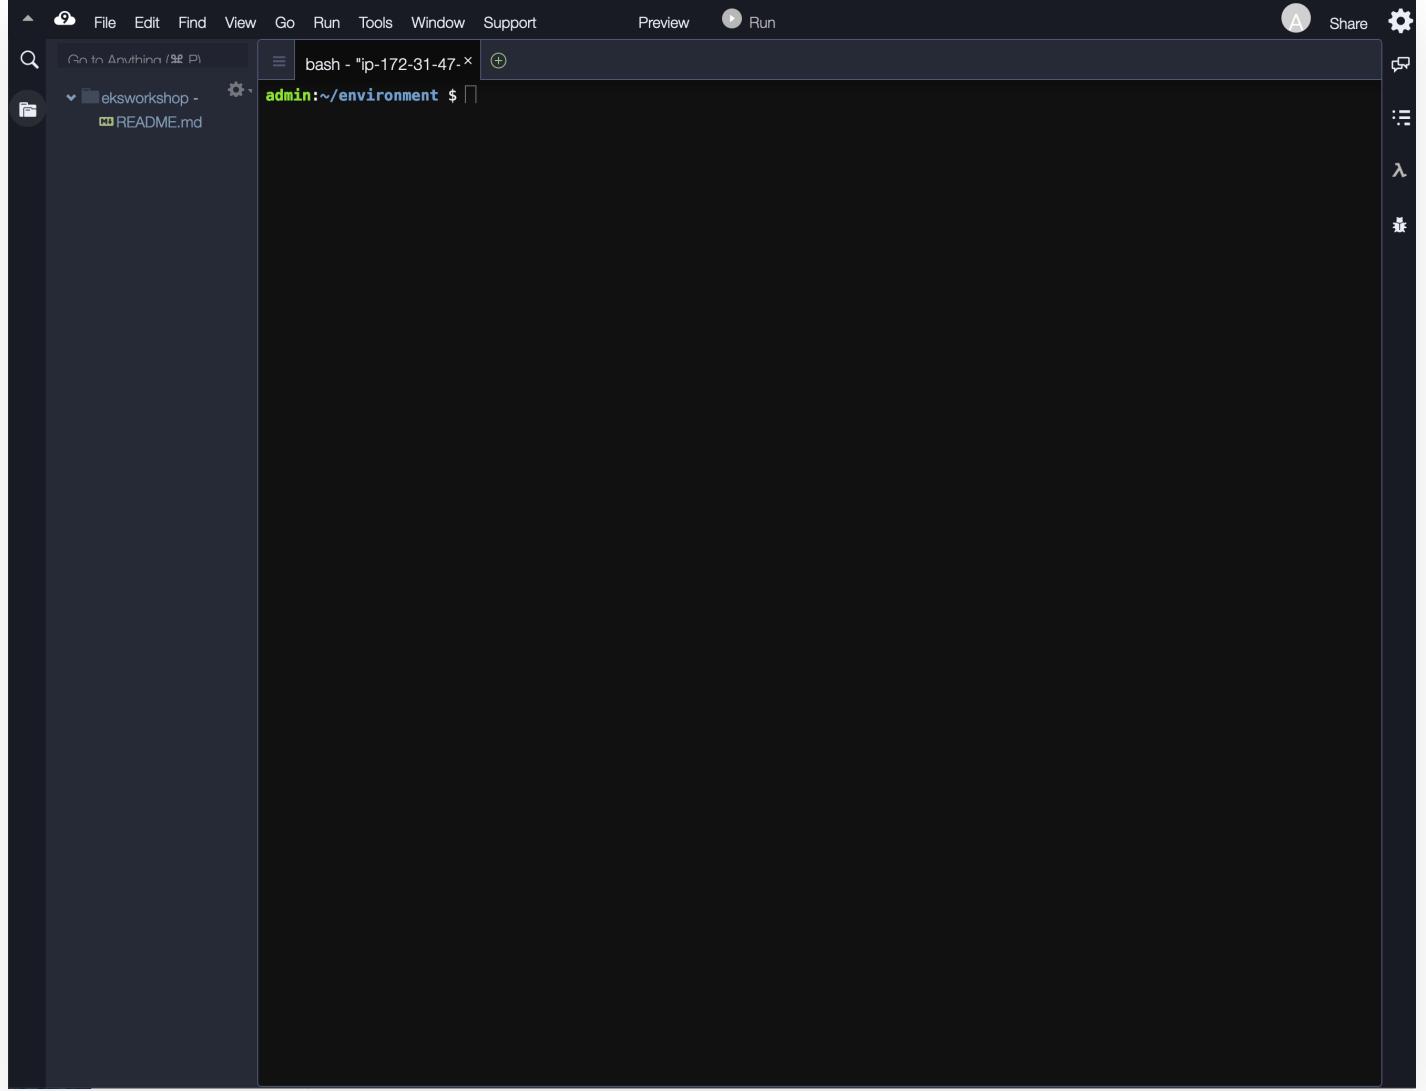
- Opening a new **Terminal** tab in the main work area



- Closing the lower work area

A screenshot of a terminal window within a dark-themed IDE. The terminal title bar reads "bash - "ip-172-31-47-x" and shows the prompt "admin:~/environment \$". The main pane is blank. Below the terminal is a status bar with tabs: "Immediate" (selected), "Preview", and "Run". The status bar also shows the command "(^ Esc)". A red arrow points to the "Run" tab in the status bar.

Your workspace should now look like this



Increase the disk size on the Cloud9 instance

Info

Some commands in this workshop require more than the default disk allocation on an AWS Cloud9 workspace. The following command adds more disk space to the root volume of the EC2 instance that Cloud9 runs on.

Copy/Paste the following code in your Cloud9 terminal (you can paste the whole block at once). Once the command completes, we reboot the instance and it could take a minute or two for the IDE to come back online.

```
# Ensure we have newest boto3 installed
pip3 install --user --upgrade boto3

# Identify instance ID of the Cloud9 environment
export instance_id=$(curl -s http://169.254.169.254/latest/meta-data/instance-id)
```

```

# Use API to identify attached volume and increase size
python -c "import boto3
import os
from botocore.exceptions import ClientError
ec2 = boto3.client('ec2')
volume_info = ec2.describe_volumes(
    Filters=[
        {
            'Name': 'attachment.instance-id',
            'Values': [
                os.getenv('instance_id')
            ]
        }
    ]
)
volume_id = volume_info['Volumes'][0]['VolumeId']
try:
    resize = ec2.modify_volume(
        VolumeId=volume_id,
        Size=30
    )
    print(resize)
except ClientError as e:
    if e.response['Error']['Code'] == 'InvalidParameterValue':
        print('ERROR MESSAGE: {}'.format(e))

# Reboot - on restart the cloud-init will adjust FS size
if [ $? -eq 0 ]; then
    sudo reboot
fi

```

Update tools and dependencies



The instructions in this workshop assume you are using a bash shell in a linux-like environment. They also rely on a number of tools. Follow these instructions to install the required tools in a Cloud9 workspace:

Copy/Paste the following code in your cloud9 terminal (you can paste the whole block at once).

```

# Update to the latest stable release of npm and nodejs.
nvm install stable

# Install typescript
npm install -g typescript

```

```
# Install CDK  
npm install -g aws-cdk  
  
# Install the jq tool  
sudo yum install jq -y
```



PROVISION AWS RESOURCES

⚠ Warning

Only complete this section if you are running the workshop on your own. If you are at an AWS hosted event (such as re:Invent, Kubecon, Immersion Day, etc), these steps have already been executed for you.

Before we start running fault injection experiments we need to provision our resources in the cloud. The rest of the workshop uses these resources.

Clone the repository

```
cd ~/environment  
git clone https://github.com/aws-samples/aws-fault-injection-simulator-workshop.git
```

Deploy the resources

```
cd aws-fault-injection-simulator-workshop  
cd resources/templates  
. ./deploy-parallel.sh
```

It can take up to 30 minutes to complete.

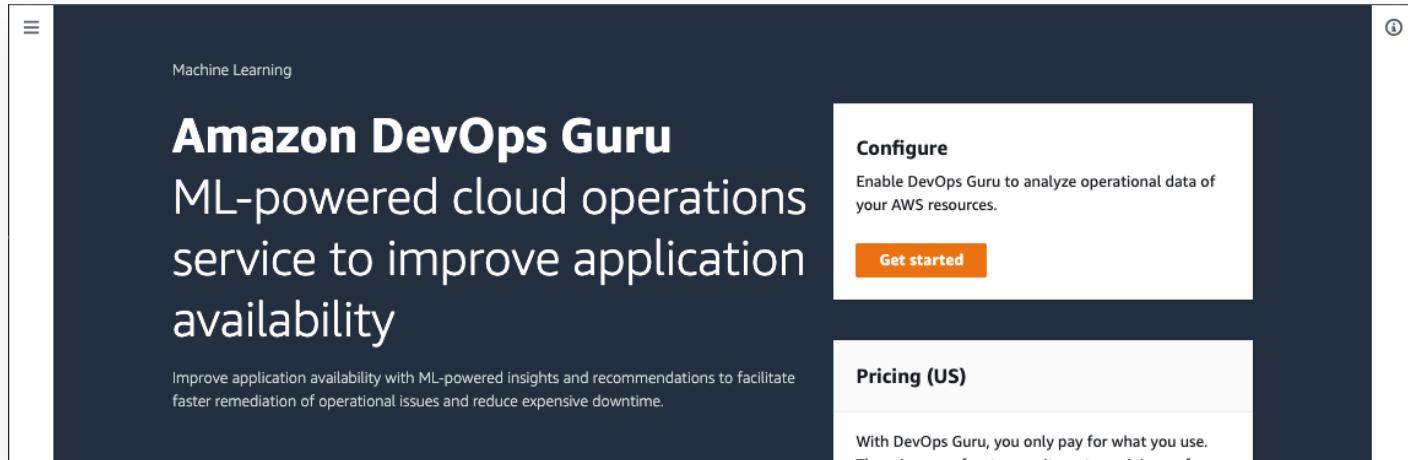


OPTIONAL: SETUP FOR DEVOPS GURU

Warning

Only complete this section if you are planning to explore the DevOps Guru section at the end of the workshop. If you are planning to explore DevOps Guru in this way please allow sufficient time for DevOps Guru to perform initial resource discovery and baselining. Depending on the number of resources in the account/region you select this may take from 2-24h.

Navigate to the [DevOps Guru console](#) and select the “**Get Started**” button:



For “Amazon DevOps Guru analysis coverage” select “**Choose later**” if you will only be exploring as part of this workshop. Otherwise you can select “Analyze all AWS resources in the current AWS account in this Region” but it may take more time and incur more cost to get started.

Amazon DevOps Guru analysis coverage

DevOps Guru analyzes the operational data for your AWS resources based on your selection. You pay for the number of AWS resource hours analyzed, for each active resource. A resource is only active if it produces metrics, events, or log entries within an hour. See the [pricing page](#) for complete details.

Choose which AWS resources to analyze by specifying the coverage boundary

Analyze all AWS resources in the current AWS account in this Region

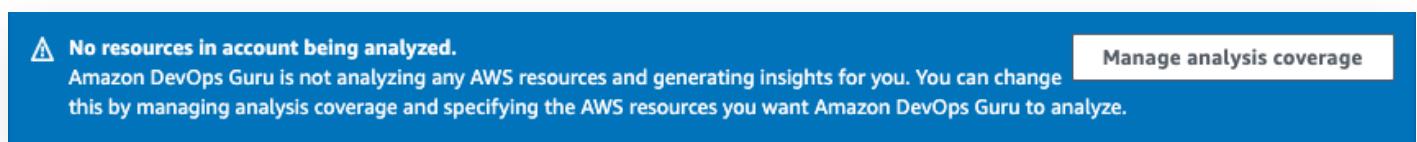
Choose later

You can specify specific AWS resources to analyze using AWS CloudFormation stacks as your coverage boundary.

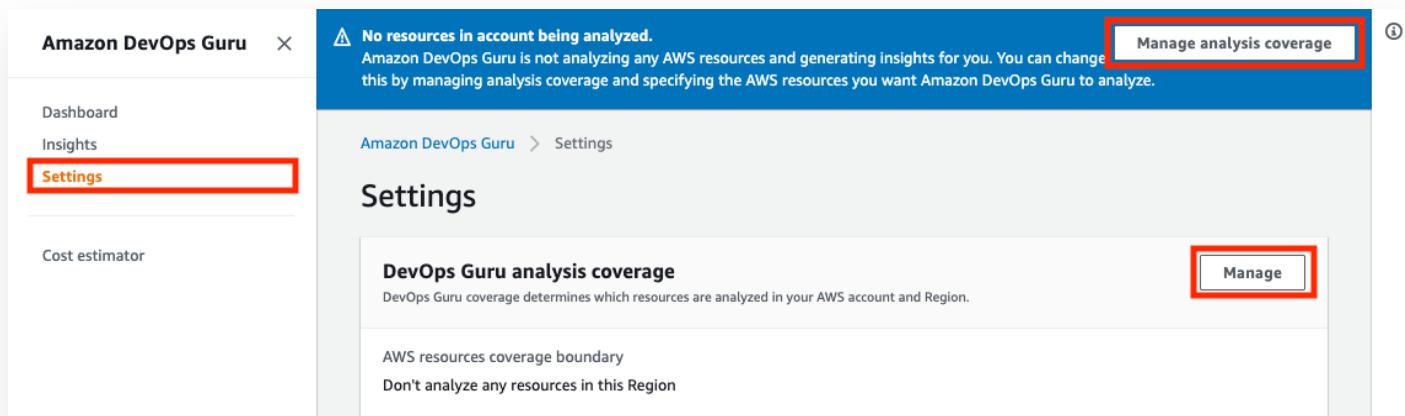
During this workshop we will not be exploring SNS notifications and thus don't need to specify an SNS topic.

Select “**Enable**”.

If you set coverage to “Choose later” you should now see an information banner notifying you that you have not yet selected resources:



Select the “**Manage analysis coverage**” option in the banner or navigate to the [DevOps Guru console](#), choose “**Settings**” and select “**Manage**” option under “DevOps Guru analysis coverage”:



Select all the stacks with names starting with **Fis**:

Manage DevOps Guru analysis coverage

AWS resource selection

Choose DevOps Guru resource coverage

DevOps Guru coverage determines which resources are analyzed in your AWS account and region.

- Analyze all AWS resources in the current AWS account in this Region
- Analyze all AWS resources in the specified CloudFormation stacks in this Region
- Don't analyze any resources in this Region

CloudFormation stacks (4/16)

You can choose up to 500 CloudFormation stacks.

X

4 matches

<

1

>



<input checked="" type="checkbox"/>	Stack name	Description	Status
<input checked="" type="checkbox"/>	FisStackAsg	-	Not enabled
<input checked="" type="checkbox"/>	FisStackRdsAurora	-	Not enabled
<input checked="" type="checkbox"/>	FisStackLoadGen	-	Not enabled
<input checked="" type="checkbox"/>	FisStackVpc	-	Not enabled

CancelSave

Select “Save”.



WORKSHOP

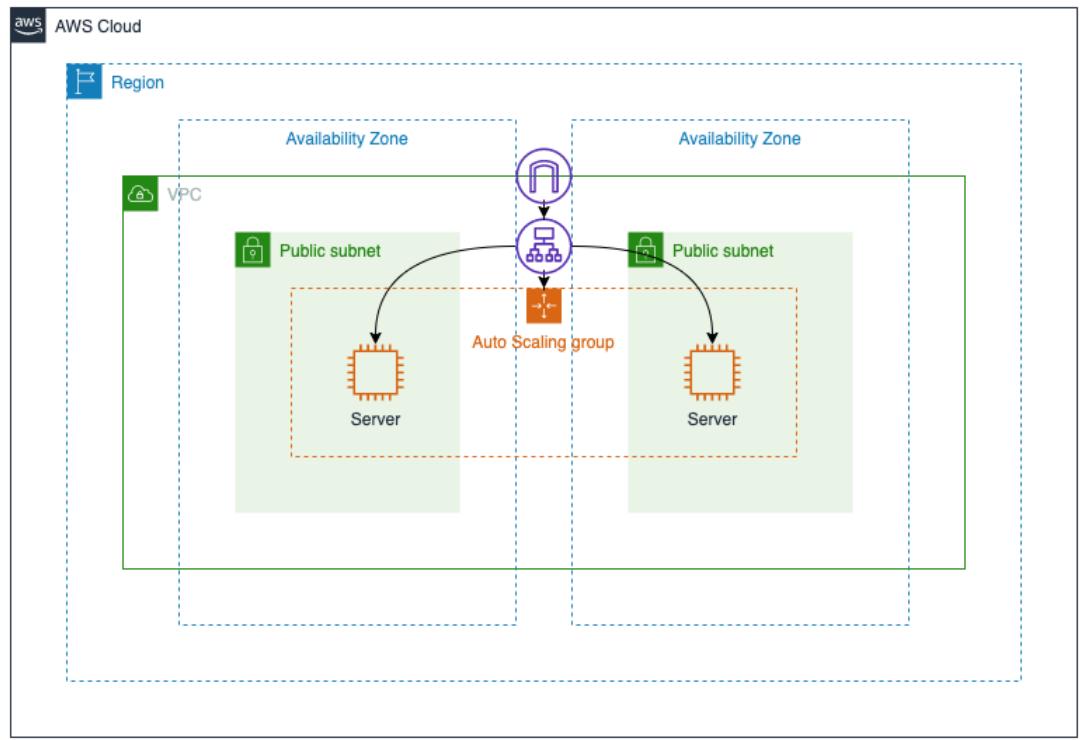
This workshop is broken into multiple chapters. The chapters are designed to be done in sequence with each chapter assuming familiarity with some concepts from previous chapters and focusing on new learnings. We include refresher links to relevant prior sections to help you skip over materials you are already familiar with.

Chapters:

- Baselining and Monitoring
- Synthetic User Experience
- First Experiment
- AWS Systems Manager Integration
- Databases
- Containers
- Serverless
- Recurrent Experiments - CI/CD
- Common scenarios
- Observability
- Cleanup

Architecture Diagrams

This workshop is focused on how to inject fault into an existing infrastructure. For this purpose the template in the **Provision AWS resources** section sets up a variety of components. Throughout this workshop we will be showing you architecture diagrams focusing on only the components relevant to the section, e.g.:



You can click on these images to enlarge them.

- › Click to expand if you are hosting a demo



BASELINING AND MONITORING

Before we start injecting faults into our system we should consider the following thought experiment:

"If a tree falls in a forest and no one is around to hear it, does it make a sound?"

For the purpose of our fault injection experiments we can rephrase this in two ways:

"If part of our system is disrupted and we do not receive any irate calls from users, did anything break?"

"If part of our system is disrupted and sysops isn't alerted, did anything break?"

Think about this for a second. There is a distinct difference between those two statements because users and ops have very different experiences.

What the users see

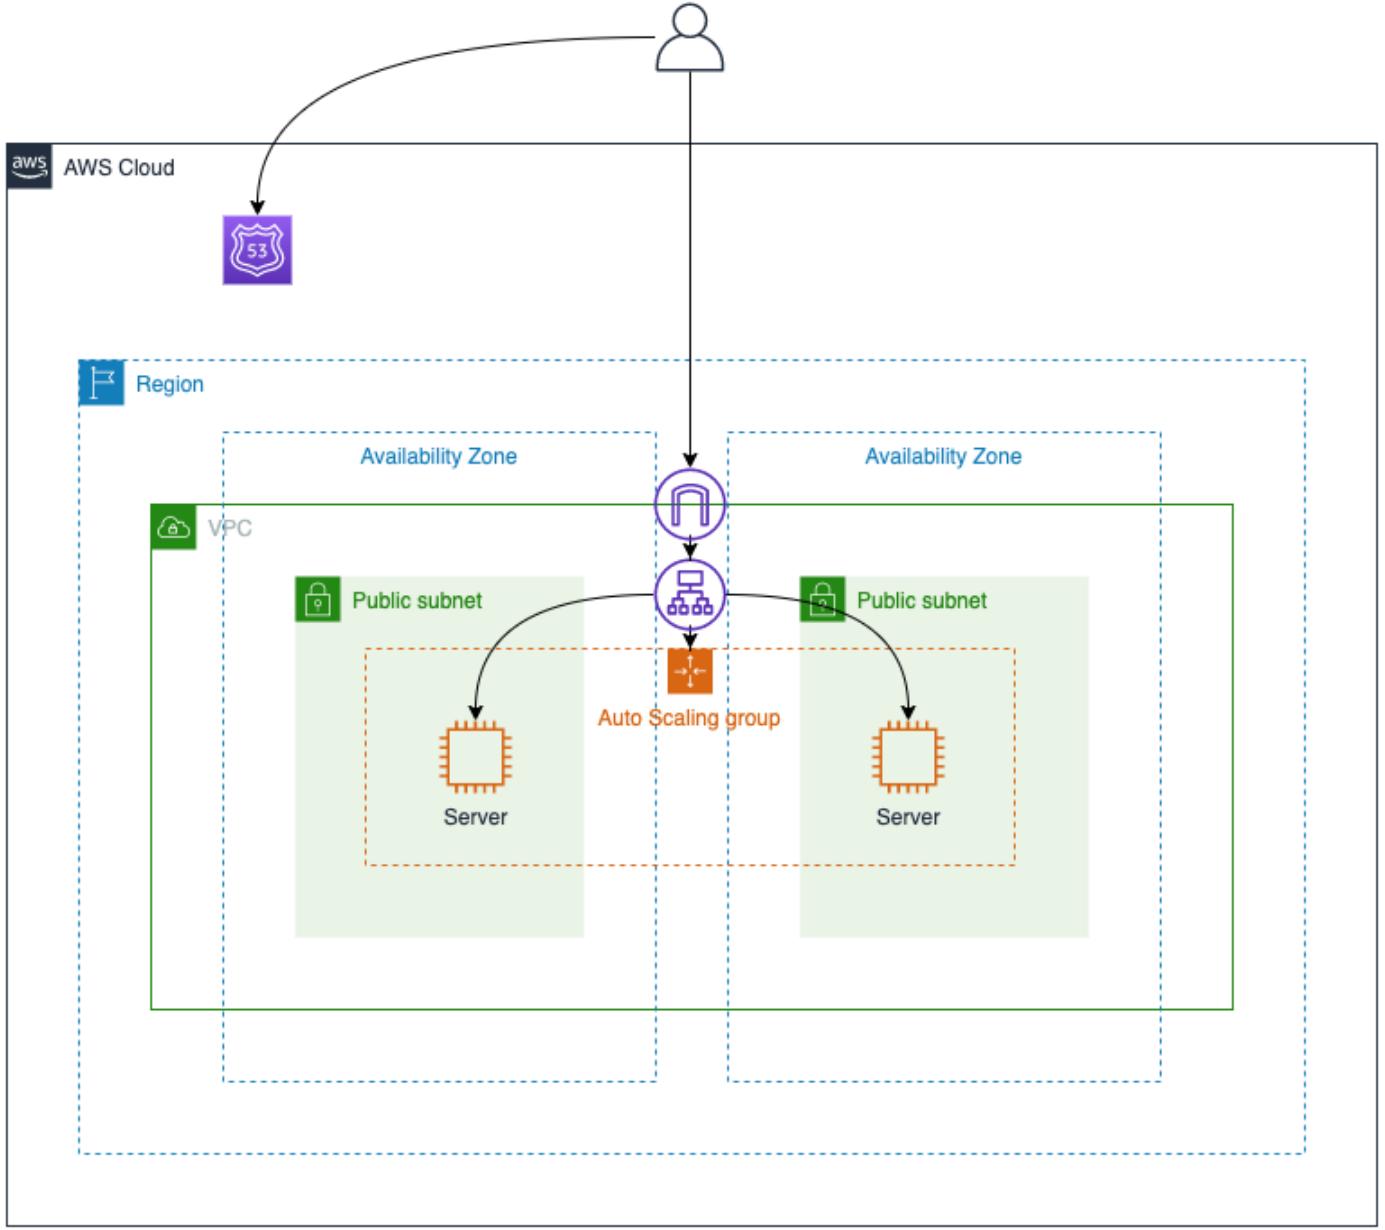
What the users see is immediate, e.g. the website not loading or loading slowly. What the users see is also an end-to-end test of all system components, and not all components of the system are in your purview, e.g. you cannot see the speed of the users' network connection or the state of their DNS caches. Finally an individual user can have an experience entirely different from all other users. For this workshop, this is particularly important for a particular edge case: developers and ops typically have better system configurations and better experiences than the average user but tend to rely on the anecdotal evidence of "it worked for me".

What sysops sees

Typically, what sysops see is a wealth of individual health and performance indicators. These often grow organically over time and especially after outages. Even where dashboards have been built with overall system health in mind, the metrics are delayed against the user experience and aggregate over the experience of many users, requiring extra effort to notice poor experiences specific to a subset of users.

Setting up for fault injection

Before starting our first fault injection experiment, let's take a look at our most basic infrastructure:

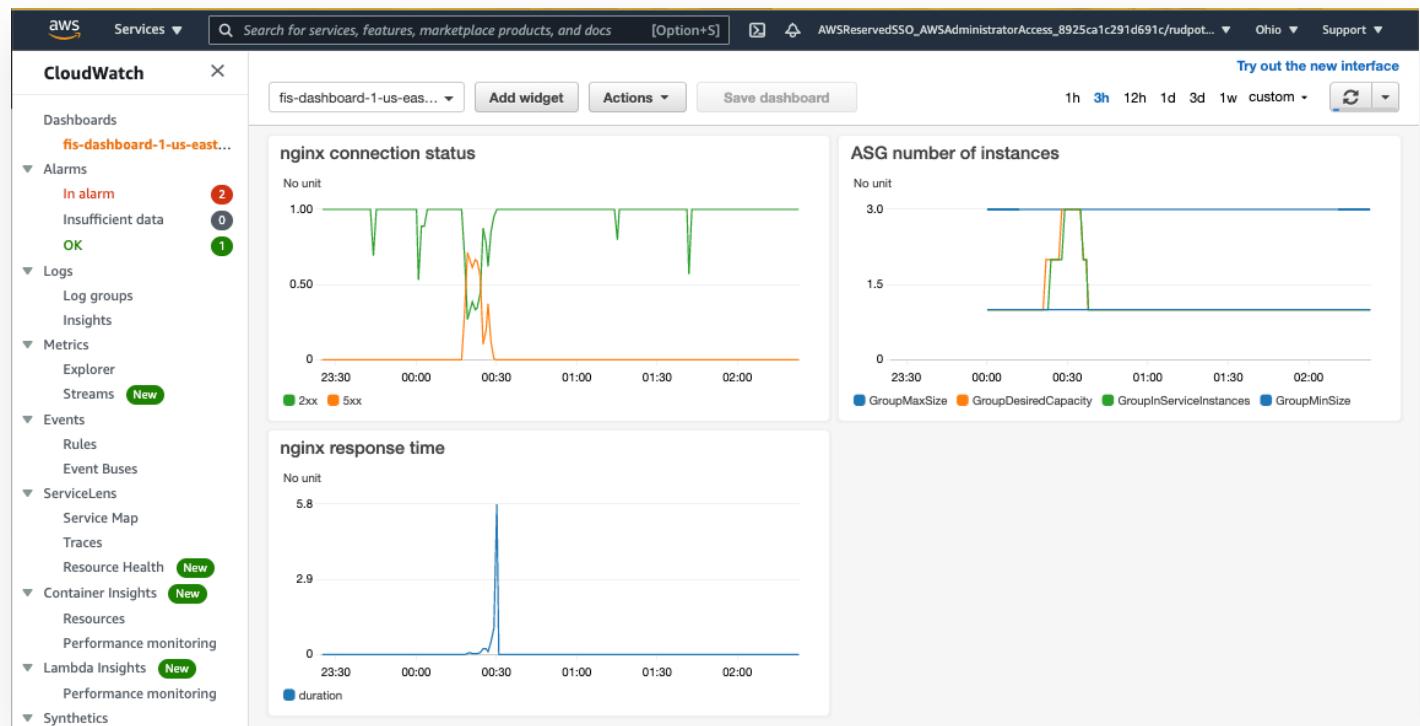


We have a user trying to access a website running on AWS. We have designed it for high availability. We used EC2 instances with an Auto Scaling group and a load balancer to ensure that users can always reach our website even under heavy load or if an instance suddenly fails.

Once you've created the resources as described in [Provision AWS resources](#) you can navigate to [CloudFormation](#), select the `FisStackAsg` stack and select the **“Outputs”** tab which will show you the server URL:

Outputs (1)			
Key		Description	Export name
URL	http://fis-a-Appli-DYENM5VTHFYU-1588384914.us-west-2.elb.amazonaws.com	The URL of the website	-

To gain visibility into the user experience from the sysops side we've used the **AWS CloudWatch agent** to export our web server logs to **AWS CloudWatch Logs** and we created **AWS CloudWatch Logs metrics filters** to track server response codes and speeds on a **dashboard**. Note that the dashboard's name is based on the region in which we deployed. If you chose a region other than **us-west-2** the dashboard's name will be different. The dashboard also shows the number of instances in our Auto Scaling Group (ASG).



› Accessing the dashboard from the console

In the next section we will cover how to measure the user experience.

<

>

SYNTHETIC USER EXPERIENCE

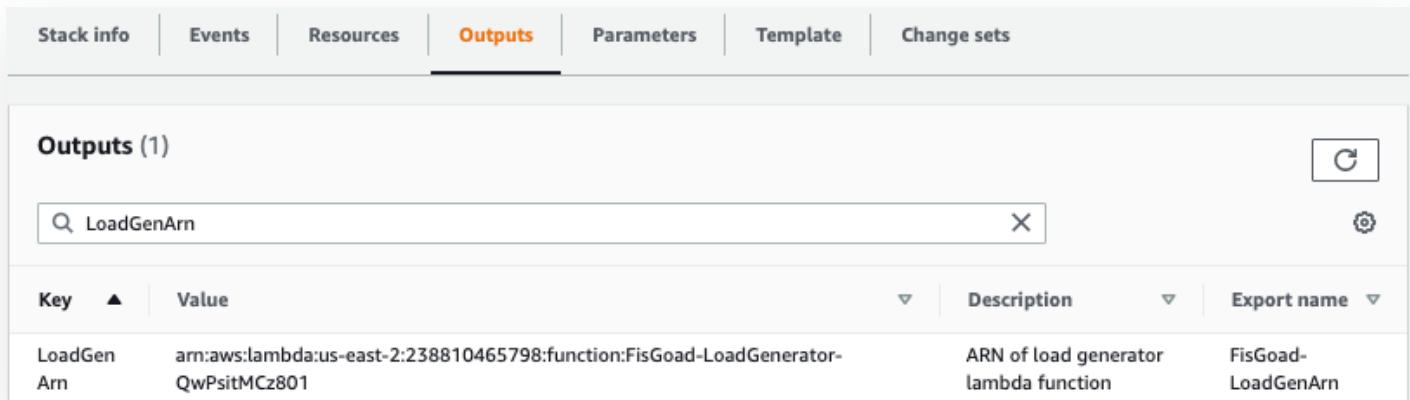
To gain insights from our fault injection experiments, we want to correlate user-experience with the sysops view from the previous section. In production, we could instrument the clients to send telemetry back to us, but in non-production we don't usually have sufficient load to do this. You also probably have better things to do than sit there clicking reload on a browser page while your experiment is running.

Generating load against our website

In the previous section, you navigated to the basic website setup as well as the sysops performance dashboard. Open a linux terminal and save the URL from the previous page in an environment variable:

```
export URL_HOME=[PASTE URL HERE]
```

Next, we need to generate load. There are many **load testing tools** available to generate a variety of load patterns. However, for the purpose of this workshop we have included a Lambda function that will make HTTP GET calls to our website and log performance data to CloudWatch. To find the Lambda function, navigate to **CloudFormation**, select the **FisStackLoadGen** stack, and click on the "**Outputs**" tab. It will show you the Lambda function ARN:



The screenshot shows the AWS CloudFormation Outputs tab. The tab bar includes Stack info, Events, Resources, Outputs (which is highlighted in orange), Parameters, Template, and Change sets. Below the tab bar, the title "Outputs (1)" is displayed. A search bar contains the text "LoadGenArn". The main table has columns for Key, Value, Description, and Export name. One row is shown with the Key "LoadGenArn" having the Value "arn:aws:lambda:us-east-2:238810465798:function:FisGoad-LoadGenerator-QwPsitMCz801", the Description "ARN of load generator lambda function", and the Export name "FisGoad-LoadGenArn".

Key	Value	Description	Export name
LoadGenArn	arn:aws:lambda:us-east-2:238810465798:function:FisGoad-LoadGenerator-QwPsitMCz801	ARN of load generator lambda function	FisGoad-LoadGenArn

Save the Lambda function ARN in another environment variable:

```
export LAMBDA_ARN=[PASTE ARN HERE]
```

Finally, invoke the Lambda function using the AWS CLI:

```
aws lambda invoke \
--function-name ${LAMBDA_ARN} \
--payload "{
  \"ConnectionTargetUrl\": \"${URL_HOME}\",
  \"ExperimentDurationSeconds\": 180,
  \"ConnectionsPerSecond\": 1000,
  \"ReportingMilliseconds\": 1000,
  \"ConnectionTimeoutMilliseconds\": 2000,
  \"TlsTimeoutMilliseconds\": 2000,
  \"TotalTimeoutMilliseconds\": 2000
}"
--invocation-type Event \
invoke.txt
```

⚠️ Warning

If you are running AWS CLI v2, you need to pass the parameter

`--cli-binary-format raw-in-base64-out` or you'll get the error "Invalid base64" when sending the payload.

Now, let's generate some load. The invocation above will generate 1000 connections per second for 3 minutes. We expect our website's performance to degrade and for Auto Scaling to kick in.

Explore impact of load

While our load is running let's explore the setup a little more.

Webserver logs and metrics

The first thing we want to look at is our webserver logs. Because we are using AWS Auto Scaling, virtual machines can be terminated and recycled which means logs written locally on the EC2 instance won't be accessible anymore. Therefore, we have installed the **Unified CloudWatch Agent** and configured our webserver to write logs to a **CloudWatch Log Group**.

› Navigating to CloudWatch Log Groups

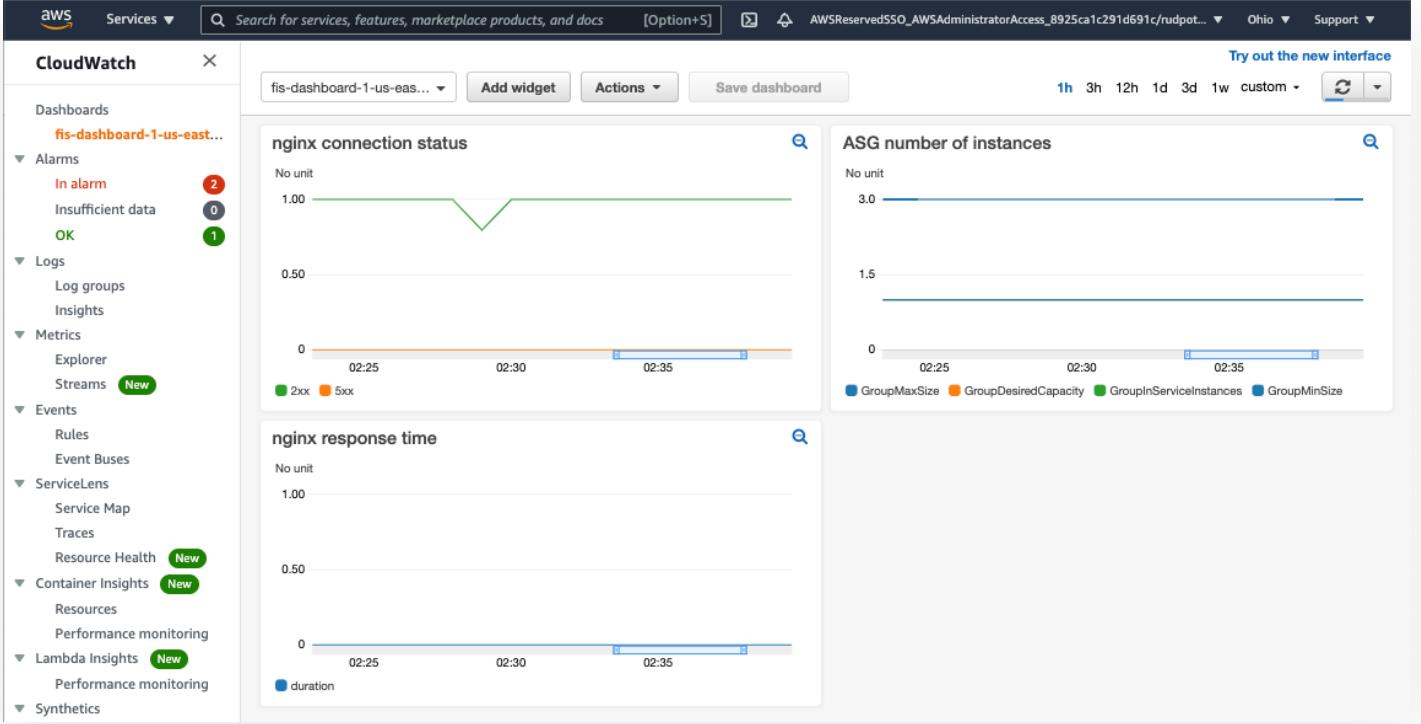
Log streams (3)		<input type="button" value="C"/>	Delete	Create log stream	Search all
<input type="text"/> Filter log streams or try prefix search			< 1 >		
<input type="checkbox"/>	Log stream	Last event time			
<input type="checkbox"/>	i-08256301c14f8b6c3	2021-05-21 19:53:21 (UTC-06:00)			
<input type="checkbox"/>	i-0ccfa12ad27166f2c	2021-05-21 18:38:13 (UTC-06:00)			
<input type="checkbox"/>	i-0f631388d3a092c74	2021-05-21 18:36:22 (UTC-06:00)			

Click through on the topmost entry and expand any of the log lines. You may notice that we've modified the Nginx output format to use JSON instead of the default format:

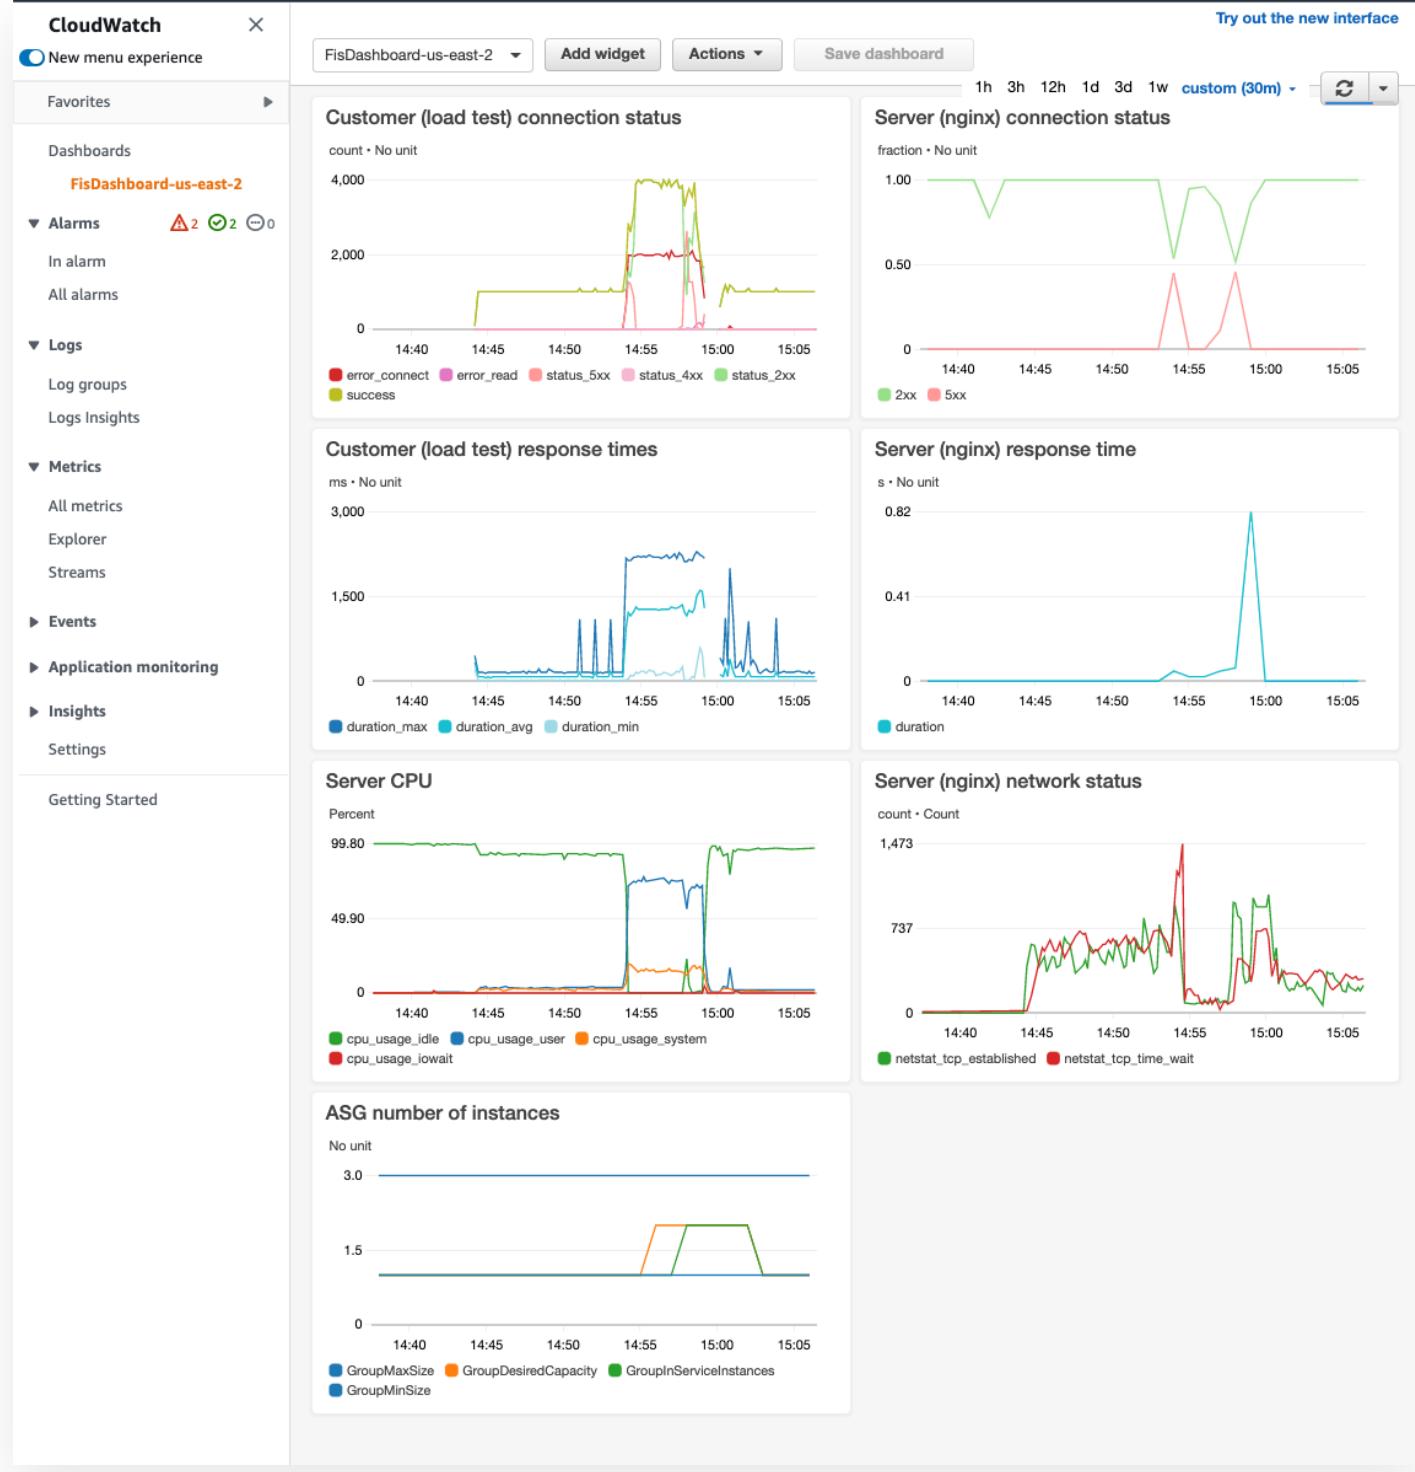
Log events	
You can use the filter bar below to search for and match terms, phrases, or values in your log events. Learn more about filter patterns	
<input type="checkbox"/> View as text	<input type="button" value="C"/>
<input type="text"/> Filter events	
<input type="button" value="Actions"/>	<input type="button" value="Create Metric Filter"/>
<input type="button" value="Clear"/>	1m 30m 1h 12h Custom <input type="button" value=""/>
Timestamp	Message
There are older events to load. Load more .	
▼ 2021-05-22T01:47:51.362Z	{"time_local": "22/May/2021:01:47:46 +0000", "remote_addr": "10.0.0.169", "remote_user": "", "request": "GET / ..."} {"time_local": "22/May/2021:01:47:46 +0000", "remote_addr": "10.0.0.169", "remote_user": "", "request": "GET / HTTP/1.1", "status": "200", "body_bytes_sent": 3520, "request_time": 0, "http_referer": "", "http_user_agent": "ELB-HealthChecker/2.0"} <input type="button" value="Copy"/>
▶ 2021-05-22T01:48:05.362Z	{"time_local": "22/May/2021:01:48:01 +0000", "remote_addr": "10.0.0.169", "remote_user": "", "request": "GET / ..."} {"time_local": "22/May/2021:01:48:01 +0000", "remote_addr": "10.0.0.169", "remote_user": "", "request": "GET / HTTP/1.1", "status": "200", "body_bytes_sent": 3520, "request_time": 0, "http_referer": "", "http_user_agent": "ELB-HealthChecker/2.0"}

While not necessary, this makes it easy to create **Metric Filters**. Navigate back to the `/fis-workshop/asg-access-log` log group and select the "**Metric filters**" tab. You will see that we have created filters to extract the count of responses with HTTP `status` codes in the `2xx` (good responses) and `5xx` (bad responses) ranges. We also created a filter to select all entries that have a `request_time` set. The resulting metrics can be found under **Metrics / All metrics / Custom Namespaces / fisworkshop**. These are also the metrics for `Server (nginx) connection status` and `Server (nginx) response time` you saw on the dashboard in the previous section.

Let's look at our dashboard:



That's odd, did anything happen? According to Nginx, it looks like nothing happened. Remember the falling tree in the forest and no one is around to hear it? We need to look at what the server CPU and the load runner. For this, we have added a more detailed dashboard:



Now, it's clearer what happened. We were requesting a small static page and Nginx is really efficient. In the [Server CPU](#) graph, we can see minimal CPU utilization correlating with the load data in the [Customer \(load test\)](#) graphs.

Increasing the load

Clearly, hitting a static page isn't a good test to validate our Auto Scaling configuration works as intended. Fortunately, the server also exposes a [phpinfo.php](#) page. Let's try loading that instead. Define another

environment variable and run the load test against the new URL. Since we want to see the Auto Scaling group adjust capacity, let's run more than one copy:

```
export URL_PHP=${URL_HOME}/phpinfo.php

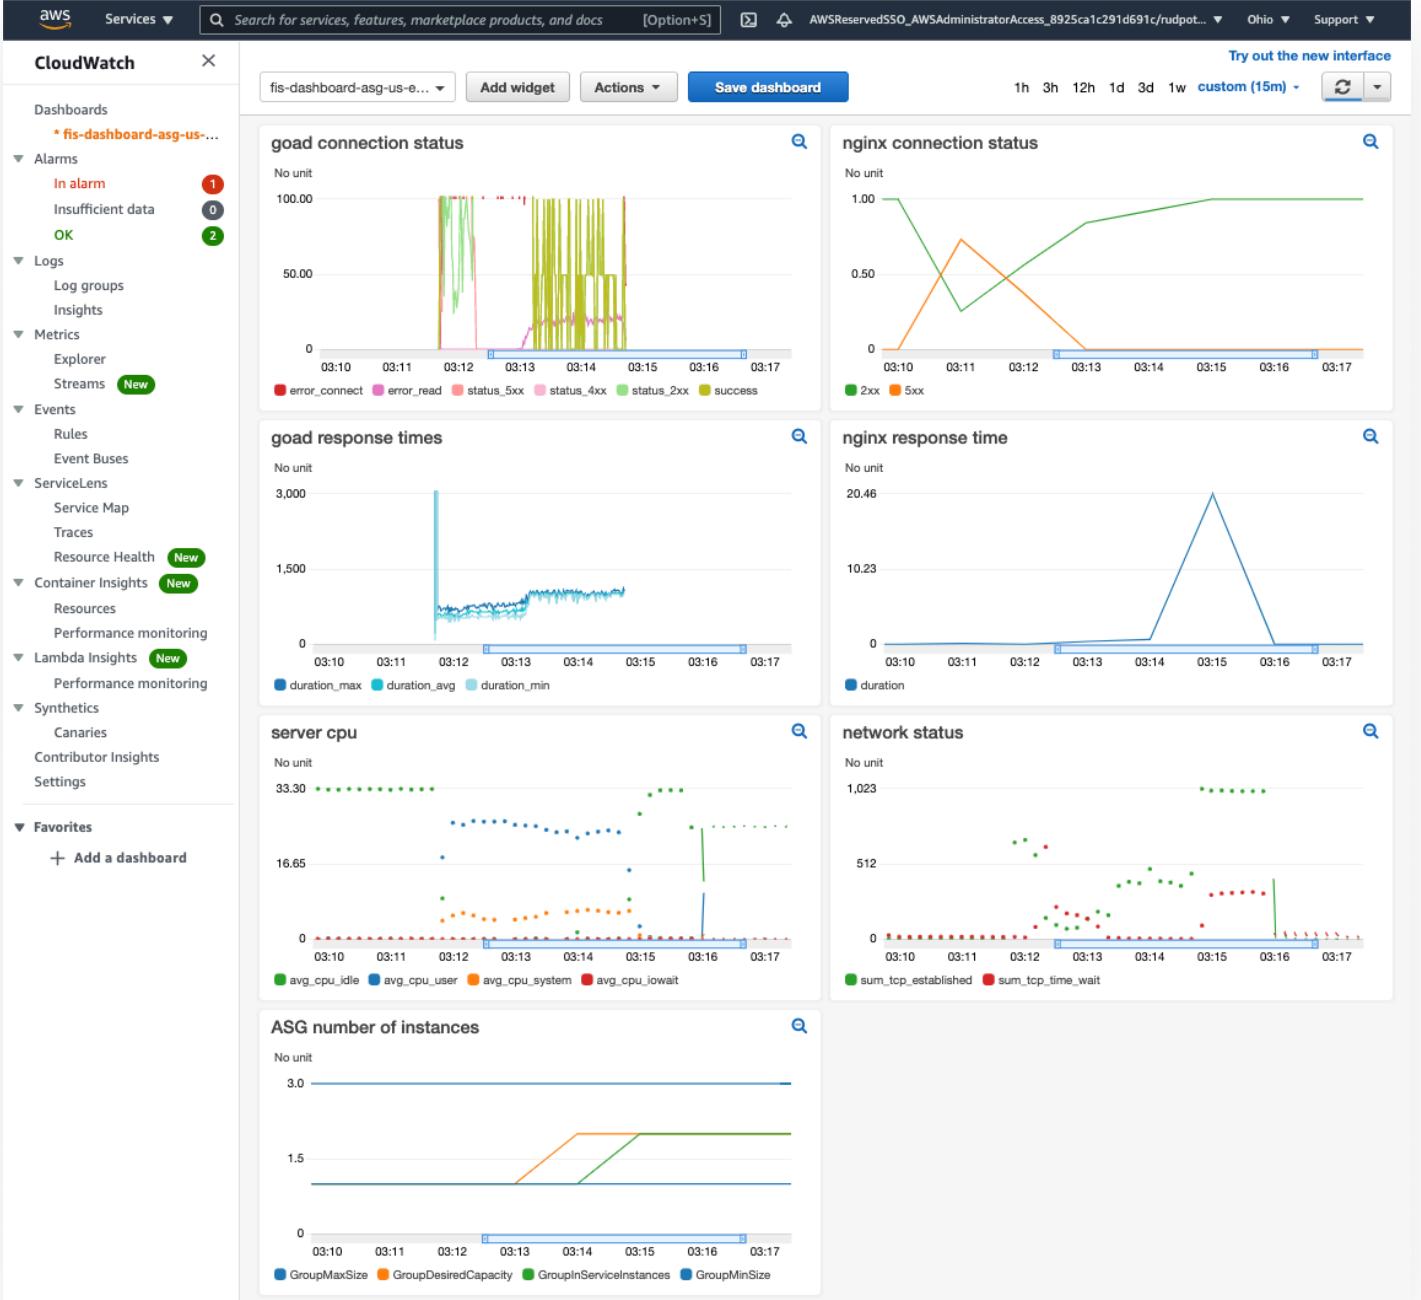
for ii in 1 2 3; do
  aws lambda invoke \
    --function-name ${LAMBDA_ARN} \
    --payload "{
      \"ConnectionTargetUrl\": \"${URL_PHP}\",
      \"ExperimentDurationSeconds\": 300,
      \"ConnectionsPerSecond\": 1000,
      \"ReportingMilliseconds\": 1000,
      \"ConnectionTimeoutMilliseconds\": 2000,
      \"TlsTimeoutMilliseconds\": 2000,
      \"TotalTimeoutMilliseconds\": 2000
    }" \
    --invocation-type Event \
    invoke-$ii.txt
done
```

Warning

If you are running AWS CLI v2, you need to pass the parameter

`--cli-binary-format raw-in-base64-out` or you'll get the error "Invalid base64" when sending the payload.

While this is executing, we encourage you to explore CloudWatch logs and create some dashboard views of your own.



According to the dashboards, we've now generated enough load to force a scaling event. We can also see how different the user experience is from the Nginx report. Requests timeout after 2s, substantially affecting user experiences, and rendering the website unavailable. Nginx, in contrast, doesn't report this as an error because the connection was terminated by the client before being served. We will leave it as an exercise to the reader to figure out more details and will move on to fault injection experiments.



FIRST EXPERIMENT

In this section, we will cover the setup required for using AWS FIS to run our first fault injection experiment

Experiment idea

In the previous section, we ensured that we can measure the user experience. We also have configured an Auto Scaling group that should make sure we can “always” provide a good experience to the customer. Let’s validate this:

- **Given:** we have an Auto Scaling group with multiple instances
- **Hypothesis:** Failure of a single EC2 instance may lead to slower response times but should not affect service availability for our customers.



CONFIGURING PERMISSIONS

The AWS FIS security model uses two IAM roles. The first IAM role, the one you used to log into the console, controls access to AWS FIS service. It governs whether you are able to see, modify, and run AWS FIS experiments.

The second role governs what resources an AWS FIS experiment can affect during a fault injection experiment. For the purposes of this workshop, we will create one generic role. However, you can create fine grained IAM roles for each fault injection experiment.

Create FIS service role

We need to create an **IAM role for the AWS FIS service** to grant it permissions to inject faults into the system. While we could have pre-created this IAM role for you, we think it is important to review its scope with you.

Navigate to the **IAM console** and create a new policy called **FisWorkshopServicePolicy**. On the "Create Policy" page select the **JSON** tab

The screenshot shows the 'Create policy' page in the AWS IAM console. At the top, there are three tabs: 'Visual editor' (disabled), 'JSON' (selected), and 'Import managed policy'. Below the tabs, a JSON code editor displays a policy document. The code is as follows:

```
1 "Version": "2012-10-17",
2 "Statement": [
3     {
4         "Sid": "AllowFISExperimentRoleReadOnly",
5         "Effect": "Allow",
6         "Action": [
7             "ec2:DescribeInstances",
8             "ecs:DescribeClusters"
9         ]
10    }
11]
```

and paste the following policy. Take the time to look at how broad these permissions are:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "AllowFISExperimentRoleReadOnly",
      "Effect": "Allow",
      "Action": [
        "ec2:DescribeInstances",
        "ecs:DescribeClusters",
        "lambda:ListFunctions",
        "lambda:InvokeFunction",
        "lambda:GetFunctionConfiguration",
        "lambda:UpdateFunctionConfiguration",
        "lambda:UpdateFunctionCode",
        "lambda:DeleteFunction",
        "lambda:CreateFunction",
        "lambda:ListTags",
        "lambda:TagResource",
        "lambda:UntagResource",
        "lambda:ListResolvers",
        "lambda:CreateResolver",
        "lambda:UpdateResolver",
        "lambda:DeleteResolver",
        "lambda:ListResolversByFunction",
        "lambda:ListTagsForResource",
        "lambda:ListEventSources",
        "lambda:PutEventSourceMapping",
        "lambda:DeleteEventSourceMapping",
        "lambda:ListEventSourceMappings",
        "lambda:ListEventTypes",
        "lambda:CreateEventSourceMapping",
        "lambda:UpdateEventSourceMapping",
        "lambda:DeleteEventSource",
        "lambda:ListEventSources",
        "lambda:ListFunctions",
        "lambda:InvokeFunction",
        "lambda:GetFunctionConfiguration",
        "lambda:UpdateFunctionConfiguration",
        "lambda:UpdateFunctionCode",
        "lambda:DeleteFunction",
        "lambda:CreateFunction",
        "lambda:ListTags",
        "lambda:TagResource",
        "lambda:UntagResource",
        "lambda:ListResolvers",
        "lambda:CreateResolver",
        "lambda:UpdateResolver",
        "lambda:DeleteResolver",
        "lambda:ListResolversByFunction",
        "lambda:ListTagsForResource",
        "lambda:ListEventSources",
        "lambda:PutEventSourceMapping",
        "lambda:DeleteEventSourceMapping",
        "lambda:ListEventSourceMappings",
        "lambda:ListEventTypes",
        "lambda:CreateEventSourceMapping",
        "lambda:UpdateEventSourceMapping",
        "lambda:DeleteEventSource",
        "lambda:ListEventSources"
      ]
    }
  ]}
```

```
        "ecs>ListContainerInstances",
        "eks>DescribeNodegroup",
        "iam>ListRoles",
        "rds>DescribeDBInstances",
        "rds>DescribeDbClusters",
        "ssm>ListCommands"
    ],
    "Resource": "*"
},
{
    "Sid": "AllowFISExperimentRoleEC2Actions",
    "Effect": "Allow",
    "Action": [
        "ec2>RebootInstances",
        "ec2>StopInstances",
        "ec2>StartInstances",
        "ec2>TerminateInstances"
    ],
    "Resource": "arn:aws:ec2:*:*:instance/*"
},
{
    "Sid": "AllowFISExperimentRoleECSActions",
    "Effect": "Allow",
    "Action": [
        "ecs>UpdateContainerInstancesState",
        "ecs>ListContainerInstances"
    ],
    "Resource": "arn:aws:ecs:*:*:container-instance/*"
},
{
    "Sid": "AllowFISExperimentRoleEKSActions",
    "Effect": "Allow",
    "Action": [
        "ec2>TerminateInstances"
    ],
    "Resource": "arn:aws:ec2:*:*:instance/*"
},
{
    "Sid": "AllowFISExperimentRoleFISActions",
    "Effect": "Allow",
    "Action": [
        "fis>InjectApiInternalError",
        "fis>InjectApiThrottleError",
        "fis>InjectApiUnavailableError"
    ],
    "Resource": "arn:*fis*experiment/*"
},
{
    "Sid": "AllowFISExperimentRoleRDSReboot",
    "Effect": "Allow",
    "Action": [
        "rds>RebootDBInstance"
    ],
    "Resource": "arn:aws:rds:*:*:db:*
```

```

    "Sid": "AllowFISExperimentRoleRDSFailOver",
    "Effect": "Allow",
    "Action": [
        "rds:FailoverDBCluster"
    ],
    "Resource": "arn:aws:rds:*.*:cluster:*
```

},

{

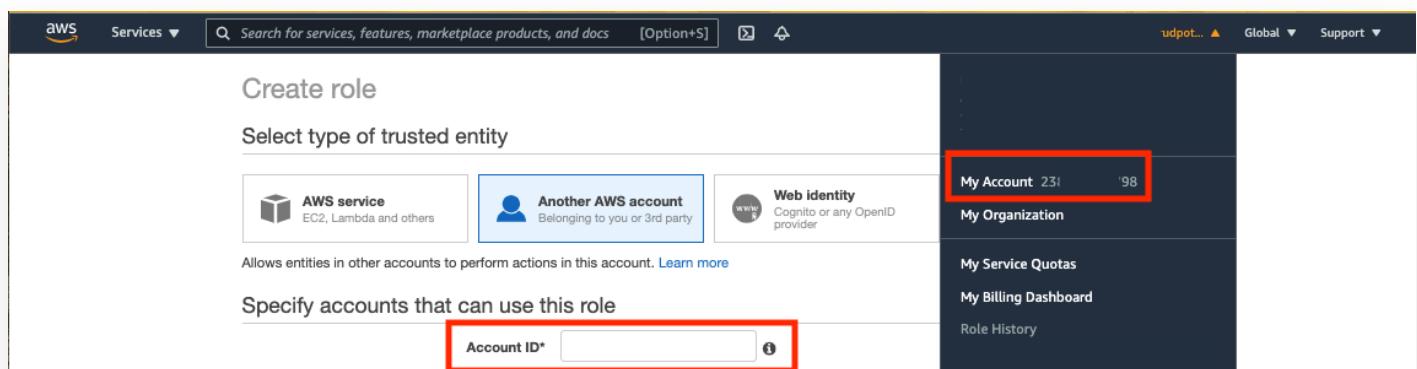
```

    "Sid": "AllowFISExperimentRoleSSMSendCommand",
    "Effect": "Allow",
    "Action": [
        "ssm:SendCommand"
    ],
    "Resource": [
        "arn:aws:ec2:.*.*:instance/*",
        "arn:aws:ssm:.*.*:document/*"
    ]
},
{
    "Sid": "AllowFISExperimentRoleSSMCancelCommand",
    "Effect": "Allow",
    "Action": [
        "ssm:CancelCommand"
    ],
    "Resource": "*"
}
]
}

```

Navigate to the **IAM console** and create a new role called **FisworkshopServiceRole**.

On the "Select type of trusted entity" page AWS FIS does not exist as a trusted service yet. Select "**Another AWS Account**" and add the current account number. You can find the account number in the drop-down menu as shown:



On the "Attach permissions" page search for the **FisWorkshopServicePolicy** we just created and check the box beside it to attach it to the role.

Create role

1 2 3 4

Attach permissions policies

Choose one or more policies to attach to your new role.

Create policy

Filter policies ▾

Showing 1 result

	Policy name ▾	Used as
<input checked="" type="checkbox"/>	FisWorkshopServicePolicy	None

Back in the **IAM Roles** page, find and edit the **FisWorkshopServiceRole**. Select “**Trust relationships**” and the “**Edit trust relationship**” button.

Identity and Access Management (IAM)

Roles > FisWorkshopServiceRole

Summary

Delete role

Role ARN: arn:aws:iam::23 98:role/FisWorkshopServiceRole

Role description: FIS service role | Edit

Instance Profile ARNs:

Path: /

Creation time: 2021-05-28 14:09 MDT

Last activity: Not accessed in the tracking period

Maximum session duration: 1 hour | Edit

Give this link to users who can switch roles in the console: <https://signin.aws.amazon.com/switchrole?roleName=FisWorkshopServiceRole&account=23>

Permissions Trust relationships Tags Access Advisor Revoke sessions

You can view the trusted entities that can assume the role and the access conditions for the role. Show policy document

Edit trust relationship

Trusted entities

The following trusted entities can assume this role.

Trusted entities

The account 23 98

Conditions

The following conditions define how and when trusted entities can assume the role.

There are no conditions associated with this role.

Replace the policy document with the following:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Principal": {
        "Service": [
          "fis.amazonaws.com"
        ]
      },
      "Action": "sts:AssumeRole",
      "Condition": {}
  }
```

]

}



EXPERIMENT (CONSOLE)

In this section, we will learn how to create an AWS FIS experiment template using the AWS Console.

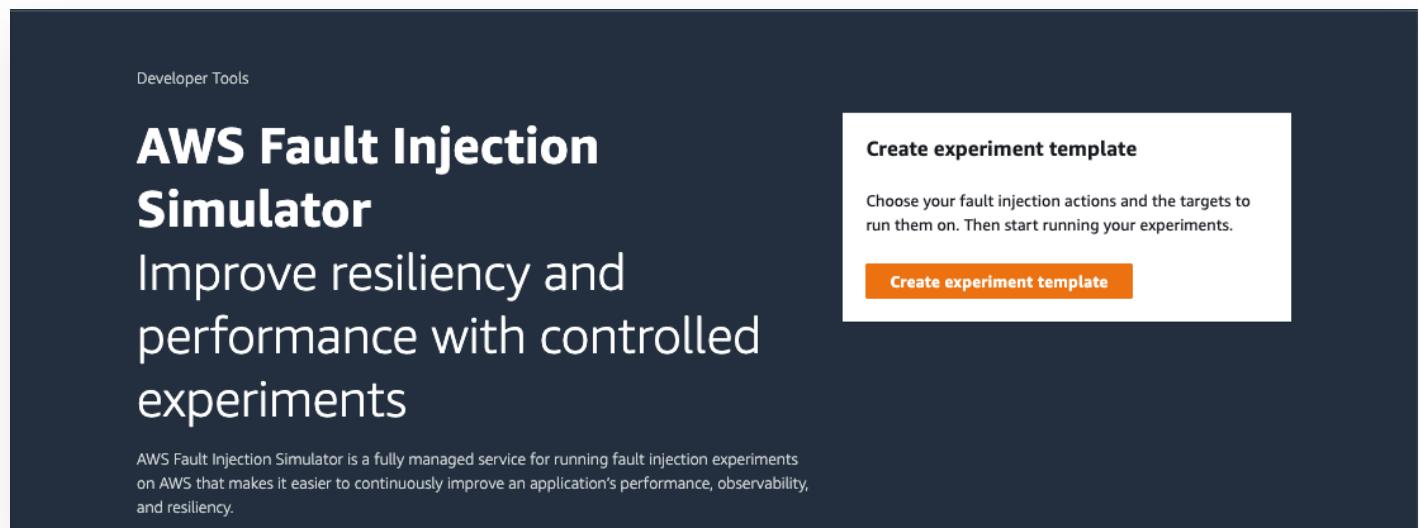
Experiment setup

To create a fault injection experiment, we first need to create an AWS FIS template defining:

- Name (optional)
- Description (optional)
- Template permissions
- **Actions**
- **Targets**
- **Stop Conditions** (optional but strongly recommended)
- Tags

Create an AWS FIS experiment template

Navigate to the [FIS console](#) and select “Create experiment template”.



Note: if you've used AWS FIS before you may not see the splash screen. In that case select "Experiment templates" in the burger menu on the left and access "**Create experiment template**" from there.

Template description and permissions

Let's write a description for our experiment template and select an IAM role to use when performing the experiment. Go to the "Description and permission" section. For "Description" enter **Terminate half of the instances in the auto scaling group** and for "IAM Role" select the **FisWorkshopServiceRole** role you created previously.

Description and permission

Description
Add a description for your experiment.

Terminate half of the instances in the autoscaling group

Enter a description of up to 512 characters.

IAM role
Select an IAM role to grant it permission to run the experiment. [Learn more](#)

FisWorkshopServiceRole
▼

Action definition

Here we select the type of fault we wish to inject, the action to take. To test the hypothesis that we can safely impact half the instances in our Auto Scaling group, we will terminate those instances. Go to the "Actions" section and select "**Add Action**".

Actions (0)

Specify one or more actions to run on your target resources. Decide how long to run each action (in minutes), and when to start the action during the experiment. [Learn more](#)

Add action

For "Name" enter **FisworkshopAsg-TerminateInstances** and add a "Description" like **Terminate instances**. For "Action type" select **aws:ec2:terminate-instances**.

We will leave the "Start after" section blank since we are only taking a single action in this experiment template.

Leave the default "Target" **Instances-Target-1** and select "**Save**".

Actions (1)

Specify one or more actions to run on your target resources. Decide how long to run each action (in minutes), and when to start the action during the experiment. [Learn more](#)

▼ New action		Save	Remove
Name	<input type="text" value="FisWorkshopAsg-TerminateInstances"/>	Description - optional	<input type="text" value="Terminate instances"/>
Action type	Select the action type to run on your targets. Learn more	Start after - optional	Select actions to run before this action. Otherwise, this action runs as soon as the experiment begins.
	<input type="text" value="aws:ec2:terminate-instances"/>		<input type="text" value="Select an action"/>
Target	A target will be automatically created for this action if one does not already exist. Additional targets can be created below.	<input type="text" value="Instances-Target-1"/>	
Add action			

Note

Instances-Target-1 was auto-generated for us because no appropriate target type existed in the experiment template. If one or more targets already exist, e.g. because we added actions before, then we will be presented with a drop down selector for existing targets instead.

Target definition

For our action we are choosing to terminate EC2 instances. In the target section we define which instances to terminate. As a reminder, for this first experiment we want to prove the hypothesis that we can safely impact half the instances in our Auto Scaling group.

Go to the "Targets" section, select the **Instances-Target-1** section, and select "**Edit**".

Targets (1)

Specify the target resources on which to run your selected actions.

► Instances-Target-1 (aws:ec2:instance)

Actions: FisWorkshopAsg-TerminateInstances

[Edit](#)

[Remove](#)

[Add target](#)

You may leave the default name `Instances-Target-1` but for maintainability we recommend using descriptive target names. Change the name to `FisWorkshopAsg-50Percent` (this will automatically update the name in the action as well) and make sure “Resource type” is set to `aws:ec2:instances`. For “Target method” we will dynamically select resources based on an associated tag. Select the `Resource tags and filters` checkbox. Pick `Percent` from “Selection mode” and enter `50`. Under “Resource tags” enter `Name` in the “Key” field and `FisStackAsg/ASG` for “Value” to select only from instances associated with the desired Auto Scaling group. Under filters enter `State.Name` in the “Attribute path” field and `running` under “Values” to ensure we do not consider instances that are starting or stopping due to unrelated events. For more information on filters see the [documentation](#). Select “**Save**”.

Add target

X

Specify the target resources on which to run your selected actions. [Learn more](#)

Name

FisWorkshopAsg-50Percent

Resource type

aws:ec2:instance

▼

Target method

- Resource IDs
- Resource tags and filters

Selection mode

Percentage (%)

Percent

50

Resource tags

Key

Name

Value - optional

FisStackAsg/ASG

Remove

Add new tag

Resource filters - optional

Filter resources by the attributes you specify. [Learn more](#)

Attribute path

State.Name

Values

running

Remove

Separate multiple values with commas.

Add new filter

Cancel

Save

Stop conditions

AWS FIS provides stop conditions tied to **Amazon CloudWatch alarms** as a safeguard to minimize the impact of experiments that do not perform as expected. In this experiment we are performing a single action that cannot be reverted so we will leave this empty.

Stop conditions

Select stop conditions - optional

If these alarms are triggered, the experiment is stopped. [Learn more](#)

Select a CloudWatch alarm

Template name and tags

Finally we can attach tags to our template. Tags can be used in IAM policy **condition keys** to control access to the experiment template. AWS FIS also recognizes the special tag `Name` which is displayed in the "Name" field of the experiment template list view.

For this experiment we will give our template a short name to be used on the list page. To do this scroll to the "Tags" section at the bottom, select "**Add new tag**", then enter `Name` in the "Key" field and `FisworkshopExp1` for "Value"

Tags

Key	Value - optional	
<input type="text" value="Name"/> X	<input type="text" value="FisWorkshopExp1"/> X	Remove

Add new tag

You can add 49 more tags.

Creating template without stop conditions

Scroll to the bottom of the template definition page and select "*Create experiment template*".

Since we didn't specify a stop condition we receive a warning. This is ok, for this experiment we won't use a stop condition. Type `create` in the text box as indicated and select "**Create experiment template**".

Create experiment template

X

⚠ You have not specified a stop condition for your experiment template. A stop condition can help to prevent your experiment from going out of bounds by stopping it automatically. [Learn more ↗](#)

To confirm that you want to create an experiment template without a stop condition, enter *create* in the field:

create

Cancel

Create experiment template

Validation procedure

We will be using the AWS CloudWatch dashboard from the previous sections for validation, no additional setup required.

Run FIS experiment

As [previously discussed](#), we should collect both customer and ops metrics. For larger experiments we would add the load generator into our experiment.

However, for this experiment we will manually trigger load generation on the system before starting the experiment, similar to what we did in the previous section. Here we have increased the run time to 5 minutes by setting `ExperimentDurationSeconds` to `300`:

```
# Please ensure that LAMBDA_ARN and URL_HOME are still set from previous section
aws lambda invoke \
--function-name ${LAMBDA_ARN} \
--payload "{
  \"ConnectionTargetUrl\": \"${URL_HOME}\",
  \"ExperimentDurationSeconds\": 300,
  \"ConnectionsPerSecond\": 1000,
  \"ReportingMilliseconds\": 1000,
  \"ConnectionTimeoutMilliseconds\": 2000,
  \"TlsTimeoutMilliseconds\": 2000,
```

```
  \"TotalTimeoutMilliseconds\": 2000  
}" \  
--invocation-type Event \  
invoke.txt
```

⚠ Warning

If you are running AWS CLI v2, you need to pass the parameter

`--cli-binary-format raw-in-base64-out` or you'll get the error "Invalid base64" when sending the payload.

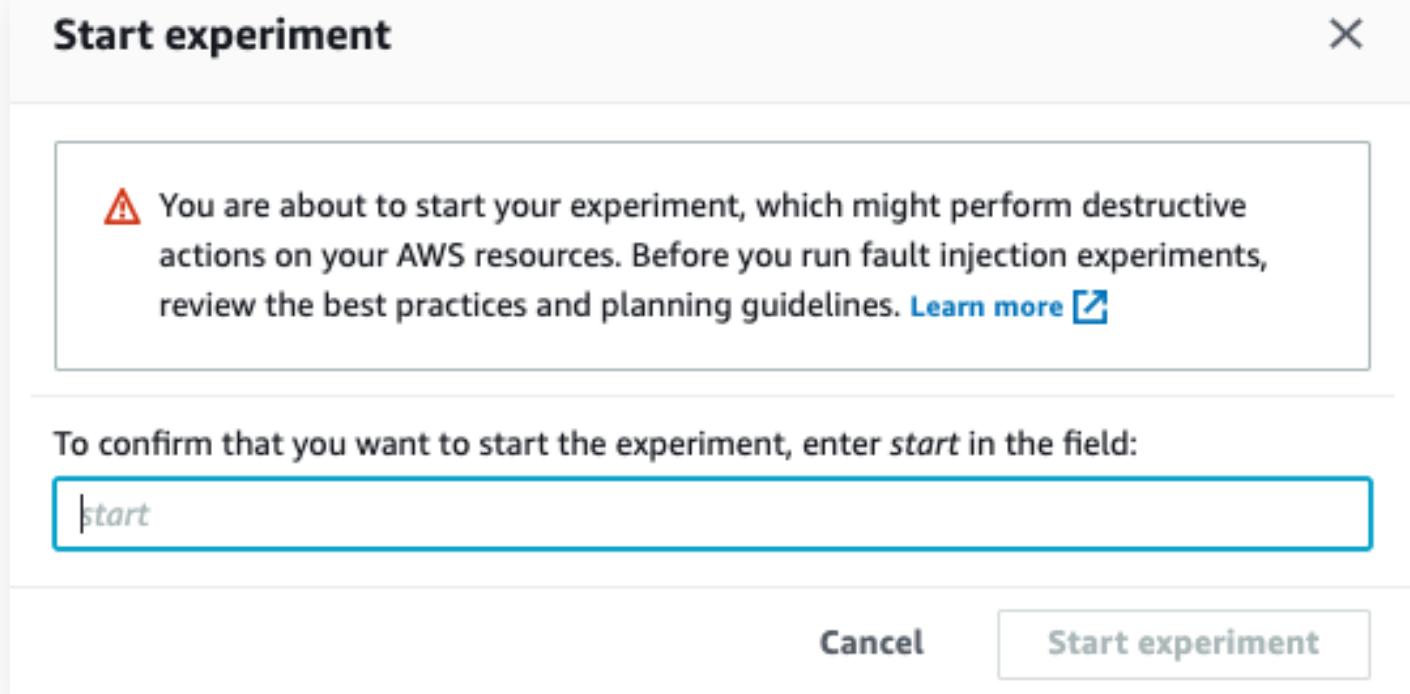
To start the experiment navigate to the **FIS console**, select the `FisWorkshopExp1` template we just created. Under **"Actions"** select **"Start experiment"**.

The screenshot shows the AWS FIS Experiment templates page. At the top, there is a breadcrumb navigation: AWS FIS > Experiment templates. Below the breadcrumb, a table lists one experiment template: FisWorkshopExp1. The table has columns for Name and Experiment template ID. On the right side of the table, there is a 'Actions' dropdown menu with several options: View details, Update experiment template, Start experiment, Manage tags, and Delete experiment template. The 'Start experiment' option is highlighted with a red box. The 'Actions' button itself is also highlighted with a red box.

Let's give the experiment run a friendly name. It will make it easier to find it from the list page. Under "Experiment tags" enter `Name` for "Key and `FisWorkshopExp1Run1` then select **"Start experiment"**.

The screenshot shows the AWS FIS Experiment tags page. At the top, there is a section titled "Experiment tags" with the sub-instruction "Associate tags with this experiment". Below this, there is a table for managing tags. The table has two columns: "Key" and "Value - optional". A single tag is listed: Key is "Name" and Value is "FisWorkshopExp1Run1". There are "Remove" and "Add new tag" buttons next to the tag entry. A note at the bottom says "You can add 49 more tags."

Because you are about to start a potentially destructive process, you will be asked to confirm that you really want to do this. Type **start** as directed and select “**Start experiment**”.



Review results

Navigate to the **FIS console**, select “Experiments”, and click the experiment ID for the experiment you just started.

Look at the “State” entry. If this still shows pending, feel free to select the “**Refresh**” button a few times until you see a result. If you followed the above steps carefully there is a good chance that your experiment state will be **Failed**.

The screenshot shows the AWS FIS Experiments page. In the top navigation bar, it says "AWS FIS > Experiments > EXPg2fM6y1MfemRxmo". Below the navigation, the experiment ID "EXPg2fM6y1MfemRxmo" is displayed with an "Info" link. To the right of the experiment ID are two buttons: "Refresh" (which is highlighted with a red box) and "Actions ▾". The main area shows a table with experiment details:

Details		State	Experiment template ID
Experiment ID EXPg2fM6y1MfemRxmo	Start time June 04, 2021, 15:11:00 (UTC-06:00)	Failed	EXT3HAoiwxmDmWqn
Creation time June 04, 2021, 15:10:59 (UTC-06:00)	End time June 04, 2021, 15:11:02 (UTC-06:00)	IAM role FisWorkshopServiceRole	Stop conditions -

Click on the failed result to get more information about why it failed. The message should say **Target resolution returned empty set**. To see why this would happen, have a look at the auto scaling group from which we tried to select instances. Navigate to the **EC2 console**, select **“Auto Scaling Groups”** on the bottom of the burger menu, and search for **FisStackAsg -**:

Auto Scaling groups (3)			
<input type="checkbox"/>	Name	Launch template/configuration	Instances
<input type="checkbox"/>	FisStackAsg-ASG46ED3070-OL3BET7A77OV	FisStackAsg-ASGLaunchConfigC00A...	1

Learning and improving

It looks like our ASG was configured to scale down to just one instance while idle. Since we can't shut down half of one instance our 50% selector came up empty and the experiment failed.

Great! While this wasn't really what we expected, we just found a flaw in our configuration that would severely affect our system's resilience! Let's fix it and try again!

Click on the auto scaling group name and **“Edit”** the “Group Details” to raise both the “Desired capacity” and “Minimum capacity” to **2**.

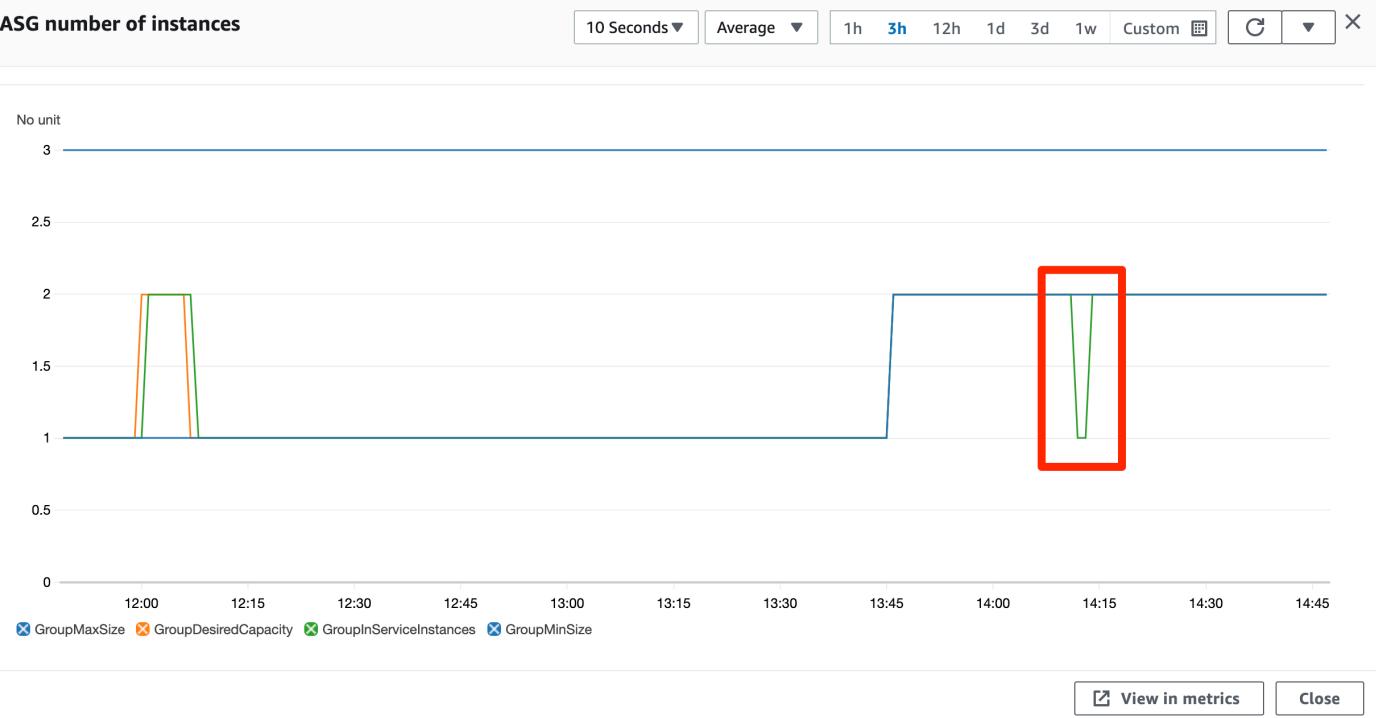
Group details	
Desired capacity	<input type="text" value="1"/> Edit
Minimum capacity	<input type="text" value="1"/> Date created Fri May 21 2021 17:08:24 GMT-0600 (Mountain Daylight Time)
Maximum capacity	<input type="text" value="3"/> Amazon Resource Name (ARN) arn:aws:autoscaling:us-east-2:238810465798:autoScalingGroup:7aa2afdc-f249-4494-9404-a1c24ec5bc65:autoScalingGroupName/FisStackAsg-WebServerGroup-1LXEZRDXBRVJ5

Check the ASG details or the CloudWatch Dashboard we explored in the previous section to make sure the active instances count has come up to 2.

To repeat the experiment, repeat the steps above:

- restart the load
- navigate back to the **FIS Experiment Templates Console**, start the experiment adding a **Name** tag of **FisWorkshopExp1Run2**
- check to make sure the experiment succeeded

Finally navigate to the **CloudWatch Dashboard** from the previous section. Review the number of instances in the ASG going down and then up again and review the error responses reported by the load test.



Findings and next steps

From this experiment we learned:

- Carefully choose the resource to affect and how to select them. If we had originally chosen to terminate a single instance (**COUNT**) rather than a fraction (**PERCENT**), we would have severely affected our service.
- Spinning up instances takes time. To achieve resilience, Auto Scaling groups should be set to have at least two instances running at all times

From here you can explore how to set up experiments programatically using the [AWS CLI](#) or [AWS CloudFormation](#), or move on exploring [more fault types](#) to inject.



EXPERIMENT (CLI)

In this section we will show you how to create an experiment using AWS FIS templates. For clarity, we will replicate the same experiment as we previously did via the AWS console.

Template overview

Experiment templates are JSON files containing Actions, Targets, an IAM role, and optional Stop Conditions, and Tags:

```
{  
  "experimentTemplate": {  
    "description": "...",  
    "actions": {},  
    "targets": {},  
    "roleArn": "arn:aws:iam:...",  
    "stopConditions": [],  
    "tags": {}  
  }  
}
```

Actions

Actions specify an action name and `description`, an `actionId` and matching `parameters` picked from the [AWS FIS Action reference](#), and a list of `targets` which references the target section in the same template:

```
"ActionName": {  
  "description": "ActionDescription",  
  "actionId": "aws:ec2:terminate-instances",  
  "parameters": {},  
  "targets": {}  
}
```

Targets

Targets specify a name, a `resourceType` from which to select by `resourceArn`, `resourceTags` or `filters`, and `selectionMode` for sampling from the eligible resources by `COUNT()` or `PERCENT()`.

```
"TargetGroupName": {
    "resourceType": "aws:ec2:instance",
    "resourceArns": [],
    "resourceTags": {
        "TagName1": "TagValue1",
        "TagName2": "TagValue2",
        ...
    },
    "filters": [
        {
            "path": "State.Name",
            "values": [
                "running"
            ]
        }
    ],
    "selectionMode": "COUNT(1)"
}
```

A note on finding the `path` and `values` for `filters`: as described under "**Resource filters**" in the AWS documentation, filter paths are based on API output. E.g.: if we want to only target running EC2 instances we could use the AWS CLI to list instances:

```
aws ec2 describe-instances
```

To find the relevant `path` and `values` start in the `Instances` block of the API output and identify entries you would like to filter on:

```
{
    "Reservations": [
        {
            "Groups": [],
            "Instances": [
                {
                    "ImageId": "ami-00c36fdebc0d948bd",
                    "InstanceType": "t2.micro",
                    "Placement": {
                        "AvailabilityZone": "us-east-2a",
                        "GroupName": "",
                        "Tenancy": "default"
                    }
                }
            ]
        }
    ]
}
```

```

},
  "State": {
    "Code": 16,
    "Name": "running"
  },
  "SubnetId": "subnet-0e912694b51e205d6",
  "VpcId": "vpc-0d4c31ce84606e7eb",
  "Tags": [
    {
      "Key": "Name",
      "Value": "FisStackAsg/ASG"
    },
    ...
  ],
  ...
},
  ...
]
}

```

E.g.: to select an instance that is `running` in `us-east-2a` we would add the following filters:

```

"filters": [
  {
    "path": "State.Name",
    "values": [
      "running"
    ]
  },
  {
    "path": "Placement.AvailabilityZone",
    "values": [
      "us-east-2a"
    ]
  }
],

```

Stop conditions

Stop conditions use a list of Amazon CloudWatch alarms to prematurely stop the experiment if it does not proceed along expected lines:

```

"stopConditions": [
  {

```

```
        "source": "aws:cloudwatch:alarm",
        "value": "arn:aws:cloudwatch:..."
    }
]
```

Finished template

Using the above, this would be the finished template.

Note

Before using this template, please ensure that you replace the ARN for the FIS execution role on the last line with the ARN of the role you created in the [Configuring permissions](#) section.

```
{
  "description": "Terminate 50% of instances based on Name Tag",
  "tags": {
    "Name": "FisWorkshop-Exp1-CLI"
  },
  "actions": {
    "FisWorkshopTerminateAsg-1-CLI": {
      "actionId": "aws:ec2:terminate-instances",
      "description": "Terminate 50% of instances based on Name Tag",
      "parameters": {},
      "targets": {
        "Instances": "FisWorkshopAsg-50Percent"
      }
    },
    "Wait": {
      "actionId": "aws:fis:wait",
      "parameters": {
        "duration": "PT3M"
      }
    }
  },
  "targets": {
    "FisWorkshopAsg-50Percent": {
      "resourceType": "aws:ec2:instance",
      "resourceTags": {
        "Name": "FisStackAsg/ASG"
      },
      "selectionMode": "PERCENT(50)"
    }
  },
  "stopConditions": [
    {
      "source": "none"
    }
  ]
}
```

```
],
  "roleArn":  
"arn:aws:iam::YOUR_ACCOUNT_NUMBER_HERE:role/FisWorkshopServiceRole"  
}
```

Working with templates

The rest of this section uses the [AWS CLI](#). If you are using [AWS Cloud9](#) this should work out of the box. Otherwise please ensure you have configured AWS credentials for the CLI.

Creating templates

To create an experiment template, copy the above JSON into a file named `fis.json` and ensure you have changed the `roleArn` entry to be the ARN of the role you created earlier. Then use the CLI to create the template in AWS:

```
aws fis create-experiment-template --cli-input-json file://fis.json
```

You should now be able to see the newly created experiment template in the [AWS Console](#).

Listing templates

This command

```
aws fis list-experiment-templates
```

will list all the templates. If you happened to run the `create-experiment-template` command above multiple times you might notice that it is possible to have multiple copies of a template only differentiated by the `id` field.

While it is possible to update an existing experiment template via the `update-experiment-template` command, and while the content of the template at execution time is saved with the experiment data, this may make it harder to establish what happened during an experiment.

Exporting / saving templates

Once you have established the `id` of an experiment template you can dump the template. This can be a good way of learning how to write templates as well:

```
export EXPERIMENT_TEMPLATE_ID=<YOUR_EXPERIMENT_TEMPLATE_ID_HERE>
aws fis get-experiment-template --id $EXPERIMENT_TEMPLATE_ID
```

You will note that the result is wrapped into an `experimentTemplate: {}` block. You may also notice that there are some additional fields that are not used during experiment template creation. You can generate reusable JSON like so:

```
aws fis get-experiment-template --id $EXPERIMENT_TEMPLATE_ID | jq
'.experimentTemplate | del(.id) | del(.creationTime) | del(.lastUpdateTime)'
```

Running the experiment

Finally we want to run the experiment:

```
aws fis start-experiment --experiment-template-id $EXPERIMENT_TEMPLATE_ID --tags
Name=FisWorkshopTerminateAsg-1-CLI | jq '.experiment.id'
```

Using the returned `id` field you can check on the outcome of the experiment:

```
aws fis get-experiment --id YOUR_EXPERIMENT_ID_HERE
```

Findings and next steps

The learnings here should be the same as for the console section:

- Carefully choose the resource to affect and how to select them. If we had originally chosen to terminate a single instance (`COUNT`) rather than a fraction (`PERCENT`), we would have severely affected our service.
- Spinning up instances takes time. To achieve resilience, ASGs should be set to have at least two instances running at all times

From here you can explore how to set up experiments using [AWS CloudFormation](#) or move on exploring [more fault types](#) to inject.



EXPERIMENT

(CLOUDFORMATION)

In this section we will cover how to define and update experiment templates using [CloudFormation](#).

CFN template format

The AWS CloudFormation template uses the same format as the API but capitalizes the first letter of section names. As such the FIS experiment template from the previous section would become:

```
{  
  "Type" : "AWS::FIS::ExperimentTemplate",  
  "Properties" : {  
    "Description": "Terminate 50% of instances based on Name Tag",  
    "Tags": {  
      "Name": "FisWorkshop-Exp1-CFN-v1.0.0"  
    },  
    "Actions": {  
      "FisWorkshopTerminateAsg-1-CFN": {  
        "ActionId": "aws:ec2:terminate-instances",  
        "Description": "Terminate 50% of instances based on Name Tag",  
        "Parameters": {},  
        "Targets": {  
          "Instances": "FisWorkshopAsg-50Percent"  
        }  
      },  
      "Wait": {  
        "ActionId": "aws:fis:wait",  
        "Parameters": {  
          "duration": "PT3M"  
        }  
      }  
    },  
    "Targets": {  
      "FisWorkshopAsg-50Percent": {  
        "ResourceType": "aws:ec2:instance",  
        "ResourceTags": {  
          "Name": "FisStackAsg/ASG"  
        }  
      },  
    }  
  }  
}
```

```

        "SelectionMode": "PERCENT(50)"
    }
},
"StopConditions": [
    {
        "Source": "none"
    }
],
"RoleArn": {
    "Fn::Sub": "arn:aws:iam::YOUR_ACCOUNT_ID:role/FisWorkshopServiceRole"
}
}
}

```

We can wrap this into the `Resources` section of a **CloudFormation template**. Additionally CloudFormation allows us to use **pseudo parameters** which we can use to automatically insert the account number into the role definition using the `AWS::AccountId` parameter in conjunction with the `Fn::Sub` function. Thus, a simple CFN template would become:

```

{
"Resources" : {
    "FisExperimentDemo" : {
        "Type" : "AWS::FIS::ExperimentTemplate",
        "Properties" : {
            "Description": "Terminate 50% of instances based on Name Tag",
            "Tags": {
                "Name": "FisWorkshop-Exp1-CFN-v1.0.0"
            },
            "Actions": {
                "FisWorkshopTerminateAsg-1-CFN": {
                    "ActionId": "aws:ec2:terminate-instances",
                    "Description": "Terminate 50% of instances based on Name
Tag",
                    "Parameters": {},
                    "Targets": {
                        "Instances": "FisWorkshopAsg-50Percent"
                    }
                },
                "Wait": {
                    "ActionId": "aws:fis:wait",
                    "Parameters": {
                        "duration": "PT3M"
                    }
                }
            },
            "Targets": {
                "FisWorkshopAsg-50Percent": {
                    "ResourceType": "aws:ec2:instance",
                    "ResourceTags": {
                        "Name": "FisStackAsg/ASG"
                    },

```

```
        "SelectionMode": "PERCENT(50)"
    }
},
"StopConditions": [
{
    "Source": "none"
}
],
"RoleArn": {
    "Fn::Sub": "$arn:aws:iam::${AWS::AccountId}:role/FisworkshopServiceRole"
}
}
}
}
```

Using the CFN template

A deep dive into **AWS CloudFormation** is beyond the scope of this workshop, so we will only cover how to create and update stacks via the CLI.

Create a new template / experiment

To create a stack, and thus the contained FIS experiment template, copy the above JSON into a file named `cfn-fis-experiment.json` then run this AWS CLI command:

```
aws cloudformation create-stack --stack-name FisWorkshopExperimentTemplate --  
template-body file://cfn-fis-experiment.json
```

If you navigate to the **CloudFormation console** you should now see a new stack named `FisWorkshopExperimentTemplate` and navigating to the **FIS console** should show an experiment named `FisWorkshop-Exp1-CFN-v1.0.0`

Update template / experiment

To update the experiment template you will need to update the CFN template. Let's change the `Name` tag from `FisWorkshop-Exp1-CFN-v1.0.0` to `FisWorkshop-Exp1-CFN-v2.0.0` and save the file.

Then run the AWS CLI command:

```
aws cloudformation update-stack --stack-name FisWorkshopExperimentTemplate --template-body file://cfn-fis-experiment.json
```

This should update the name of your experiment template in the FIS console. Obviously this is most useful if you make actual changes to the template itself too.

Findings and next steps

The learnings here should be the same as for the console section:

- Carefully choose the resource to affect and how to select them. If we had originally chosen to terminate a single instance (**COUNT**) rather than a fraction (**PERCENT**), we would have severely affected our service.
- Spinning up instances takes time. To achieve resilience, ASGs should be set to have at least two instances running at all times

In the next section we will explore more fault injection options.



AWS SYSTEMS MANAGER INTEGRATION

In this section, we will demonstrate how you can use AWS Systems Manager (SSM) along with AWS Fault Injection Simulator (FIS) to emulate faults within an EC2 Instance.

AWS FIS does not need an agent for actions affecting the AWS control plane like the ones we have worked with thus far, such as stopping instances or failing over RDS Databases. However, there are actions that require us to initiate commands within the operating system of the EC2 Instance, such as affecting CPU or Memory consumption, or terminating processes. For these types of actions AWS FIS can use

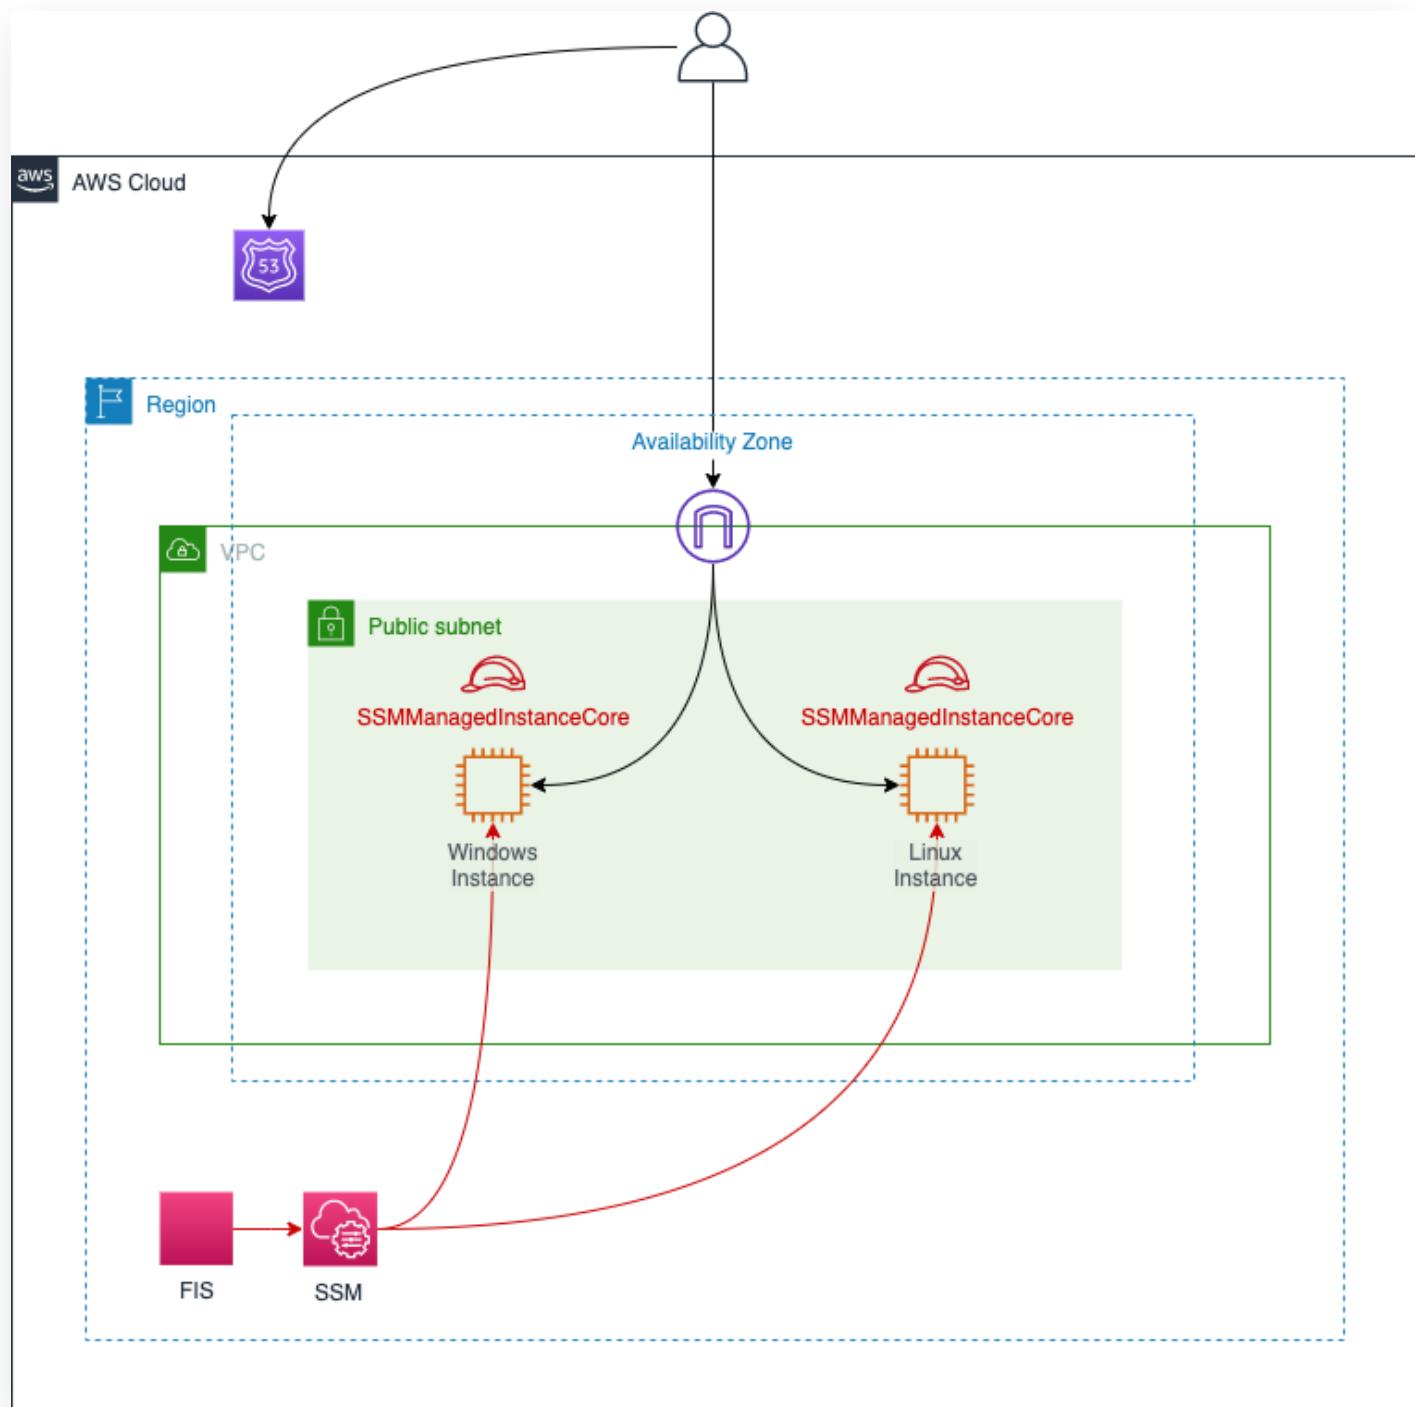
AWS Systems Manager (SSM) and the **SSM Agent**. This approach provides you with the access controls to grant FIS limited access to your instances under the **shared responsibility model**.

In the following sections we will show you how to use the built-in SSM actions and how to build your own SSM documents to create custom actions.



FIS SSM SEND COMMAND SETUP

For this section we will use Linux and Windows instances created specifically for the purpose of enabling FIS SSM commands. As shown in the diagram below, SSM access to the instance **requires an instance role** with the `AmazonSSMManagedInstanceCore` policy attached. Additionally FIS access to SSM is controlled via the execution policy as shown in the **First Experiment** section.



The resources above have been created as part of the [Start the workshop](#) section. If you would like to examine how these resources were defined you can examine the [AWS Cloud Formation template](#).



LINUX CPU STRESS EXPERIMENT

Experiment idea

In this section we are exploring tooling so we will start without a hypothesis. However, we will provide some learnings and next steps at the end.

Specifically, in this section we will run a CPU Stress test using AWS Fault Injection Simulator against an Amazon Linux EC2 Instance. The Linux **CPU stress** test is an out of the box FIS action. We will do the following:

1. Create experiment template to stress CPU.
2. Connect to a Linux EC2 Instance and run the `top` command.
3. Start experiment and observe results.

Experiment setup

Note

We are assuming that you know how to set up a basic FIS experiment and will focus on things specific to this experiment. If you need a refresher see the previous **First Experiment** section.

General template setup

- Create a new experiment template
 - Add `Name` tag of `LinuxBurnCPUviaSSM`
 - Add `Description` of `Inject CPU stress on Linux`
 - Select `FisCpuStress-FISRole` as execution role

Action definition

In the “Actions” section select the **“Add Action”** button.

Name the Action as `StressCPUViaSSM`, and under "Action Type" select `aws:ssm:send-command/AWSFIS-Run-Cpu-Stress`. This is an out of the box action to run stress test on Linux Instances using the **stress-ng** tool. Set the "documentParameters" field to `{"DurationSeconds":120}` which is passed to the script and the "duration" field to 2 minutes which tells FIS how long to wait for a result. Leave the default "Target" `Instances-Target-1` and select "**Save**".

New action

Name `LinuxBurnCPUviaSSM`

Description - optional Inject CPU stress on Linux

Action type Select the action type to run on your targets. [Learn more](#)

`aws:ssm:send-command/AWSFIS-Run-CPU-St...`

Start after - optional Select actions to run before this action. Otherwise, this action runs as soon as the experiment begins.

`Select an action`

Target A target will be automatically created for this action if one does not already exist. Additional targets can be created below.

`Instances-Target-1`

Action parameters Specify the parameter values for this action. [Learn more](#)

documentArn The ARN of the SSM document to run.

`arn:aws:ssm:us-west-2::document/AWSFIS-Run-CPU-`

documentParameters The JSON string of the parameters to pass to the document that is run.

`{"DurationSeconds":120}`

documentVersion - optional The version of the document to run. If not specified, the document's default version will be used.

duration The length of time to monitor the SSM command (ISO 8601 duration).

`Minutes` `2` `PT2M`

This action will use **AWS Systems Manager Run Command** to run the `AWSFIS-Run-Cpu-Stress` command document against our targets for two minutes.

Target selection

For this action we need to designate EC2 instance targets on which to run the commands. Go to the "Targets" section, select the `Instances-Target-1` section, and select "**Edit**".

You may leave the default name `Instances-Target-1` but for maintainability we recommend using descriptive target names. Change the name to `FisWorkshop-StressLinux` (this will automatically update the name in the action as well) and make sure “Resource type” is set to `aws:ec2:instances`. To select our target instances by tag select “Resource tags and filters” and keep selection mode `ALL`. Select **“Add new tag”** and enter a “Key” of `Name` and a “Value” of `FisLinuxCPUSTress`. Finally select **“Save”**.

Edit target

Specify the target resources on which to run your selected actions. [Learn more](#)

Name	Resource type
<code>FisWorkshop-StressLinux</code>	<code>aws:ec2:instance</code>

Actions
`LinuxBurnCPUviaSSM`

Target method
 Resource IDs
 Resource tags and filters

Selection mode
All

Resource tags

Key	Value - optional
Name	<code>FisLinuxCPUSTress</code>

[Remove](#)

[Add new tag](#)

Resource filters - optional
Filter resources by the attributes you specify. [Learn more](#)

No resource filters are associated with the target.

[Add new filter](#)

[Cancel](#) [Save](#)

Creating template without stop conditions

Select **“Create experiment template”** and confirm that you wish to create a template without stop conditions.

Validation procedure

We will use the linux `top` system command to observe the increased CPU load. To do this we now need to connect to our EC2 Instance so we can observe the CPU being stressed. Head over to the **EC2 Console**.

- Once at the EC2 Console lets select our instance named `FisLinuxCpuStress` and click on the "Connect" button.

The screenshot shows the AWS EC2 Instances page. At the top, it displays "Instances (1/9)" and a "Filter instances" search bar. Below this, there are two buttons: "Instance state: running" (with an "X" icon) and "Clear filters". The main table lists one instance: "FisLinuxCPUStress" with a checked checkbox. To the right of the instance name is a "Connect" button, which has a large red arrow pointing directly at it. The table also includes columns for "Name" and "Instance ID".

- Select "**Session Manager**" and select "**Connect**".

The screenshot shows the "Connect to instance" dialog. At the top, it says "Connect to instance" and "Info". Below that, it says "Connect to your instance i-0d1e1d3b5258b34b7 (CPUSTressTest) using any of these options". There are four tabs at the bottom: "EC2 Instance Connect", "Session Manager" (which is highlighted with a yellow arrow), "SSH client", and "EC2 Serial Console". Under the "Session Manager" tab, there is a section titled "Session Manager usage:" with the following bullet points:

- Connect to your instance without SSH keys or a bastion host.
- Sessions are secured using an AWS Key Management Service key.
- You can log session commands and details in an Amazon S3 bucket or CloudWatch Logs log group.
- Configure sessions on the Session Manager [Preferences](#) page.

At the bottom right of the dialog, there are "Cancel" and "Connect" buttons, with a yellow arrow pointing to the "Connect" button.

This will open a session to the EC2 instance in another tab. In the new tab enter:

top

You should now see a continuously updating display similar to the next screenshot. Initially the CPU percentage should be at or close to zero as this instance is not doing anything. Keep this tab open, we will come back once we have started our experiment.

The screenshot shows the output of the 'top' command. At the top, it displays system load average (0.00, 0.00, 0.00), tasks (86 total, 1 running, 47 sleeping, 0 stopped, 0 zombie), CPU usage (%Cpu(s): 0.0 us, 0.2 sy, 0.0 ni, 99.8 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st), memory usage (KiB Mem: 3977792 total, 3574032 free, 122444 used, 281316 buff/cache), and swap usage (KiB Swap: 0 total, 0 free, 0 used). Below this is a detailed process list with columns: PID, USER, PR, NI, VIRT, RES, SHR, S, %CPU, %MEM, TIME+, and COMMAND. The processes listed are: 1 root (systemd), 2 root (kthreadd), 4 root (kworker/0:0H), 6 root (mm_percpu_wq), 7 root (ksoftirqd/0), 8 root (rcu_sched), 9 root (rcu_bh), and several other low-priority background tasks.

Run CPU Stress Experiment

Note

We are assuming that you know how to set up a basic FIS experiment and will focus on things specific to this experiment. If you need a refresher see the previous [First Experiment](#) section.

Keep the EC2 instance session with `top` running. In a new browser window navigate to the [AWS Fault Injection Simulator Console](#) and start the experiment:

- use the `LinuxBurnCPUviaSSM`
- add a `Name` tag of `FisWorkshopLinuxStress1`
- confirm that you want to start the experiment
- ensure that the "State" is `Running`

Details

Experiment ID	Start time	State
[REDACTED]	July 08, 2021, 22:35:20 (UTC-04:00)	Running

In the EC2 terminal window watch the CPU percentage displayed by `top`: it should hit 100% for a few minutes and then return back to 0%. Once we have observed the action we can click the `Terminate` button to terminate our Session Manager session.

```
top - 01:50:32 up 1:48, 0 users,  Load average: 1.07, 0.28, 0.10
Tasks: 94 total, 3 running, 50 sleeping, 0 stopped, 0 zombie
%Cpu(s): 100.0 us, 0.0 sy, 0.0 ni, 0.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem : 3977792 total, 3536452 free, 145152 used, 296188 buff/cache
KiB Swap: 0 total, 0 free, 0 used. 3625256 avail Mem

 PID USER      PR  NI    VIRT    RES    SHR S %CPU %MEM     TIME+ COMMAND
 2909 root      20   0  69196  2672  1820 R 100.0  0.1  0:32.83 stress-ng-cpu
 2910 root      20   0  69196  2672  1820 R 100.0  0.1  0:33.47 stress-ng-cpu
  1 root      20   0  43648  5388  3900 S  0.0  0.1  0:01.26 systemd
  2 root      20   0      0      0      0 S  0.0  0.0  0:00.00 kthreadd
  4 root      0 -20      0      0      0 I  0.0  0.0  0:00.00 kworker/0:0H
  6 root      0 -20      0      0      0 I  0.0  0.0  0:00.00 mm_percpu_wq
  7 root      20   0      0      0      0 S  0.0  0.0  0:00.04 ksoftirqd/0
```

Congratulations for completing this lab! In this lab you walked through running an experiment that took action within a Linux EC2 Instance using AWS Systems Manager. Using the integration between Fault Injection Simulator and AWS Systems Manager you can run scripted actions within an EC2 Instance. Through this integration you can script events against your applications or run other chaos engineering tools and frameworks.

Learning and improving

Since this instance wasn't doing anything, there aren't any actions. To think about how to use this to test a hypothesis and make an improvement, consider running the same experiment against the ASG instances from the **First Experiment** section. Maybe you could use this to tune the optimal CPU levels for scaling up or down?

<

>



WORKING WITH SSM DOCUMENTS

Pre-configured SSM documents

The linux CPU stress experiment we saw in the previous section used one of the **pre-configured SSM documents** to run a script on our Linux instance.

To find the script, navigate to the **AWS Systems Manager console**, scroll down in the left-hand menu all the way to the bottom to “**Documents**”, select “**Owned by Amazon**”, and search for **AWSFIS**. Note that this search may take a few seconds to display results.

The screenshot shows the AWS Systems Manager console with the left navigation bar expanded. The 'Shared Resources' section is selected, and the 'Documents' tab is highlighted with a red box. The main content area shows a list of SSM documents under the 'Owned by Amazon' tab, also highlighted with a red box. A search bar at the top of the document list has 'Search: AWSFIS' entered, with another red box around it. The search results show three documents: 'AWSFIS-Run-CPU-Stress', 'AWSFIS-Run-Kill-Process', and 'AWSFIS-Run-Memory-Stress'. Each document card provides details like Document type (Command), Owner (Amazon), Platform types (Linux), and Default version (2, 2, and 3 respectively). A 'Create document' button is visible at the top right of the document list.

To inspect the script, click on the script name, i.e. **AWSFIS-Run-CPU-Stress**, then select the “**Content**” tab.

The content of this document is as follows:

```

1  ---
2  description: I
3  ## Document name - AWSFIS-Run-CPU-Stress
4
5  ## What does this document do?
6  It runs CPU stress on an instance via stress-ng tool.
7
8  ## Input Parameters
9  * DurationSeconds: (Required) The duration - in seconds - of the CPU stress.
10 * CPU: Specify the number of CPU stressors to use (Default 0 = all)
11 * InstallDependencies: If set to True, Systems Manager installs the required dependencies on the target instances. (Default False)
12
13 ## Output Parameters
14 None.
15
16 schemaVersion: '2.2'
17 parameters:
18   DurationSeconds:
19     type: String
20     description: "(Required) The duration - in seconds - of the CPU stress."
21     allowedPattern: "^[0-9]+$"

```

The document is a YAML file defining two `aws:runShellScript` actions: `InstallDependencies` to install the `stress-ng` package, and `ExecuteStressNg` to inject CPU stress.

Custom SSM documents

Currently AWS does not provide a CPU stress document for Windows, but we can create our own as shown in the next section. For more information on writing SSM documents please see these resources:

- [AWS Systems Manager documentation](#)
- [Writing your own SSM documents blog](#)
- [AWS SSM workshop](#)

If you want to see an example how one might inject stress, you can have a look at the `WinStressDocument` resource in the [CloudFormation template](#). Alternatively you can follow the same search procedure as for the AWS owned documents but search the “Owned by me” or “Shared by me” tabs instead of “Owned by AWS”.

For additional SSM sample documents relating to FIS see these resources

- <https://github.com/adhorn/chaos-ssm-documents>

Working with custom SSM documents in FIS

While writing custom SSM documents is outside the scope of this workshop, there are a few aspects of SSM documents you should be aware of:

- **Document ARN** - FIS requires the full SSM document ARN. The ARN can easily be constructed from the document name (and the owner account ID if the template is shared with you) using this format string:
`arn:${AWS::Partition}:ssm:${AWS::Region}:${AWS::AccountId}:document/${WinStressDocumentName}`
- **Exit status** - Shell script convention is to signal success with a return/exit value of `0` and a failure with any non-zero numeric value. If FIS detects a non-zero exit status on an SSM script it will mark the action as "Failed", terminate all running actions, cancel queued actions, invoke any outstanding roll-back actions, cancel experiment execution, and mark the overall experiment as "Failed".
- **Duration** - FIS actions have a "Duration" setting and will stop action execution if the action has not finished within this time period. For SSM actions this will "**cancel**" the command. If the command has a sequence of steps, this will result in only some of the steps being executed.
- **onCancel / onFailure** - SSM provides you with a means to ensure that automation can fail / clean-up safely by providing `onCancel` and `onFailure` properties on each step. These properties allow designating clean-up steps to perform.



WINDOWS CPU STRESS EXPERIMENT

Warning

This section requires that you have an RDP client on your local machine. This section cannot be performed from a cloud9 instance.

Experiment idea

In this section we are exploring tooling so we will start without a hypothesis. However, we will provide some learnings and next steps at the end.

Specifically, in this section we will run a CPU Stress test using AWS Fault Injection Simulator against an Amazon Windows EC2 Instance. The Windows CPU stress test will use a custom SSM command document. We will do the following:

1. Create experiment template to stress CPU.
2. Reset password on Windows Instance.
3. Connect to Windows EC2 Instance and run task manager.
4. Start experiment and observe results.

Experiment Setup

Note

We are assuming that you know how to set up a basic FIS experiment and will focus on things specific to this experiment. If you need a refresher see the previous [First Experiment](#) section.

General template setup

- Create a new experiment template
 - Add **Name** tag of `WindowsBurnCPUviaSSM`
 - Add **Description** of `Inject CPU stress on Windows`
 - Select `FisCpuStress-FISRole` as execution role

Action definition

In the "Actions" section select the "**Add Action**" button.

"Name" the action as `StressCPUViaSSM`, and under "Action Type" select the `aws:ssm:send-command` action. Currently there is no out of box Action for Windows CPU Stress Testing, so we are using the send-command action along with a command document that was deployed by our CloudFormation template. To view this document please reference the `WinStressDocument` resource in the [CloudFormation template](#).

To find the ARN of the document that was created by the template, open a new tab and browse to the [CloudFormation console](#), select "**Stacks**", select the stack named "**FisCpuStress**", then select "**Outputs**". Copy the value of the `WinStressDocumentArn` entry as you will need it in the next step.

Return to the FIS console and enter the ARN you copied into the "documentArn" field. Then set the "documentParameters" field to `{"durationSeconds":120}` which is passed to the script and the "duration" field to `2` minutes which tells FIS how long to wait for a result. Leave the default "Target" `Instances-Target-1` and select "**Save**".

▼ StressCPUViaSSM / aws:ssm:send-command (2 min)

SaveCancel

Name

Description - optional

Action type

Select the action type to run on your targets. [Learn more](#)

Start after - optional

Select actions to run before this action. Otherwise, this action runs as soon as the experiment begins.

Target

A target will be automatically created for this action if one does not already exist. Additional targets can be created below.

Action parameters

Specify the parameter values for this action. [Learn more](#)

documentArn

The ARN of the SSM document to run.

documentParameters - optional

The JSON string of the parameters to pass to the document that is run.

documentVersion - optional

The version of the document to run. If not specified, the document's default version will be used.

duration

The length of time to monitor the SSM command (ISO 8601 duration).

This action will use **AWS Systems Manager Run Command** to run the `FisCpuStress-WinStressDocument` document against our targets for two minutes.

Target selection

For this action we need to designate EC2 instance targets on which to run the commands. Go to the "Targets" section, select the `Instances-Target-1` section, and select "**Edit**".

You may leave the default name `Instances-Target-1` but for maintainability we recommend using descriptive target names. Change the name to `Fisworkshop-StressWindows` (this will automatically update the name in the action as well) and make sure "Resource type" is set to `aws:ec2:instances`. To select our target instances by tag select "Resource tags and filters" and keep selection mode `ALL`. Select "**Add new tag**" and enter a "Key" of `Name` and a "Value" of `FisWindowsCPUSTress`. Finally select "**Save**".

Edit target

X

Specify the target resources on which to run your selected actions. [Learn more](#)

Name

`FisWorkshop-StressWindows`

Resource type

aws:ec2:instance

▼

Actions

LinuxBurnCPUviaSSM

Target method

- Resource IDs
 Resource tags and filters

Selection mode

All

▼

Resource tags

Key

Name

Value - optional

FisWindows|CPUSTress

Remove

Add new tag

Resource filters - optional

Filter resources by the attributes you specify. [Learn more](#)

No resource filters are associated with the target.

Add new filter

Cancel

Save

Creating template without stop conditions

Select “**Create experiment template**” and confirm that you wish to create a template without stop conditions.

Validation procedure

We will use the Windows task manager to observe increased CPU load. To do this we now need to connect to our EC2 Instance so we can observe the CPU being stressed.

Use AWS Systems Manager Run Command to reset Password

When we deployed the instance we didn't use SSH Keys, and we don't know the password. However, with the SSM Agent along with the right IAM privileges we have a break glass scenario where we can reset the password. Please adjust the value of `TMP_PASSWORD` and use the commands below to find the InstanceId of the `FisWindowsCPUSTress` instance and help you reset the admin password to the password of choice.

```
# For readability - passing passwords this way is not secure
TMP_PASSWORD=ENTER_NEW_PASSWORD_HERE

# For readability and convenience
TMP_INSTANCE=$( aws ec2 describe-instances --filter
Name=tag:Name,Values=FisWindowsCPUSTress --query
'Reservations[*].Instances[0].InstanceId' --output text )

# Reset password on instance - this is not a secure method,
# in real life use AWS-PasswordReset document
aws ssm send-command \
--document-name "AWS-RunPowerShellScript" \
--document-version "1" \
--targets '[{"Key":"InstanceIds","Values":["'$TMP_INSTANCE'""]}]' --parameters
'{"workingDirectory":[""], "executionTimeout":["3600"], "commands":["net user
administrator '$TMP_PASSWORD'"]}' \
--timeout-seconds 600 \
--max-concurrency "50" \
--max-errors "0" \
--cloud-watch-output-config '{"CloudWatchOutputEnabled":false}'
```

Use AWS Systems Manager Session Manager to connect to Target Instance

We now need to connect to our EC2 Instance so we can observe the CPU being stressed. We are going to do this by using the port forwarding capability of AWS Systems Manager Session Manager and using RDP.

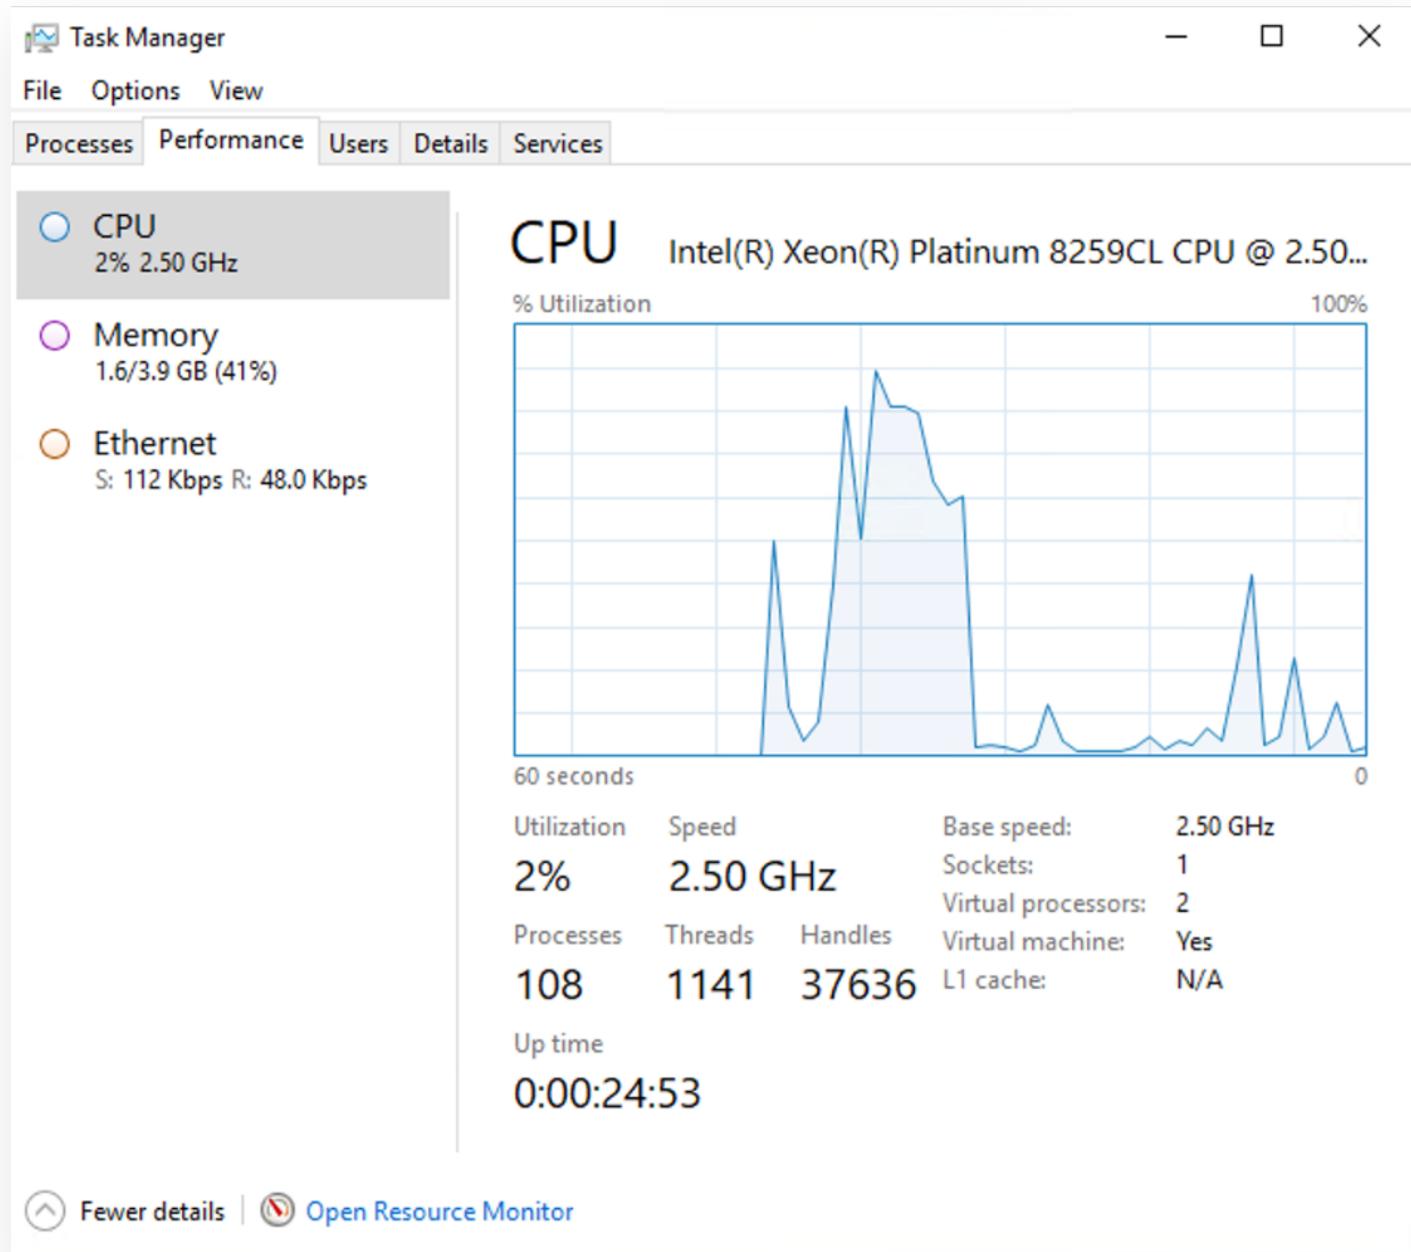
1. First make sure that the **Session Manager plugin for the AWS CLI** is installed on your local machine.
2. Run the following command first, this will securely forward local port 56788 to port 3389 on the Windows EC2 Instance. Note that we are targeting a specific instance by passing the `TMP_INSTANCE` variable from above.

```
# This presumes you set TMP_INSTANCE (see above)
aws ssm start-session --target ${TMP_INSTANCE} --document-name AWS-
StartPortForwardingSession --parameters '{"portNumber":'
["3389"], "localPortNumber":["56788"]}'
```

3. Once the command says **waiting for connections** you can launch the RDP client and enter **localhost:56788** for the server name and login as **administrator** with the password you set in the previous section.

› Troubleshooting connectivity

4. Once you have RDP'ed into the Windows Instance, launch task manager by right clicking on the menu bar and selecting "Task Manager" (or by using the SHIFT-CTRL-ESC keyboard shortcut). Click on "**More details**" button and then on the "**Performance**" tab so you can see the CPU graph as shown below.



Run CPU Stress Experiment

Note

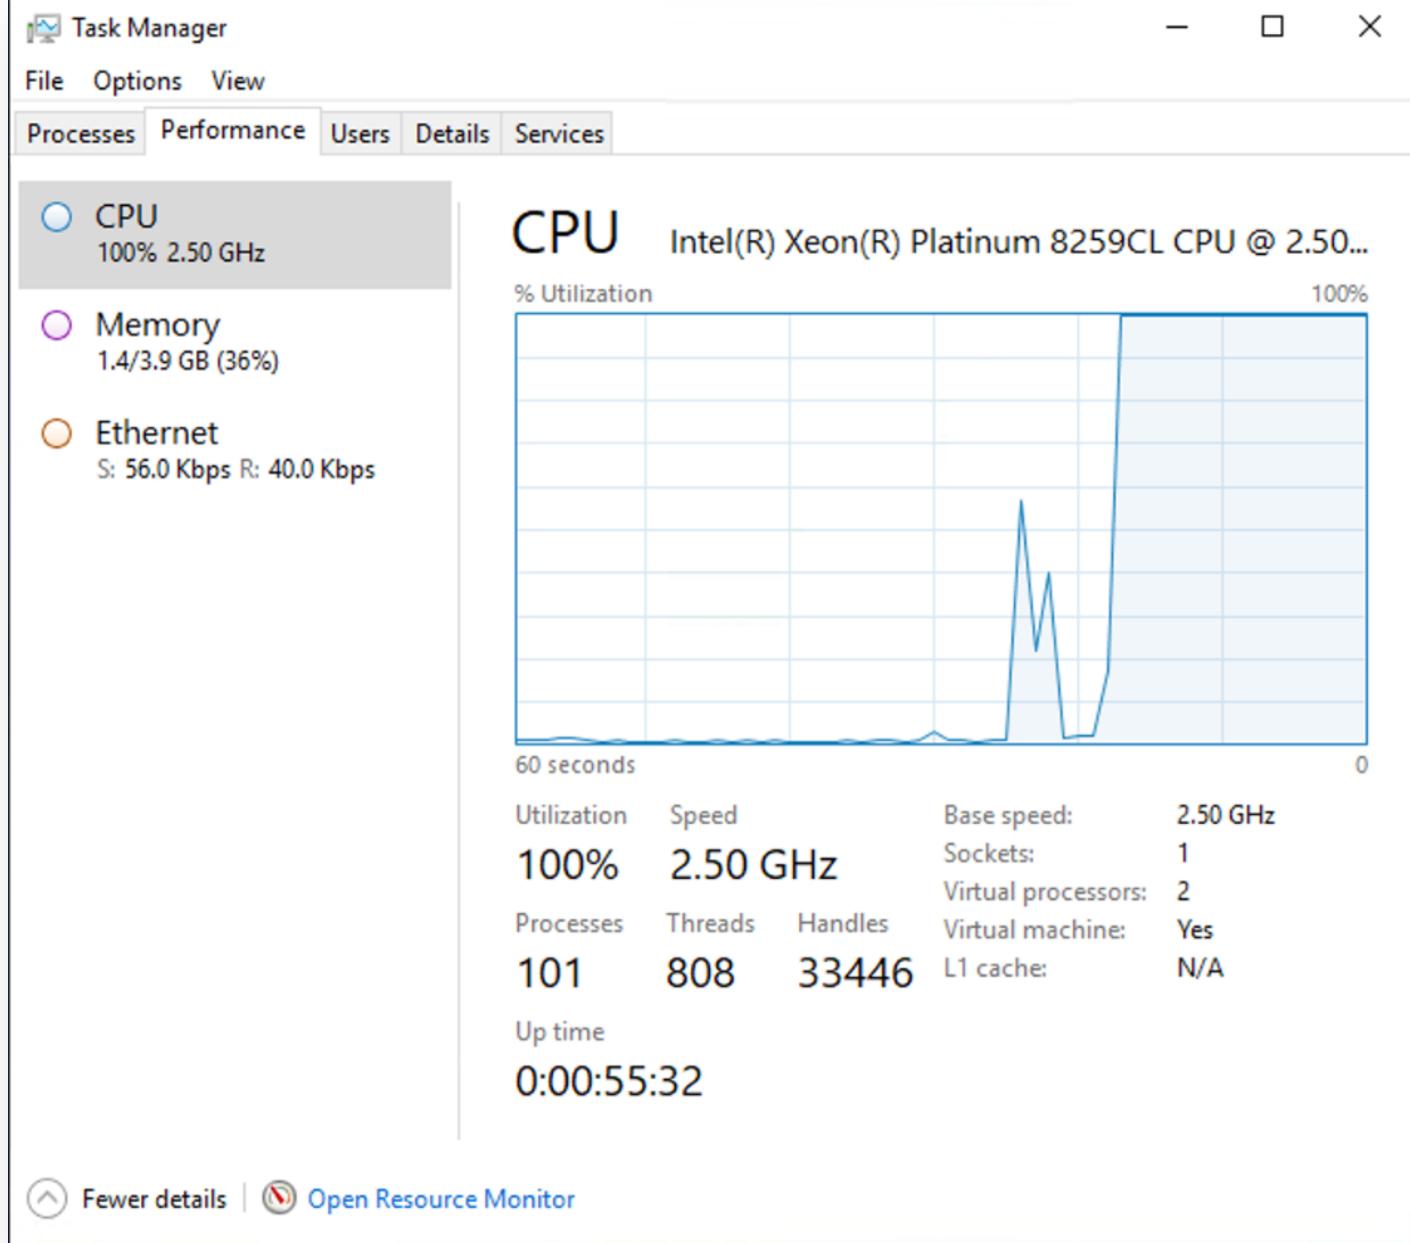
We are assuming that you know how to set up a basic FIS experiment and will focus on things specific to this experiment. If you need a refresher see the previous [First Experiment](#) section.

Keep the RDP session with "Task Manager" running. In a new browser window navigate to the [AWS Fault Injection Simulator Console](#) and start the experiment:

- use the `WindowsBurnCPUviaSSM`
- add a `Name` tag of `FisWorkshopWindowsStress1`
- confirm that you want to start the experiment
- ensure that the "State" is `Running`

Once the experiment is running, lets go back to the RDP session and observe the task manager graph.

Watch the CPU percentage, it should hit 100% for a few minutes and then return back to 0%. Once we have observed the action we can logout of the Windows Instance and hit CTRL + C on the window you ran the port forwarding command to close the session.



Congratulations for completing this lab! In this lab you walked through running an experiment that took action within a Windows EC2 Instance using AWS Systems Manager and a custom run command. Using the integration between Fault Injection Simulator and AWS Systems Manager you can run scripted actions within an EC2 Instance. Through this integration you can script events against your applications or run other chaos engineering tools and frameworks.

Learning and improving

Since this instance wasn't doing anything there aren't any actions. To think about how to use this to test a hypothesis and make an improvement consider building custom SSM scripts to run custom scenarios. We will cover some of these in the [Common Scenarios](#) section.

Cleanup

If you created an additional **CpuStress** CloudFormation stack in the **FIS SSM Setup** section, make sure to delete that stack to avoid incurring additional costs.



SSM ADDITIONAL RESOURCES

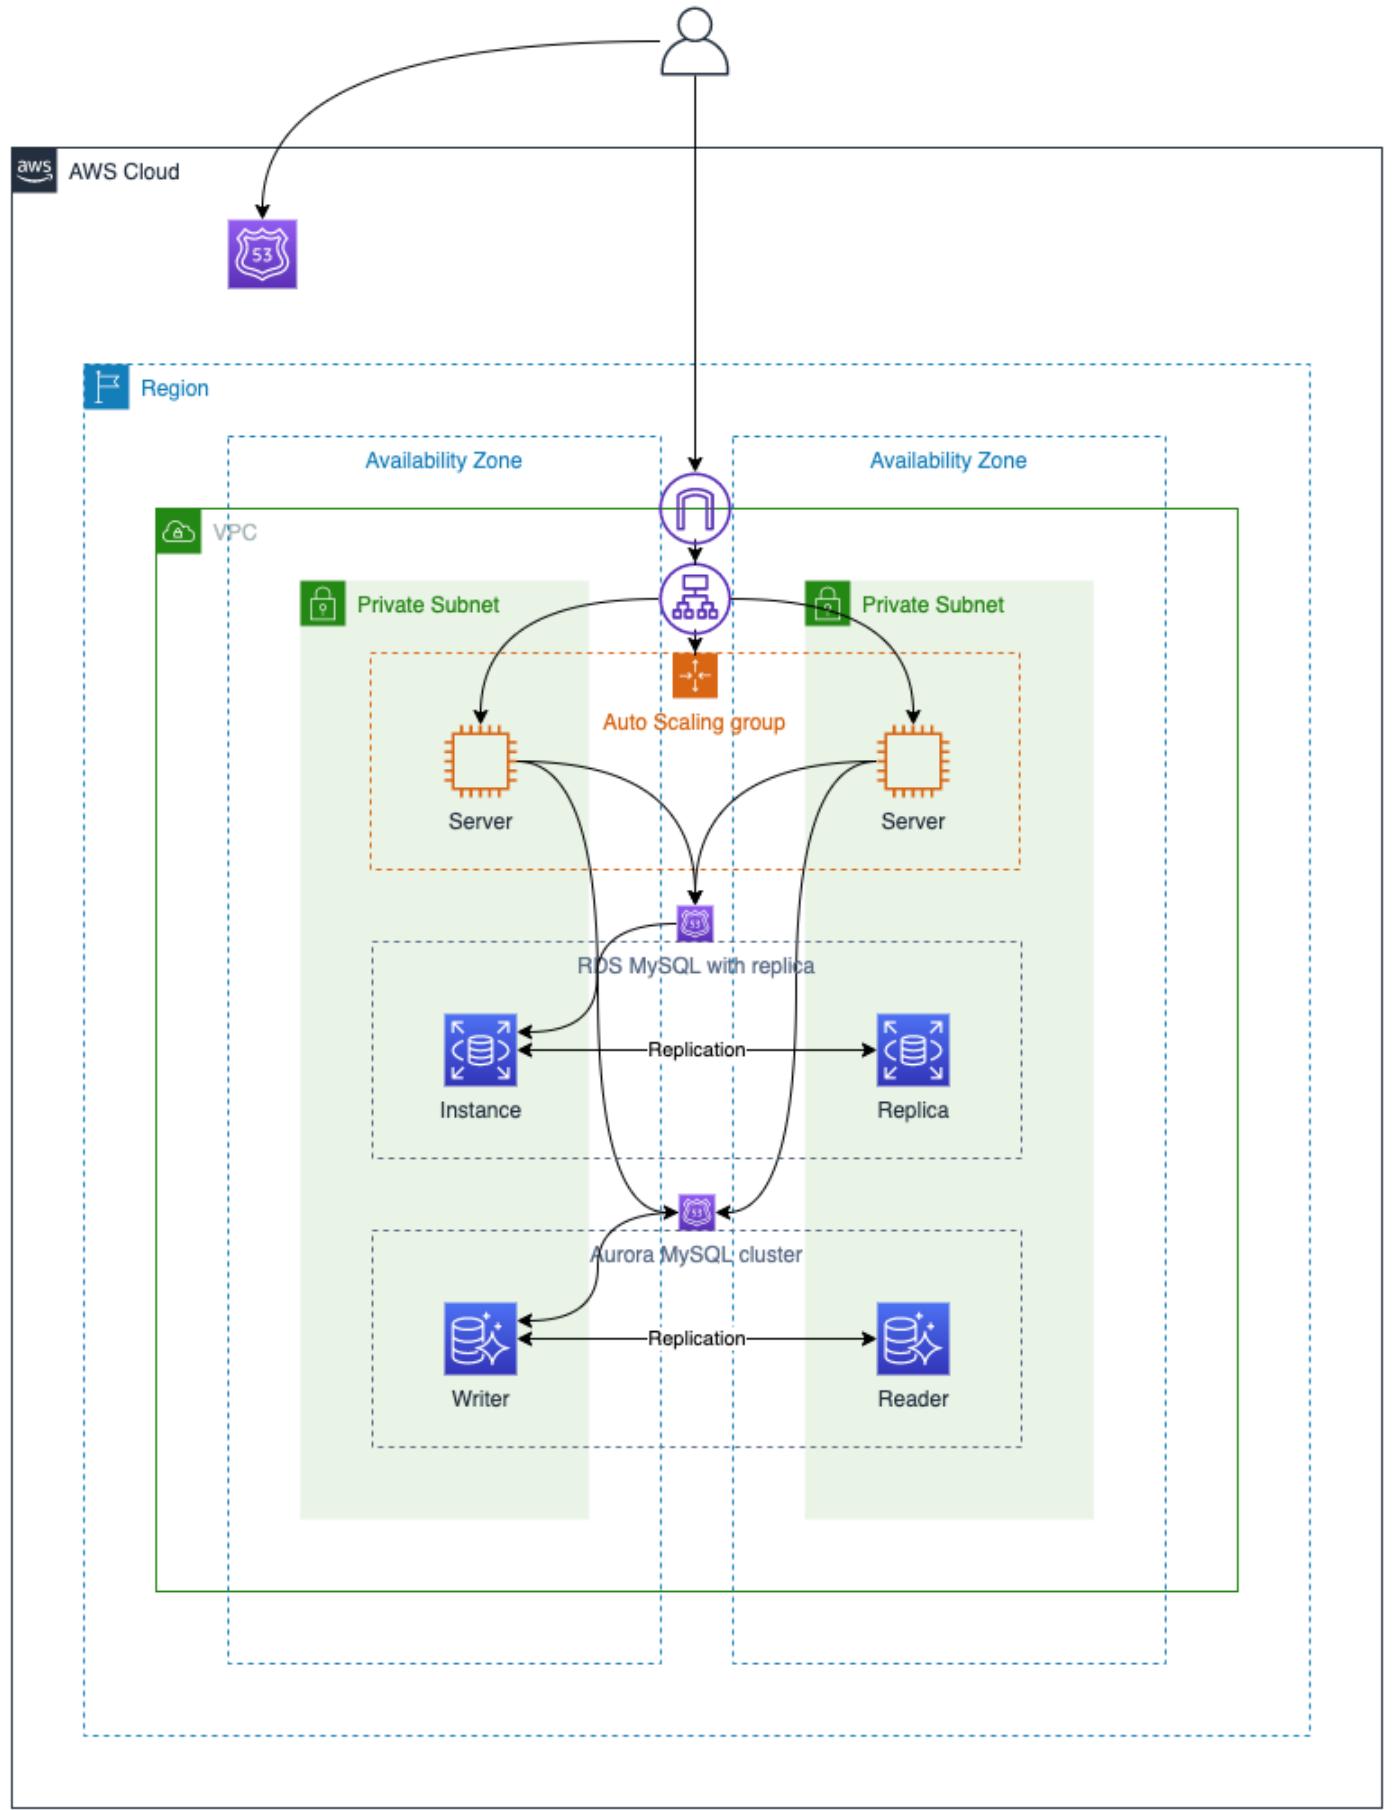
For additional SSM automation resources see:

- [SSM Workshop](#)
- [SSM Chaos Documents](#)
- [SSM Documents AWS documentation](#)
- [SSM working with inputs and outputs](#)



DATABASES

In this section we will cover working with databases. For this setup we are adding RDS MySQL and Aurora for MySQL to our test architecture:



Both RDS MySQL and Aurora for MySQL provide MySQL databases but they are different products. RDS MySQL is a managed service based on stock MySQL while Aurora is a custom built MySQL and PostgreSQL-compatible relational database with better performance and reliability.

Since these are different products they have slightly different failover patterns. They also use slightly different naming conventions:

- For RDS MySQL your dashboard will show "Instances" which may have "Replicas" attached for failover.
- For Aurora MySQL your dashboard will show "Clusters" with "Writers" and "Readers".

For this workshop we are using a similar configuration that replicates data across two AZs for resilience.



RDS DB INSTANCE REBOOT

Experiment idea

In the previous section we ensured that we have a resilient front end of servers in an Auto Scaling group. Typically these servers would depend on a resilient database configuration. Let's validate this:

- **Given:** we have a managed database with a replica and automatic failover enabled
- **Hypothesis:** failure of a single database instance / replica may slow down a few requests but will not adversely affect our application

Experiment setup

Note

We are assuming that you know how to set up a basic FIS experiment and will focus on things specific to this experiment. If you need a refresher see the previous [First Experiment](#) section.

General template setup

- Create a new experiment template
 - Add **Name** tag of `FisWorkshopRds1`
 - Add **Description** of `RebootRDSDInstance`
 - Select `FisworkshopServiceRole` as execution role

Action definition

In the "Actions" section select the **"Add Action"** button.

For "Name" enter `RDSInstanceReboot` and you can skip the Description. For "Action type" select `aws:rds:reboot-db-instances`.

For this experiment we are using a Multi-AZ database and we want to force a failover to the standby instance to minimize outage time. To do this, set the `forceFailover` parameter to `true`.

Leave the default "Target" `DBInstances-Target-1` and select "**Save**".

New action

Name `RDSInstanceReboot` **Description - optional**

Action type
Select the action type to run on your targets. [Learn more](#)
`aws:rds:reboot-db-instances`

Start after - optional
Select actions to run before this action. Otherwise, this action runs as soon as the experiment begins.
`Select an action`

Target
A target will be automatically created for this action if one does not already exist. Additional targets can be created below.
`DBInstances-Target-1`

Action parameters
Specify the parameter values for this action. [Learn more](#)

forceFailover - optional
If instances are Multi-AZ, force a failover from the Availability Zone to another one.
`true`

This action will reboot the main database instance and, due to the "forceFailover" setting, promote the previous replica and update the information associated with the connection string (see below).

Target selection

For this action we need to select our RDS MySQL "Instance". For this we will need to know the instance resource ID. To find this ID open a new browser window and navigate to the [RDS console](#). Note the "DB identifier" for the target DB instance, the one with "Engine" type "MySQL Community".

Filter databases

	DB identifier	Role	Engine
<input type="radio"/>	ffcsbufk247s8	Instance	MySQL Community
<input type="radio"/>	fisstackrdsaurora-fisworkshoprdsauroraeefbf768-gs9ha69qku6a	Regional cluster	Aurora MySQL
<input type="radio"/>	ffa9mkjyj1wpx	Writer instance	Aurora MySQL
<input type="radio"/>	ffhxkk5la659ca	Reader instance	Aurora MySQL

Return to the FIS experiment setup, scroll to the "Targets" section, select **DBInstances-Target-1** and select "**Edit**".

You may leave the default name **Instances-Target-1** but for maintainability we recommend using descriptive target names. Change "Name" to **FisWorkshopRDSDB** for name (this will automatically update the name in the action as well) and make sure "Resource type" is set to **aws:rds:db**.

For "Target method" we will select resources based on the ID. Select the "Resource IDs" checkbox. Under "Resource IDs" pick the target DB instance matching the "DB Identifier" you noted above, then select **All** from "Selection mode". Select "**Save**".

Edit target X

Specify the target resources on which to run your selected actions. [Learn more](#)

Name	Resource type
<input type="text" value="FisWorkshopRDSDB"/>	<input type="text" value="aws:rds:db"/>
Actions	
RDSInstanceReboot	
Target method	
<input checked="" type="radio"/> Resource IDs	
<input type="radio"/> Resource tags and filters	
Resource IDs	Selection mode
<input type="text" value="Select a resource ID"/>	<input type="text" value="ALL"/>
<input type="text" value="ffcsbufk247s8"/> X	
Cancel Save	

Creating template without stop conditions

Select “**Create experiment template**” and confirm that you wish to create a template without stop conditions.

Validation procedure

Before running the experiment we should consider how we will define success. How will we know that our failover was in fact non-impacting. For this workshop we have installed a python script that will read and write data to the database, conceptually like this but with some added safeguards (see [full code in GitHub](#)):

```
import mysql.connector
mydb = mysql.connector.connect(....)
cursor = mydb.cursor()
while True:
    cursor.execute("insert into test (value) values (%d)" %
int(32768*random.random()))
    cursor.execute("select * from test order by id desc limit 10")
    for line in cursor:
        cursor.append("%-30s" % str(line))
```

We would expect that this would keep writing output while the DB is available, stop while it's failing over and restart when the DB has successfully failed over.

Additionally, because the DB connection does a DNS lookup, our script will also print the IP address of the database it's currently connected to. A healthy output should look like this:

AURORA	RDS
10.0.89.224	10.0.95.247
(7711, 2282)	(5419, 15189)
(7710, 5964)	(5418, 15841)
(7709, 10634)	(5417, 8071)
(7708, 4834)	(5416, 21948)
(7707, 20291)	(5415, 27256)
(7706, 9343)	(5414, 8187)
(7705, 5496)	(5413, 9359)
(7704, 30985)	(5412, 6058)
(7703, 21808)	(5411, 26174)
(7702, 20243)	(5410, 21155)

Starting the validation procedure

Connect to one of the EC2 instances in your auto scaling group. In a new browser window - we need to be able to see this side-by-side with the FIS experiment later - navigate to your **EC2 console** and search for instances named `FisStackAsg/ASG`. Select one of the instances and select the “**Connect**” button:

The screenshot shows the AWS EC2 Instances page. At the top, there's a header with "Instances (1/1)" and a "Info" link. Below the header are buttons for "Connect" (which is highlighted with a red box), "Instance state", "Actions", and "Launch instances". There's also a "Filter instances" search bar and some filter options like "Instance state: running" and "search: FisStackAsg/ASG". The main table lists one instance: "FisStackAsg/ASG" with Instance ID "i-054562513f4f2bb4f", status "Running", type "t3.large", and 2/2 checks passed. A checkbox next to the instance name is also highlighted with a red box.

On the next page select “**Session Manager**” and “**Connect**”:

The screenshot shows the "Connect to instance" page for the instance "i-054562513f4f2bb4f (FisStackAsg/ASG)". It has tabs for "EC2 Instance Connect", "Session Manager" (which is highlighted with a red box), "SSH client", and "EC2 Serial Console". Below the tabs, it says "Session Manager usage:" followed by a bulleted list: "Connect to your instance without SSH keys or a bastion host.", "Sessions are secured using an AWS Key Management Service key.", "You can log session commands and details in an Amazon S3 bucket or CloudWatch Logs log group.", and "Configure sessions on the Session Manager Preferences page.". At the bottom right, there are "Cancel" and "Connect" buttons, with "Connect" highlighted with a red box.

This will open a linux terminal session. In this session sudo to assume the `ec2-user` identity:

```
sudo su - ec2-user
```

If this is the first time you are doing this run the `create_db.py` script (review **code in GitHub**) to ensure we can connect to the DB and we have created the required tables:

```
./create_db.py
```

If all went well you should see output similar to this:

AURORA 10.0.89.224 done	RDS 10.0.95.247
-------------------------------	--------------------

Now start the test script and leave it running:

```
./test_mysql_connector_curses.py
```

Run FIS experiment



We are assuming that you know how to set up a basic FIS experiment and will focus on things specific to this experiment. If you need a refresher see the previous [First Experiment](#) section.

Record current RDS state

Navigate to the [RDS console](#), select “**Databases**” on the left menu, and select the “MySQL Community” instance. Note that the current instance state is “Available”:

The screenshot shows the Amazon RDS console interface. On the left, there's a sidebar with links: Dashboard, Databases (which is highlighted in orange), Query Editor, Performance Insights, Snapshots, Automated backups, Reserved instances, and Proxies. The main area shows the RDS > Databases > ffc809i9ltvodd page. The instance name "ffc809i9ltvodd" is displayed prominently. Below it is a "Summary" table with the following data:

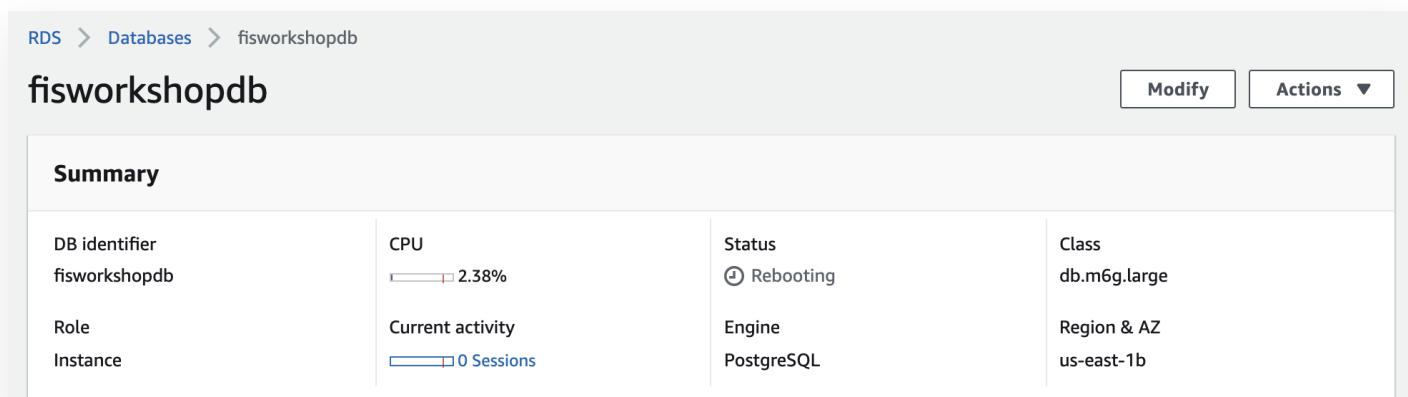
DB identifier	CPU	Status	Class
ffc809i9ltvodd	2.62%	Available	db.t3.micro
Role	Current activity	Engine	Region & AZ
Instance	0 Connections	MySQL Community	us-east-2a

Start the experiment

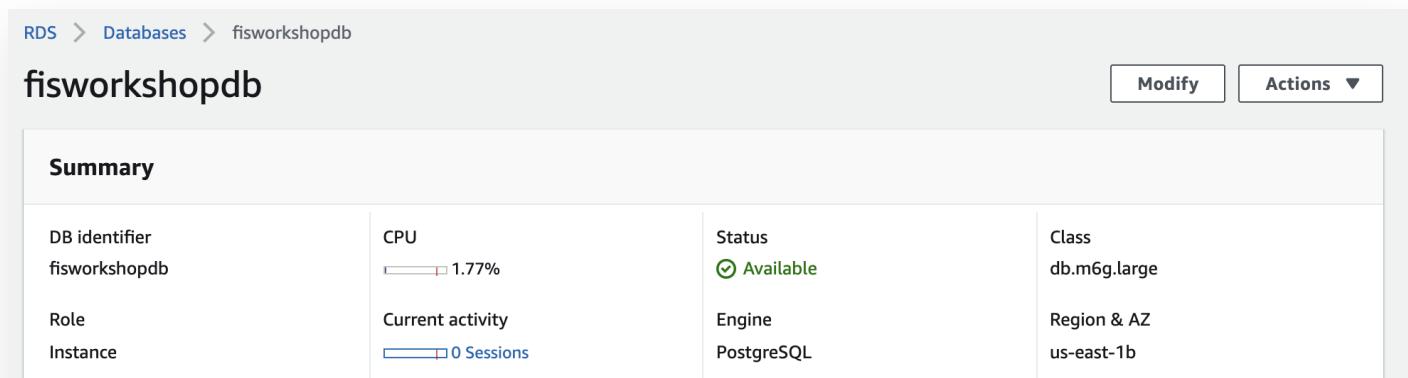
- Select the `FisWorkshopRds1` experiment template you created above
- Select start experiment
- Add a `Name` tag of `FisWorkshopMysql1Run1`
- Confirm that you want to start an experiment
- Watch the output of your test script
- Check the state of your database in the **RDS console**

Review results

If all went "well" the status of the database in the RDS console should have changed from "Available" to "Rebooting"



and back to "Available".



However, even though your database failed over successfully, your script should have locked up during the failover - no more updates to your data and it didn't recover even after the DB successfully failed over.

Learning and Improving

What happened is that our script used a common MySQL database connector library that does not have a `read_timeout` setting. The database successfully failed over but the `INSERT` or `SELECT` statement that was in flight during the failover never timed out and locked our code into waiting forever.

Fortunately there is another common library that has very similar configuration and does implement `read_timeout`. For your convenience we have provided an updated script (review [code in GitHub](#)). CTRL-C out of the hung script and repeat the experiment but this time running

```
./test_pymysql_curses.py
```

This time you should see almost no interruption in your code's ability to interact with the database.

To end the session, hit CTRL+C to stop the script, and click "**Terminate**" button.



AURORA CLUSTER FAILOVER

Experiment idea

In the previous section we ensured that we have a resilient front end of servers in an Auto Scaling group. Typically these servers would depend on a resilient database configuration. Let's validate this:

- **Given:** we have a managed database with a replica and automatic failover enabled
- **Hypothesis:** failure of a single database instance / replica may slow down a few requests but will not adversely affect our application

Experiment setup

Note

We are assuming that you know how to set up a basic FIS experiment and will focus on things specific to this experiment. If you need a refresher see the previous [First Experiment](#) section.

General template setup

- Create a new experiment template
 - Add **Name** tag of `FisWorkshopAurora1`
 - Add **Description** of `FailoverAuroraCluster`
 - Select `FisworkshopServiceRole` as execution role

Action definition

In the "Actions" section select the **"Add Action"** button.

For "Name" enter `FisWorkshopFailoverAuroraCluster` and add a "Description" like `Failover Aurora Cluster`. For "Action type" select `aws:rds:failover-db-cluster`.

Leave the default "Target" `Clusters-Target-1` and select **"Save"**.

▼ New action

SaveRemove**Name**

FisWorkshopFailoverAuroraCluster

Description - optional

Failover Aurora Cluster

Action typeSelect the action type to run on your targets. [Learn more](#)

aws:rds:failover-db-cluster

Start after - optional

Select actions to run before this action. Otherwise, this action runs as soon as the experiment begins.

Select an action

Target

A target will be automatically created for this action if one does not already exist. Additional targets can be created below.

Clusters-Target-1

Target selection

For this action we need to select our Aurora "Cluster". For this we will need to know the instance resource ID. To find this ID open a new browser window and navigate to the [RDS console](#)). Note the "DB identifier" for the target cluster, the one with "Engine" type "Aurora MySQL" and "Role" "Regional Cluster".

Databases		<input checked="" type="checkbox"/> Group resources		Modify	Actions ▾	Restore from
		<input type="text"/> Filter databases				
	DB identifier		▲	Role	▼	Engine
<input type="radio"/>	ffcslnbufk247s8			Instance		MySQL Community
<input type="radio"/>	fisstackrdsaurora-fisworkshoprdsaurorae7bf768-gs9ha69qku6a			Regional cluster		Aurora MySQL
<input type="radio"/>	ffa9mkjyj1wpx			Writer instance		Aurora MySQL
<input type="radio"/>	ffhxkk5la659ca			Reader instance		Aurora MySQL

Return to the FIS experiment setup, scroll to the "Targets" section, select **Clusters-Target-1** and select "Edit".

You may leave the default name **Clusters-Target-1** but for maintainability we recommend using descriptive target names. Change "Name" to **FisWorkshopAuroraCluster** for name (this will automatically update the name in the action as well) and make sure "Resource type" is set to **aws:rds:cluster**.

For “Target method” we will select resources based on the ID. Select the “Resource IDs” checkbox. Under “Resource IDs” pick the target DB instance matching the “DB Identifier” you noted above, then select **All** from “Selection mode”. Select **“Save”**.

The screenshot shows the 'Add target' dialog box. At the top, it says 'Specify the target resources on which to run your selected actions.' with a 'Learn more' link. Below that, there are two main sections: 'Name' and 'Resource type'. The 'Name' field contains 'FisWorkshopAuroraCluster'. The 'Resource type' field is set to 'aws:rds:cluster'. Under 'Target method', the radio button for 'Resource IDs' is selected, while 'Resource tags and filters' is unselected. In the 'Resource IDs' section, a dropdown menu shows 'Select a resource ID'. Below it, a list box contains the resource ID 'fisstackrdsaurora-fisworkshoprdsauroraee7bf768-gs9ha69qku6a / cluster-JPGL5RHN6JFB2LI72UP5TT73HQ'. To the right, under 'Selection mode', the dropdown is set to 'All'. At the bottom right, there are 'Cancel' and 'Save' buttons.

Creating template without stop conditions

Select **“Create experiment template”** and confirm that you wish to create a template without stop conditions.

Validation procedure

The validation procedure is identical to what we did in the **RDS DB Instance Reboot** section. If you have not explored that section before, perform the steps as described there under the “Validation Procedure” heading and return here when you reach the “Run FIS experiment” heading.

Run FIS experiment

We are assuming that you know how to set up a basic FIS experiment and will focus on things specific to this experiment. If you need a refresher see the previous [First Experiment](#) section.

Record current Aurora state

Navigate to the [RDS console](#), select “**Databases**” on the left menu, and search for “fisworkshop”. Take a screenshot or write down the “Reader” and “Writer” AZ information, e.g.:

DB identifier	Role	Engine	Region & AZ
fisworkshop	Regional	Aurora PostgreSQL	us-east-1
fisworkshop-instance-1	Writer	Aurora PostgreSQL	us-east-1d
fisworkshop-instance-1-us-east-1a	Reader	Aurora PostgreSQL	us-east-1a

Start the experiment

- Select the `FisWorkshopAurora1` experiment template you created above
- Select start experiment
- Add a `Name` tag of `FisWorkshopAurora1Run1`
- Confirm that you want to start an experiment
- Watch the output of your test script

Review results

Verify that the experiment worked. If you are not already on the pane viewing your experiment, navigate to the [FIS console](#), select “**Experiments**”, and select the experiment ID for the experiment you just started. This should show “success”.

Verify that the failover actually happened. Navigate to the RDS console again and about a minute after you started the experiment you’ll see the “Reader” and “Writer” instances flipped to the other AZ:

DB Identifier	Role	Engine	Region & AZ
fisworkshop	Regional	Aurora PostgreSQL	us-east-1
fisworkshop-instance-1-us-east-1a	Writer	Aurora PostgreSQL	us-east-1a
fisworkshop-instance-1	Reader	Aurora PostgreSQL	us-east-1d

If all went well, the “Reader” and “Writer” instances should have traded places.

If you were watching the output of your test script carefully you might also have noticed that for a short period of time DNS returns no value for Aurora. To address this our code already contains an additional try/except block for DB reconnection (see [code in GitHub](#)).

Learning and improving

As this was essentially the same as the previous [RDS DB Instance Reboot](#) section there are no new learnings here.

However, you may want to experiment further with built-in Aurora [fault injection queries](#).

To access your Aurora database, you can extract the connection information from the [AWS Secrets Manager console](#) by selecting the `FisAuroraSecret` and selecting “**Retrieve secret value**”:



Using the information you can open another terminal, e.g. from the same instance you were using for testing, and connect to your Aurora database with the retrieved secret values:

```
mysql -h HOST -u USERNAME -p DBNAME
```

you can then run fault injection queries as further explained in this [blog post](#) and observe the effect on the test script, e.g.:

```
ALTER SYSTEM CRASH NODE;
```

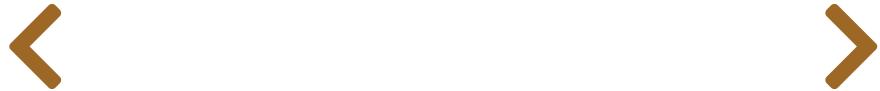
<

>

CONTAINERS

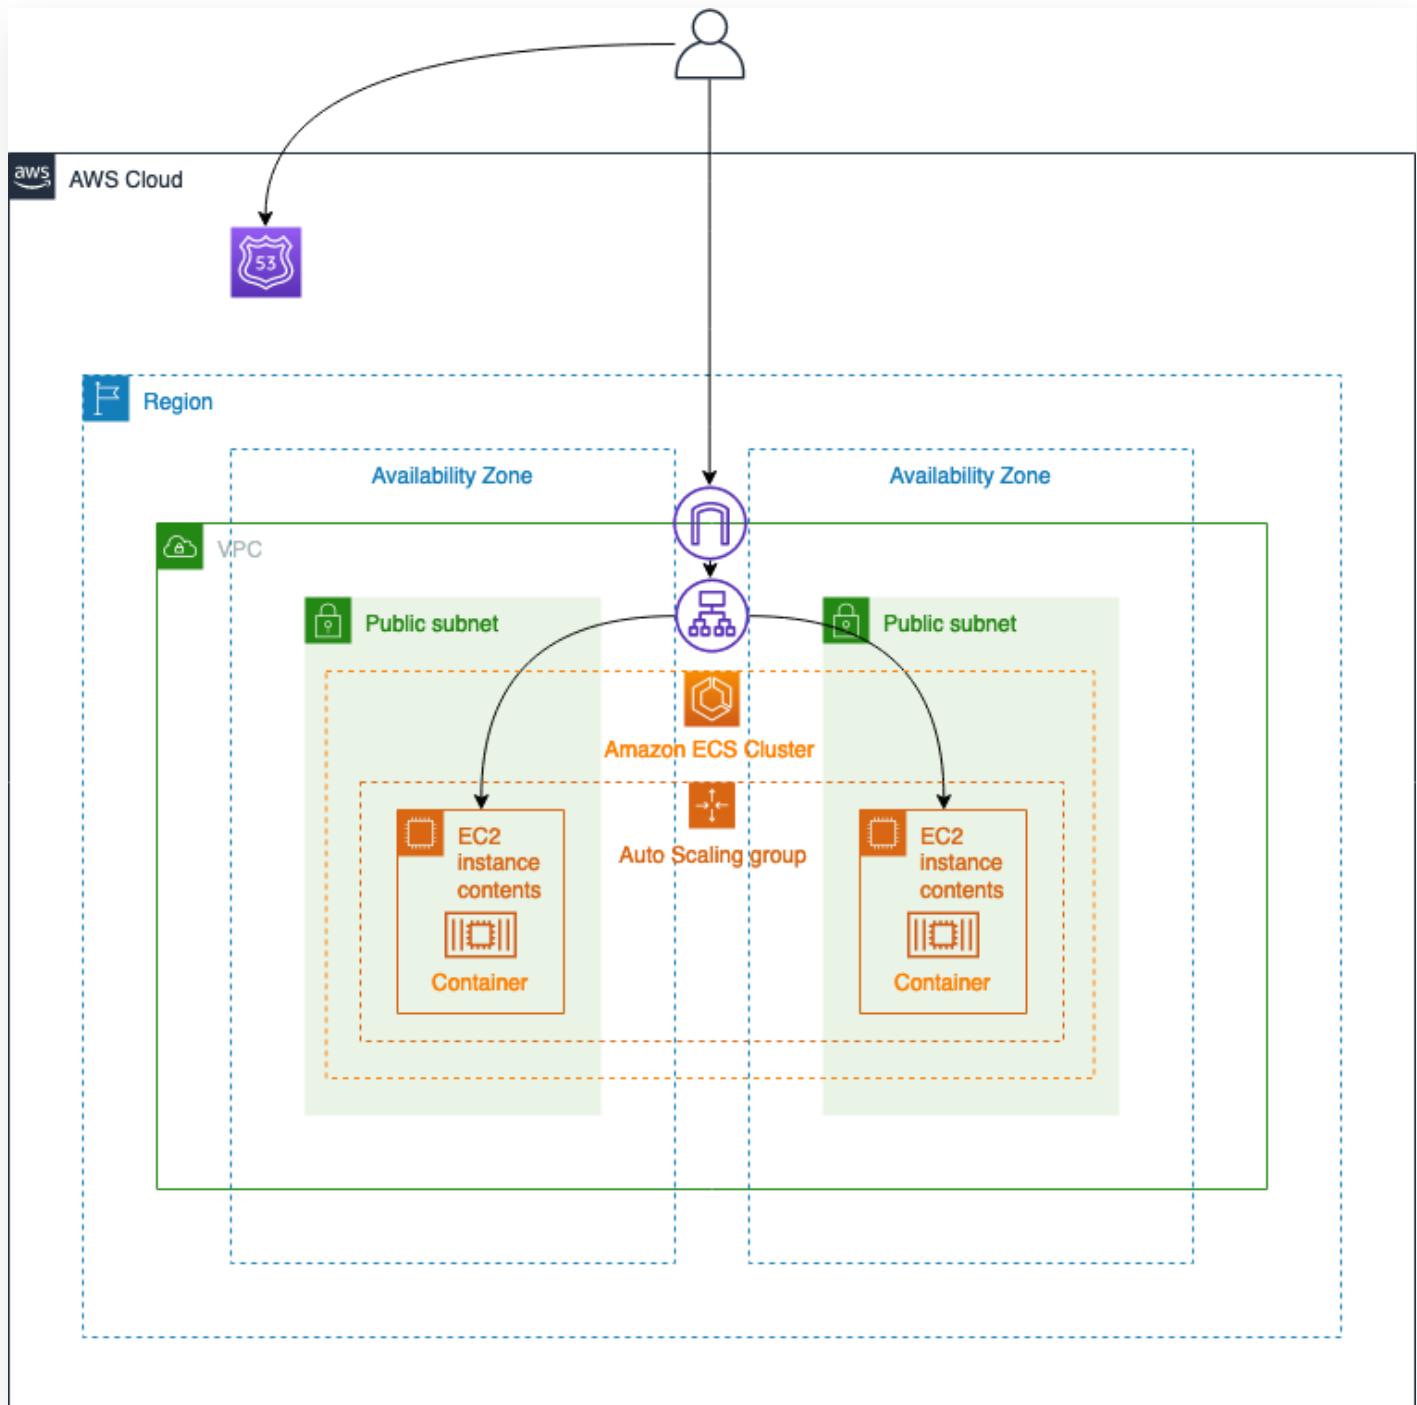
In addition to the ability to run containers on virtual machines, AWS provides managed services to run containers at scale. This section covers fault injection experiments on:

- [Amazon ECS](#)
- [Amazon EKS](#)



AMAZON ECS

In this section we will cover working with containers running on **Amazon Elastic Container Service** (Amazon ECS). For this setup we'll be using the following test architecture:



Amazon ECS is a fully managed container orchestration service that helps you easily deploy, manage, and scale containerized applications. It deeply integrates with the rest of the AWS platform to provide a secure and easy-to-use solution for running container workloads in the cloud and now on your infrastructure with Amazon ECS Anywhere.



HYPOTHESIS & EXPERIMENT

Experiment idea

In this section we want to ensure that our containerized application running on Amazon ECS is designed in a fault tolerant way, so that even if an instance in the cluster fails our application is still available. Let's validate this:

- **Given:** we have a containerized application running on Amazon ECS exposing a web page.
- **Hypothesis:** failure of a single container instance will not adversely affect our application. The web page will continue to be available.

Experiment setup

Note

We are assuming that you know how to set up a basic FIS experiment and will focus on things specific to this experiment. If you need a refresher see the previous [First Experiment](#) section.

General template setup

Note

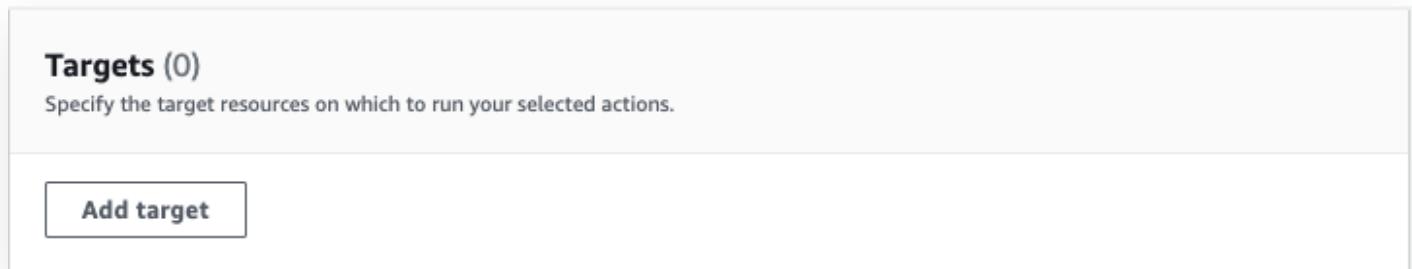
We are assuming that you have already set up an IAM role for this workshop. If you haven't, see the [Create FIS Service Role](#) section.

Create a new experiment template:

- add `Name` tag of `FisWorkshopECS`
- add `Description` of `Terminate ECS Cluster Instance`
- select `FisWorkshopServiceRole` as execution role

Target selection

Now we need to define targets. Scroll to the "Targets" section and select "Add Target"



On the "Add target" popup enter `FisWorkshopECSInstance` for name and select `aws:ec2:instance` for resource type. For "Target method" we will dynamically select resources based on an associated tag. Select the `Resource tags and filters` checkbox. Pick `Count` from "Selection mode" and enter `1`. Under "Resource tags" enter `Name` in the "Key" field and `FisStackEcs/EcsAsgProvider` for "Value". Under filters enter `State.Name` in the "Attribute path" field and `running` under "Values". Select "Save".

Add target

X

Specify the target resources on which to run your selected actions. [Learn more](#)

Name

FisWorkshopECSInstance

Resource type

aws:ec2:instance

▼

Target method

- Resource IDs
- Resource tags and filters

Selection mode

Count

Number of resources

1

Resource tags

Key

Value - optional

Name

FisStackEcs/EcsAsgProvider

Remove

Add new tag

Resource filters - optional

Filter resources by the attributes you specify. [Learn more](#)

Attribute path

Values

State.Name

running

Remove

Separate multiple values with commas.

Add new filter

Cancel

Save

Action definition

With targets defined we define the action to take. Scroll to the "Actions" section and select "Add Action"

Actions (0)

Specify one or more actions to run on your target resources. Decide how long to run each action (in minutes), and when to start the action during the experiment. [Learn more](#)

Add action

For "Name" enter `ECSInstanceTerminate` and you can skip the Description. For "Action type" select `aws:ec2:terminate-instances`.

We will leave the "Start after" section blank since the instances we are terminating are part of an auto scaling group and we can let the auto scaling group create new instances to replace the terminated ones.

Under "Target" select the `FisworkshopECSInstance` target created above. Select "Save".

Actions (1)
Specify one or more actions to run on your target resources. Decide how long to run each action (in minutes), and when to start the action during the experiment. [Learn more](#)

▼ New action

Name **Description - optional**

Action type **Start after - optional**

Select the action type to run on your targets. [Learn more](#)

Target
A target will be automatically created for this action if one does not already exist. Additional targets can be created below.

Add action

Save **Remove**

Creating template without stop conditions

Confirm that you wish to create the template without stop condition.

Create experiment template

X

⚠ You have not specified a stop condition for your experiment template. A stop condition can help to prevent your experiment from going out of bounds by stopping it automatically. [Learn more ↗](#)

To confirm that you want to create an experiment template without a stop condition, enter *create* in the field:

create

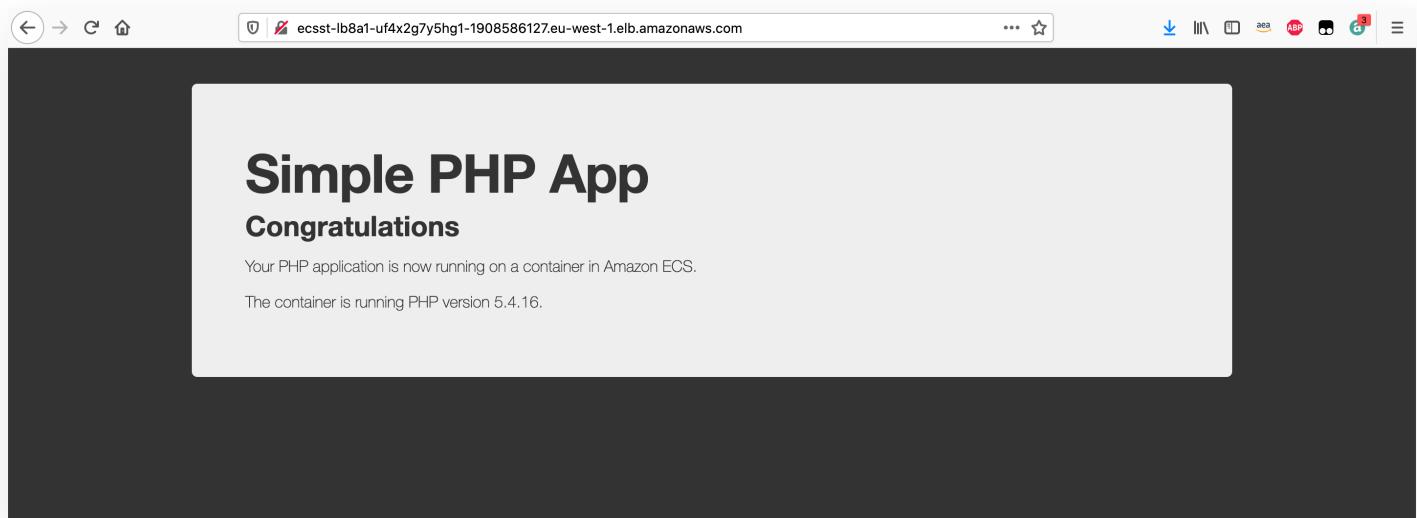
Cancel

Create experiment template

Validation procedure

Before running the experiment we should consider how we will define success. Let's check the webpage we are hosting. To find the URL of the webpage navigate to the [CloudFormation console](#), select the [FisStackEcs](#) stack, Select "Outputs", and copy the value of "FisEcsUrl".

Open the URL in a new tab to validate that our website is in fact up and running:



How will we know that our instance failure was in fact non-impacting? For this workshop we'll be using a simple Bash script that continuously polls our application.

Starting the validation procedure

In your local terminal, run the following script. For your convenience we are automating the query for the load balancer URL but you could also paste the URL you've found above:

```
# Query URL for convenience
ECS_URL=$( aws cloudformation describe-stacks --stack-name FisStackEcs --query
"Stacks[*].Outputs[?OutputKey=='FisEcsUrl'].OutputValue" --output text )

# Busy loop queries. CTRL-C to end loop
while true; do
    curl -sLo /dev/null -w 'Code %{response_code} Duration %{time_total}\n'
${ECS_URL}
done
```

We would expect that all requests will return a [HTTP 200](#) OK code with some variability in the request duration, meaning the application is still responding successfully. Healthy output should look like this:

```
Code 200 Duration 0.140314
Code 200 Duration 0.086206
Code 200 Duration 0.085946
Code 200 Duration 0.084102
Code 200 Duration 0.085972
```

Leave the script running while we run the FIS experiment next.

Run FIS experiment

Record current application state

In a new browser window navigate to the load balancer URL you copied earlier, this is your application endpoint. Notice that the application is currently running:

Simple PHP App

Congratulations

Your PHP application is now running on a container in Amazon ECS.

The container is running PHP version 5.4.16.

You can also verify the HTTP return code using this command, replacing `REPLACE_WITH_ECS_SERVICE_ALB_URL` with the load balancer DNS name you copied earlier:

```
curl -IL <REPLACE_WITH_ECS_SERVICE_ALB_URL> | grep "HTTP\/"
```

Start the experiment

- Select the `FisWorkshopECS` experiment template you created above
- Select **Start experiment** from the **Action** drop-down menu
- Add a `Name` tag of `FisWorkshopECSRun1`
- Confirm that you want to start an experiment

Start experiment



⚠ You are about to start your experiment, which might perform destructive actions on your AWS resources. Before you run fault injection experiments, review the best practices and planning guidelines. [Learn more](#)

To confirm that you want to start the experiment, enter `start` in the field:

`start`

[Cancel](#)

[Start experiment](#)



OBSERVE THE SYSTEM

Review results

Let's take a look at the output in the terminal window where your Bash script is running:

```
Code 200 Duration 0.137204
Code 200 Duration 0.080911
Code 200 Duration 0.081539
Code 200 Duration 0.077265
Code 200 Duration 0.085331
Code 200 Duration 0.081634
```

...

```
Code 503 Duration 0.083001
Code 503 Duration 0.088983
Code 502 Duration 0.085972
Code 502 Duration 0.086619
Code 502 Duration 0.086554
Code 503 Duration 0.083428
Code 502 Duration 0.084929
```

...

```
Code 200 Duration 0.082434
Code 200 Duration 0.081427
Code 200 Duration 0.087983
Code 200 Duration 0.081950
Code 200 Duration 0.082790
```

You'll notice that as not all the requests were successful. As the FIS experiment starts you should see some [HTTP 502](#) "Bad Gateway" and [HTTP 503](#) "Service Unavailable" return codes. This means our application was not available for a period of time. This is not what we were expecting, so let's dive a bit deeper to find out why it happened.

Check number of containers

In a new browser window navigate to the *Clusters* section in the [ECS console](#) and search for the cluster named [FisStackEcs-Cluster...](#), e.g. [FisStack-ClusterEB0386A7-xJ4yY19a5jLP](#). Click on the cluster name and look at the ECS services running on this cluster:

The screenshot shows the AWS ECS Cluster details page for the cluster `EcsStack-ClusterEB0386A7-xJ4yY19a5jLP`. The cluster ARN is listed as `arn:aws:ecs:eu-west-1:560846014933:cluster/EcsStack-ClusterEB0386A7-xJ4yY19a5jLP`. The status is **ACTIVE**. Registered container instances: 1. Pending tasks count: 0 Fargate, 0 EC2, 0 External. Running tasks count: 0 Fargate, 1 EC2, 0 External. Active service count: 0 Fargate, 1 EC2, 0 External. Draining service count: 0 Fargate, 0 EC2, 0 External.

Services tab selected. Table of services:

Action	Service Name	Status	Service type	Task Definition	Desired tasks...	Running task...	Launch type	Platform vers...
<input type="checkbox"/>	EcsStack-ServiceD69D759B-PsBz3nNuocPp	ACTIVE	REPLICA	<code>EcsStackTask...</code>	1	1	EC2	--

You'll notice that the service named [FisStackEcs-SampleAppService...](#), e.g.

[FisStackEcs-SampleAppServiceD69D759B-PsBz3nNuocPp](#) - i.e. our application - only has **one** desired task, meaning that only one copy of our containerized application will be running at any time.

A screenshot of the AWS ECS Services table. The table shows a single row for the service `EcsStack-ServiceD69D759B-PsBz3nNuocPp`. The **Desired tasks...** column is highlighted with a red circle.

Action	Service Name	Status	Service type	Task Definition	Desired tasks...	Running task...	Launch type	Platform vers...
<input type="checkbox"/>	EcsStack-ServiceD69D759B-PsBz3nNuocPp	ACTIVE	REPLICA	<code>EcsStackTask...</code>	1	1	EC2	--

Check number of instances

Now click on the "ECS Instances" tab. You'll see here that there's only one instance registered with our cluster.

An Amazon ECS instance is either an External instance registered using ECS Anywhere or an Amazon EC2 instance.

To register an External instance, choose Register External Instances and follow the steps. [Learn More](#)

To register an Amazon EC2 instance, you can use the Amazon EC2 console. [Learn More](#)

Register External Instances

Actions ▾

Last updated on July 26, 2021 1:03:56 PM (0m ago)



Status: [ALL](#) ACTIVE DRAINING

< 1-1 > Page size 50 ▾

Filter by attributes (click or press down arrow to view filter options)

<input type="checkbox"/> Container Instance	ECS Instance	Availability Zon...	External Insta...	Agent Connec...	Status	Running tasks...	CPU available	Memory availa
<input type="checkbox"/> 1db2d2bc84fd40e19acded...	i-0315f3d6b712...	eu-west-1a	false	true	ACTIVE	1	2048	3372

Observations

This configuration is not optimal:

- A cluster with a single instance means that if that instance fails, all the containers running on that instance will also be killed. This is what happened during our experiment and the reason why we observed some [HTTP 503 "Service Unavailable"](#) return codes. We should change this so that our cluster has more than one instance across multiple Availability Zones (AZs).
- Having an ECS Service with **one** desired task also means that if that task fails, there aren't any other tasks to continue serving requests. We can modify this by adjusting the desired task capacity to **2** (or any number greater than **1**).

Now that we have identified some issues with our current setup, let's move to the next section to fix them.



IMPROVE & REPEAT

Learning and Improving

In the previous section we have identified some issues with our current setup: our ECS cluster only had **one** instance and our application's ECS Service desired capacity was set to **1**. Now, let's improve our infrastructure setup.

Increase the number of instances

In our ECS configuration we have chosen to use EC2 with an auto scaling group as our capacity provider. To adjust desired instance capacity open a browser window and navigate to the *Auto Scaling Groups* section in the [EC2 console](#) and search for an auto scaling group named [FisStackEcs-EcsAsgProvider...](#), e.g. [FisStackEcs-EcsAsgProviderASG51CCF8BD-4L06D3044727](#). Select the check box next to our Auto Scaling group. A split pane opens up in the bottom part of the Auto Scaling groups page, showing information about the group that's selected.

The screenshot shows the AWS Auto Scaling Groups page. At the top, there is a header with 'EC2 > Auto Scaling groups'. Below the header, there is a table titled 'Auto Scaling groups (1/2)' with one item listed. The table has columns: Name, Launch template/configuration, Instances, Status, Desired capacity, Min, and Max. The item listed is 'EcsStack-ClusterDefaultAutoScaling...' with a status of '1', '1', '1', '1', and '1'. Below the table, there are tabs for Details, Activity, Automatic scaling, Instance management, Monitoring, and Instance refresh. The 'Details' tab is selected. In the lower pane, there is a section titled 'Group details' with an 'Edit' button. It contains the following information:

Desired capacity	Auto Scaling group name
1	EcsStack-ClusterDefaultAutoScalingGroupCapacityASGA16CBFC4-19 CDXR3FUOQO
Minimum capacity	Date created
1	Mon Jul 26 2021 11:33:14 GMT+0100 (British Summer Time)
Maximum capacity	Amazon Resource Name (ARN)
1	arn:aws:autoscaling:eu-

In the lower pane, in the **Details** tab and under **Group details** section, click the **Edit** button.

- Change the current settings for “minimum” to **2** to ensure we always have at least 2 instances available for redundancy. Note: if you only increase “desired” and “maximum” then the scaling policy for the auto scaling group could decrease the “desired” value back to **1** during low load periods.
- Set “desired” and “maximum” to **2** or more. Note: setting the desired value to more than the number of tasks (see below) will leave you with idle instances.
- Click **Update** to complete the changes:

Group size

Specify the size of the Auto Scaling group by changing the desired capacity. You can also specify minimum and maximum capacity limits. Your desired capacity must be within the limit range.

Desired capacity

Minimum capacity

Maximum capacity

Cancel **Update**

Increase the number of tasks

Navigate to the *Clusters* section in the [ECS console](#) and search for the cluster named

FisStackEcs-Cluster..., e.g. **FisStackEcs-ClusterEB0386A7-xJ4yY19a5jLP**. Click on the cluster name and look at the ECS service named **FisStackEcs-SampleAppService...**, e.g.

FisStackEcs-SampleAppServiceD69D759B-PsBz3nNuocPp, running on this cluster. Select the check box next to our ECS Service and click **Update**:

Services	Tasks	ECS Instances	Metrics	Scheduled Tasks	Tags	Capacity Providers
Create	Update	Delete	Actions ▾	Last updated on July 26, 2021 1:20:30 PM (0m ago)		
Filter in this page		Launch type	ALL	Service type	ALL	1 selected

Scroll to the bottom of the *Configure service* screen and change the value of the **Number of tasks** setting from **1** to **2**. Click **Skip to review** and complete the process by selecting **Update Service**.

Configure service

A service lets you specify how many copies of your task definition to run and maintain in a cluster. You can optionally use an Elastic Load Balancing load balancer to distribute incoming traffic to containers in your service. Amazon ECS maintains that number of tasks and coordinates task scheduling with the load balancer. You can also optionally use Service Auto Scaling to adjust the number of tasks in your service.

Task Definition Family [Enter a value](#)

Revision

Launch type EC2 [i](#)

[Switch to capacity provider strategy](#)

Force new deployment [i](#)

Cluster EcsStack-ClusterEB0386A7-xJ4yY... [i](#)

Service name EcsStack-ServiceD69D759B-PsBz... [i](#)

Service type* REPLICA [i](#)

Number of tasks [i](#)

© 2008 - 2021, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Repeat the experiment

Now that we have improved our configuration, let's re-run the experiment. Before starting review the ECS Cluster to ensure that the instance capacity has increased to **2** and that the number of running tasks is **2**.

This time we should observe that, even when one of the container instances gets terminated, our application is still available and successfully serving requests. In the output of the Bash script there we should no longer see the [HTTP 503 "Service Unavailable"](#) return codes.

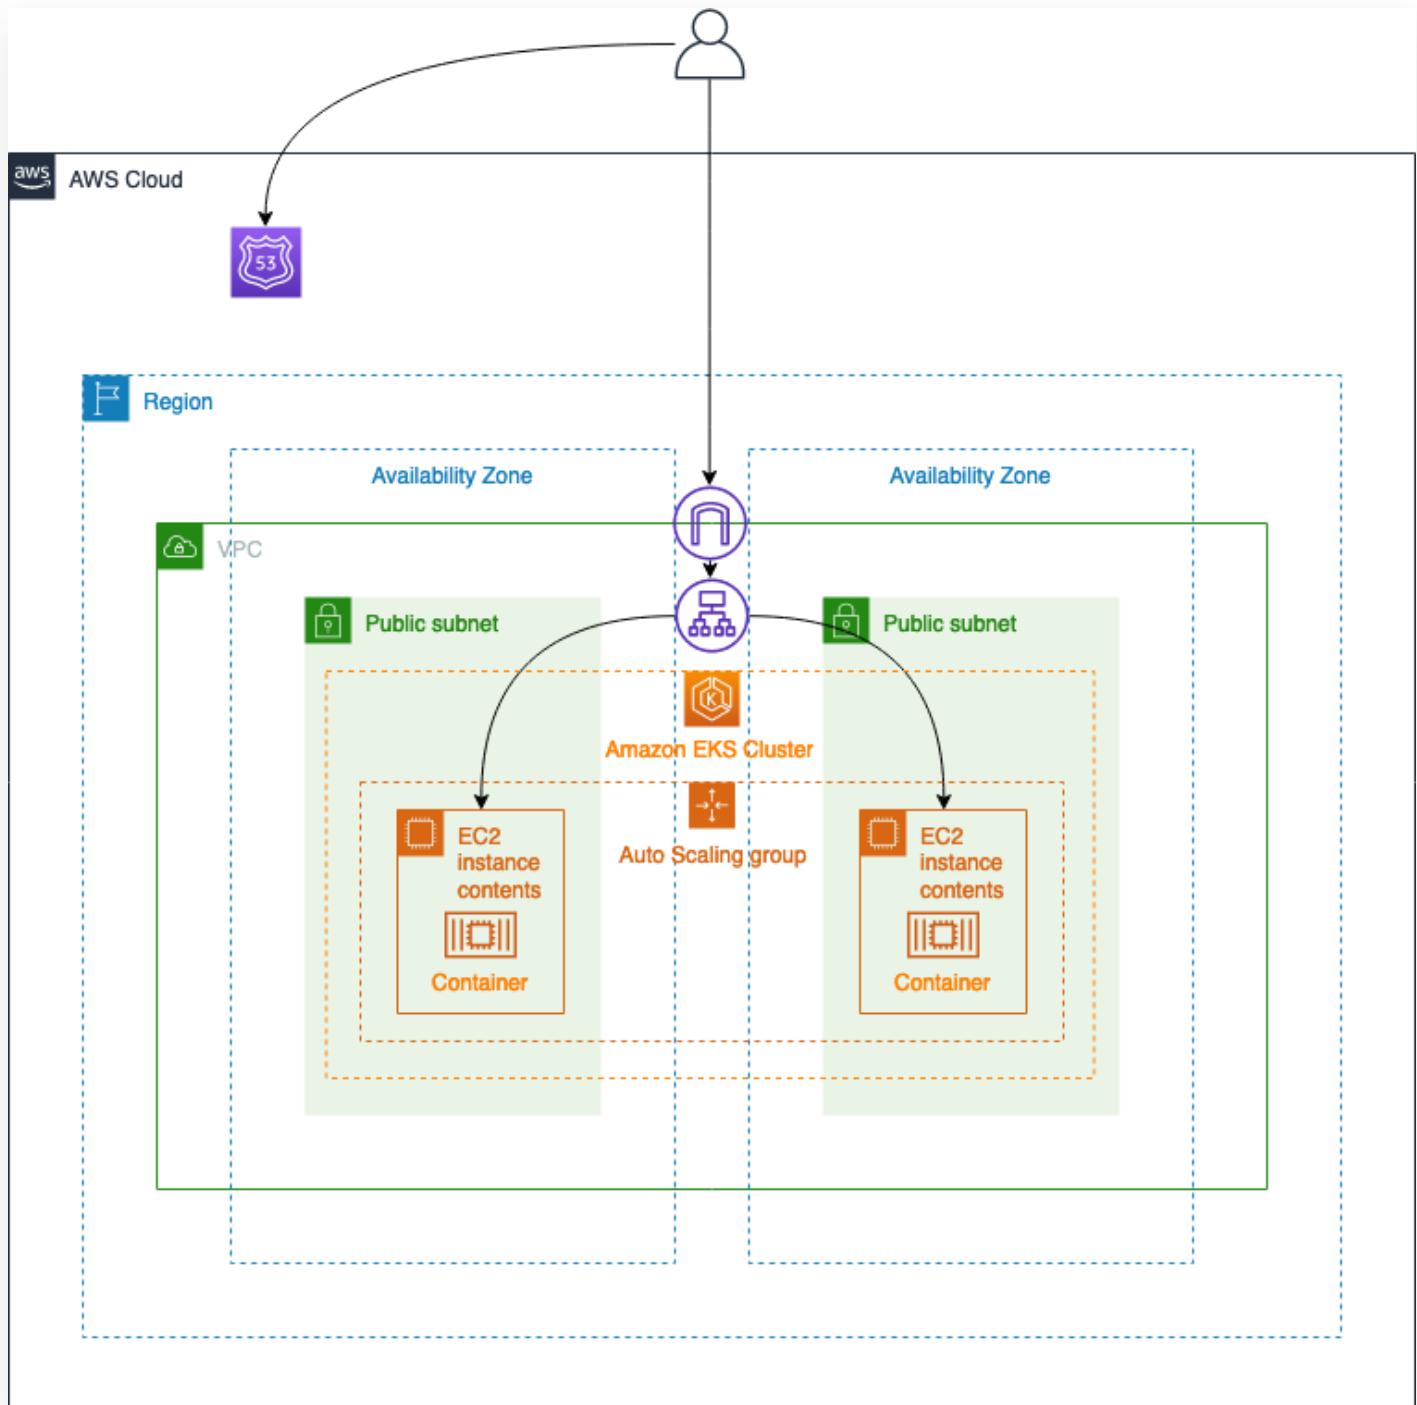
ECS FURTHER LEARNING

For more on ECS configurations see the [ECS workshop](#).



AMAZON EKS

In this section we will cover working with containers running on [Amazon Elastic Kubernetes Service \(Amazon EKS\)](#). For this setup we'll be using the following test architecture:



Amazon EKS gives you the flexibility to start, run, and scale Kubernetes applications in the AWS cloud or on-premises. Amazon EKS helps you provide highly-available and secure clusters and automates key tasks such as patching, node provisioning, and updates. EKS runs upstream Kubernetes and is certified Kubernetes conformant for a predictable experience. You can easily migrate any standard Kubernetes application to EKS without needing to refactor your code.

 Note

For this section, make sure you have `kubectl` installed in your local environment. Follow [these steps](#) if you need to install `kubectl`.



HYPOTHESIS & EXPERIMENT

Experiment idea

In this section we want to ensure that our containerized application running on Amazon EKS is designed in a fault tolerant way, so that even if an instance in the cluster fails our application is still available. Let's validate this:

- **Given:** we have a containerized application running on Amazon EKS exposing a web page.
- **Hypothesis:** failure of a single worker node instance will not adversely affect our application. The web page will continue to be available.

Experiment setup

Note

We are assuming that you know how to set up a basic FIS experiment and will focus on things specific to this experiment. If you need a refresher see the previous [First Experiment](#) section.

General template setup

Note

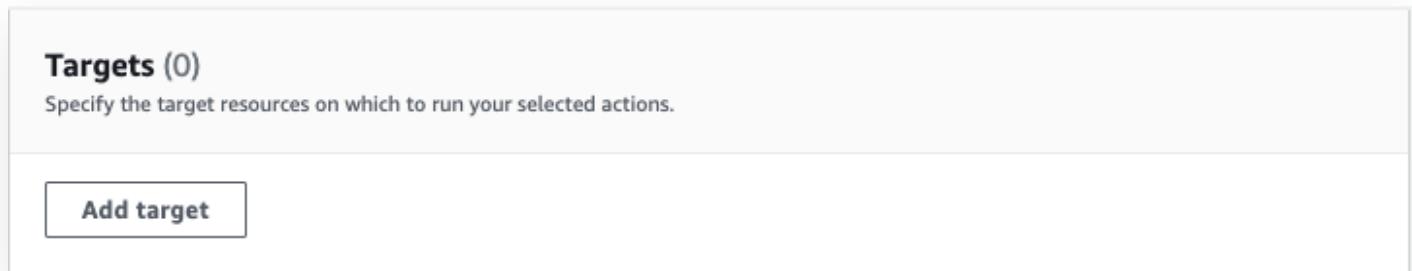
We are assuming that you have already set up an IAM role for this workshop. If you haven't, see the [Create FIS Service Role](#) section.

Create a new experiment template:

- add `Name` tag of `FisWorkshopEKS`
- add `Description` of `Terminate EKS Worker Node`
- select `FisWorkshopServiceRole` as execution role

Target selection

Now we need to define targets. Scroll to the "Targets" section and select "Add Target"



On the "Add target" popup enter `FisWorkshopEKSWorkerNode` for name and select `aws:ec2:instance`. For "Target method" we will dynamically select resources based on an associated tag. Select the `Resource tags and filters` checkbox. Pick `Count` from "Selection mode" and enter `1`. Under "Resource tags" enter `eks:nodegroup-name` in the "Key" field and `FisWorkshopNG` for "Value". Under filters enter `State.Name` in the "Attribute path" field and `running` under "Values". Select "Save".

Add target

X

Specify the target resources on which to run your selected actions. [Learn more](#)

Name

FisWorkshopEKSWorkerNode

Resource type

aws:ec2:instance

▼

Target method

- Resource IDs
- Resource tags and filters

Selection mode

Count

Number of resources

1

↑

Resource tags

Key

eks:nodegroup-name

Value - optional

FisWorkshopNG

Remove

Add new tag

Resource filters - optional

Filter resources by the attributes you specify. [Learn more](#)

Attribute path

State.Name

Values

running

Remove

Separate multiple values with commas.

Add new filter

Cancel

Save

Note: we are using the `aws:ec2:instance` action instead of the `aws:eks:nodegroup` action because currently the latter cannot terminate a single running worker node.

Action definition

With targets defined we define the action to take. Scroll to the "Actions" section and select "Add Action"

Actions (0)

Specify one or more actions to run on your target resources. Decide how long to run each action (in minutes), and when to start the action during the experiment. [Learn more](#)

Add action

For "Name" enter `EKSWorkerNodeTerminate` and you can skip the Description. For "Action type" select `aws:ec2:terminate-instances`.

We will leave the "Start after" section blank since the instances we are terminating are part of an EKS Managed Node Group and we can let the Managed Node Group create new instances to replace the terminated ones.

Under "Target" select the `FisWorkshopEKSWorkerNode` target created above. Select "Save".

Actions (1)

Specify one or more actions to run on your target resources. Decide how long to run each action (in minutes), and when to start the action during the experiment. [Learn more](#)

▼ New action

Save

Remove

Name

EKSWorkerNodeTerminate

Description - optional

Action type

Select the action type to run on your targets. [Learn more](#)

aws:ec2:terminate-instances

Start after - optional

Select actions to run before this action. Otherwise, this action runs as soon as the experiment begins.

Select an action

Target

A target will be automatically created for this action if one does not already exist. Additional targets can be created below.

FisWorkshopEKSWorkerNode

Add action

Creating template without stop conditions

Confirm that you wish to create the template without stop condition.

Create experiment template

X

⚠️ You have not specified a stop condition for your experiment template. A stop condition can help to prevent your experiment from going out of bounds by stopping it automatically. [Learn more ↗](#)

To confirm that you want to create an experiment template without a stop condition, enter *create* in the field:

create

Cancel

Create experiment template

Validation procedure

Before running the experiment we should consider how we will define success. Let's check the webpage we are hosting. To find the URL of the webpage navigate to the [CloudFormation console](#), select the [FisStackEks](#) stack, Select "Outputs", and copy the value of "FisEksUrl".

Open the URL in a new tab to validate that our website is in fact up and running:



How will we know that our instance failure was in fact non-impacting? For this workshop we'll be using a simple Bash script that continuously polls our application.

Starting the validation procedure

In your local terminal, run the following script. For your convenience we are automating the query for the load balancer URL but you could also paste the URL you've found above:

```
# Query URL for convenience
EKS_URL=$( aws cloudformation describe-stacks --stack-name FisStackEks --query
"Stacks[*].Outputs[?OutputKey=='FisEksUrl'].OutputValue" --output text )

# Busy loop queries. CTRL-C to end loop
while true; do
    curl -sLo /dev/null -w 'Code %{response_code} Duration %{time_total}\n'
${EKS_URL}
done
```

We would expect that all requests will return a [HTTP 200](#) OK code with some variability in the request duration, meaning the application is still responding successfully. Healthy output should look like this:

```
Code 200 Duration 0.140314
Code 200 Duration 0.086206
Code 200 Duration 0.085946
Code 200 Duration 0.084102
Code 200 Duration 0.085972
```

Leave the script running while we run the FIS experiment next.

Run FIS experiment

Record current application state

In a new browser window navigate to the load balancer URL you copied earlier, this is your application endpoint. Notice that the application is currently running:

pod: hello-kubernetes-ffd764cf9-csfzx
node: Linux (5.4.129-63.229.amzn2.x86_64)

You can also verify the HTTP return code using this command, replacing `REPLACE_WITH_EKS_SERVICE_ALB_URL` with the load balancer DNS name you copied earlier:

```
curl -IL <REPLACE_WITH_EKS_SERVICE_ALB_URL> | grep "HTTP\/"
```

Start the experiment

- select the `FisWorkshopEKS` experiment template you created above
- select **Start experiment** from the **Action** drop-down menu
- add a `Name` tag of `FisWorkshopEKSRun1`
- confirm that you want to start an experiment

Start experiment

⚠ You are about to start your experiment, which might perform destructive actions on your AWS resources. Before you run fault injection experiments, review the best practices and planning guidelines. [Learn more](#)

To confirm that you want to start the experiment, enter `start` in the field:

Cancel **Start experiment**

<

>

OBSERVE THE SYSTEM

Review results

Let's take a look at the output in the terminal window where your Bash script is running:

```
Code 200 Duration 0.137204  
Code 200 Duration 0.080911  
Code 200 Duration 0.081539  
Code 200 Duration 0.077265  
Code 200 Duration 0.085331  
Code 200 Duration 0.081634
```

...

```
Code 000 Duration 0.093033  
Code 000 Duration 0.088688  
Code 000 Duration 0.086454  
Code 000 Duration 0.088505  
Code 000 Duration 0.097665
```

...

```
Code 200 Duration 0.082434  
Code 200 Duration 0.081427  
Code 200 Duration 0.087983  
Code 200 Duration 0.081950  
Code 200 Duration 0.082790
```

You'll notice that not all the requests were successful. As the FIS experiment starts you should see some **000** return codes. This is not a legal HTTP response code. If we just ran curl as

```
curl $EKS_URL
```

we would see an error message indicating that the server just closed the connection on us.

```
curl: (52) Empty reply from server
```

In practice this means our application was not available for a period of time. This is not what we were expecting, so let's dive a bit deeper to find out why it happened.

Configure kubectl

Note

Make sure you have `kubectl` installed in your local environment. Follow [these steps](#) if you need to install `kubectl`.

We will follow [these steps](#) to update the `kubectl` configuration to securely connect to the EKS cluster. The cluster is named `FisWorkshop-EksCluster`. To find the ARN of the kubectl access role, navigate to the [CloudFormation console](#), select the `FisStackEks` stack, Select “Outputs”, and copy the value of “FisEksKubectlRole”.

From a local terminal, run the following command to configure kubectl:

```
# verify you have aws CLI installed
aws --version

# Retrieve the role ARN
KUBECTL_ROLE=$( aws cloudformation describe-stacks --stack-name FisStackEks --
query "Stacks[*].Outputs[?OutputKey=='FisEksKubectlRole'].OutputValue" --output
text )

# Configure kubectl with cluster name and ARN
aws eks update-kubeconfig --name FisWorkshop-EksCluster --role-arn
${KUBECTL_ROLE}
```

Note

If you get the message **“error: You must be logged in to the server (Unauthorized)”** when running `kubectl` command, please follow [these steps](#) to troubleshoot the problem.

Check number of containers

From a local terminal, run the following command to check our application service configuration:

```
kubectl get pods
```

You'll notice that there's only one pod named `hello-kubernetes-...` - e.g.

`hello-kubernetes-ffd764cf9-zwnq7` - meaning that only one copy of our containerized application is running at any time.

NAME	READY	STATUS	RESTARTS	AGE
<code>hello-kubernetes-ffd764cf9-zwnq7</code>	1/1	Running	0	8m34s

Check number of instances

In the same terminal, run the following command to check the nodes in our cluster:

```
kubectl get nodes
```

In the output you'll see that our cluster only has a single worker node.

NAME	STATUS	ROLES	AGE	VERSION
<code>ip-10-0-150-147.eu-west-1.compute.internal</code>	Ready	<none>	12m	v1.20.4-eks-6b7464

Observations

This configuration is not optimal:

- A cluster with a single worker node means that if that instance fails, all the containers running on that instance will also be killed. This is what happened during our experiment and the reason why we observed some `curl: (52) Empty reply from server` messages. We should change this so that our cluster has more than one instance across multiple Availability Zones (AZs).
- An EKS workload with **one** pod also means that if that pod fails, there aren't any other pods to continue serving requests. We can modify this by adjusting the pod count to **2** (or any number greater than **1**).

Now that we have identified some issues with our current setup, let's move to the next section to fix them.



IMPROVE & REPEAT

Learning and Improving

In the previous section we have identified some issues with our current setup: our EKS cluster only had **one** worker node and our application's pod count was set to **1**. Now, let's improve our infrastructure setup.

Increase the number of instances

In a browser window navigate to the *Clusters* section in the [EKS console](#) and search for the cluster named **FisWorkshop-EksCluster**. Click on the cluster name, select the *Configuration* tab and then the *Compute* tab. In the *Node Groups* section, select the round check box next to the group named **FisWorkshopNG** and click **Edit**.

The screenshot shows the AWS EKS Cluster Configuration page for the cluster "FisWorkshop-EksCluster". The "Compute" tab is selected. In the "Node Groups" section, there is one group named "FisWorkshopNG" with a status of "Active".

Group name	Desired size	AMI release version	Launch template	Status
FisWorkshopNG	1	1.20.4-20210722	-	Active

On the *Edit node group* page

- Change the current settings for "minimum" to **2** to ensure we always have at least 2 instances available for redundancy. Note: if you only increase "desired" and "maximum" then the scaling policy for the auto scaling group could decrease the "desired" value back to **1** during low load periods.
- Set "desired" and "maximum" to **2** or more. Note: setting the desired value to more than the number of tasks (see below) will leave you with idle instances.

Node Group scaling configuration

Minimum size

Set the minimum number of nodes that the group can scale in to.

nodes

Minimum node size must be greater than or equal to 0

Maximum size

Set the maximum number of nodes that the group can scale out to.

nodes

Maximum node size must be greater than or equal to 1 and cannot be lower than the minimum size

Desired size

Set the desired number of nodes that the group should launch with initially.

nodes

Desired node size must be greater than or equal to 0

When you're finished editing, scroll to the bottom and choose **Save changes**.

Increase the number of containers

From a local terminal, run the following command to update the application's pod count to **2**:

```
kubectl scale --current-replicas=1 --replicas=2 deployment/hello-kubernetes
```

To verify, you can run `kubectl get pods` and `kubectl get deployments`. Here's the sample output.

NAME	READY	STATUS	RESTARTS	AGE
hello-kubernetes-ffd764cf9-5v7z9	1/1	Running	0	25s
hello-kubernetes-ffd764cf9-6bdbn	1/1	Running	0	4m43s

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
hello-kubernetes	2/2	2	2	46h

Repeat the experiment

Now that we have improved our configuration, let's re-run the experiment. Before starting review the EKS Cluster to ensure that the instance capacity has increased to 2 and that the number of running containers is 2.

This time we should observe that, even when one of the container instances gets terminated, our application is still available and successfully serving requests. In the output of the Bash script there should be no curl: (52) Empty reply from server messages.

EKS/k8s cluster auto scaling

In this workshop we used manual scaling of both worker nodes and pods. In a production setup you would likely configure kubernetes / EKS to use

- a Cluster Autoscaler that is aware of scaling needs based on pod configuration.
- a Horizontal Pod Autoscaler to dynamically manage the number of pods .
- a Vertical Pod Autoscaler to dynamically manage CPU and memory allocation on your pods.

For more on EKS configurations see the [EKS workshop](#).



SERVERLESS

FIS currently does not support disrupting serverless execution in AWS Lambda. It is, however, possible to inject chaos actions by decorating the code executed within AWS Lambda. For a deeper introduction to the topic see:

- [Serverless Chaos workshop](#)
- [GitHub aws-lambda-chaos-injection](#)



RECURRENT EXPERIMENTS - CI/CD

So far we have discussed iterating through a cycle of:

- establish baseline performance data
- develop *new* hypothesis
- run experiment
- verify hypothesis
- improve based on findings

In this section we will address use cases in which we want test and *existing* hypothesis multiple times. Common examples for this are:

- ensure the system remains resilient after changes (CI/CD)
- ensure detection and recovery continue to work (Disaster Recovery)

“Experiment” or “Test”?

A deep dive into testing terminology is outside the scope of this workshop but for readers familiar with the field we want to point out some analogies and provide some considerations:

Human-led processes

The hypothesis based cycle we've discussed up to this point is very similar to "Exploratory Testing" and "Acceptance Testing" in the sense that it steps away from purely validating that something "works as intended". Just like "Exploratory Testing" and "Acceptance Testing", the human-led fault injection process should allow for human observation to adjust the "intent".

Machine-led processes

Automating fault injection based on prior validation of a hypothesis is analogous to the wide range automated and recurrent tests such as:

- unit tests
- regression tests
- integration tests
- load tests

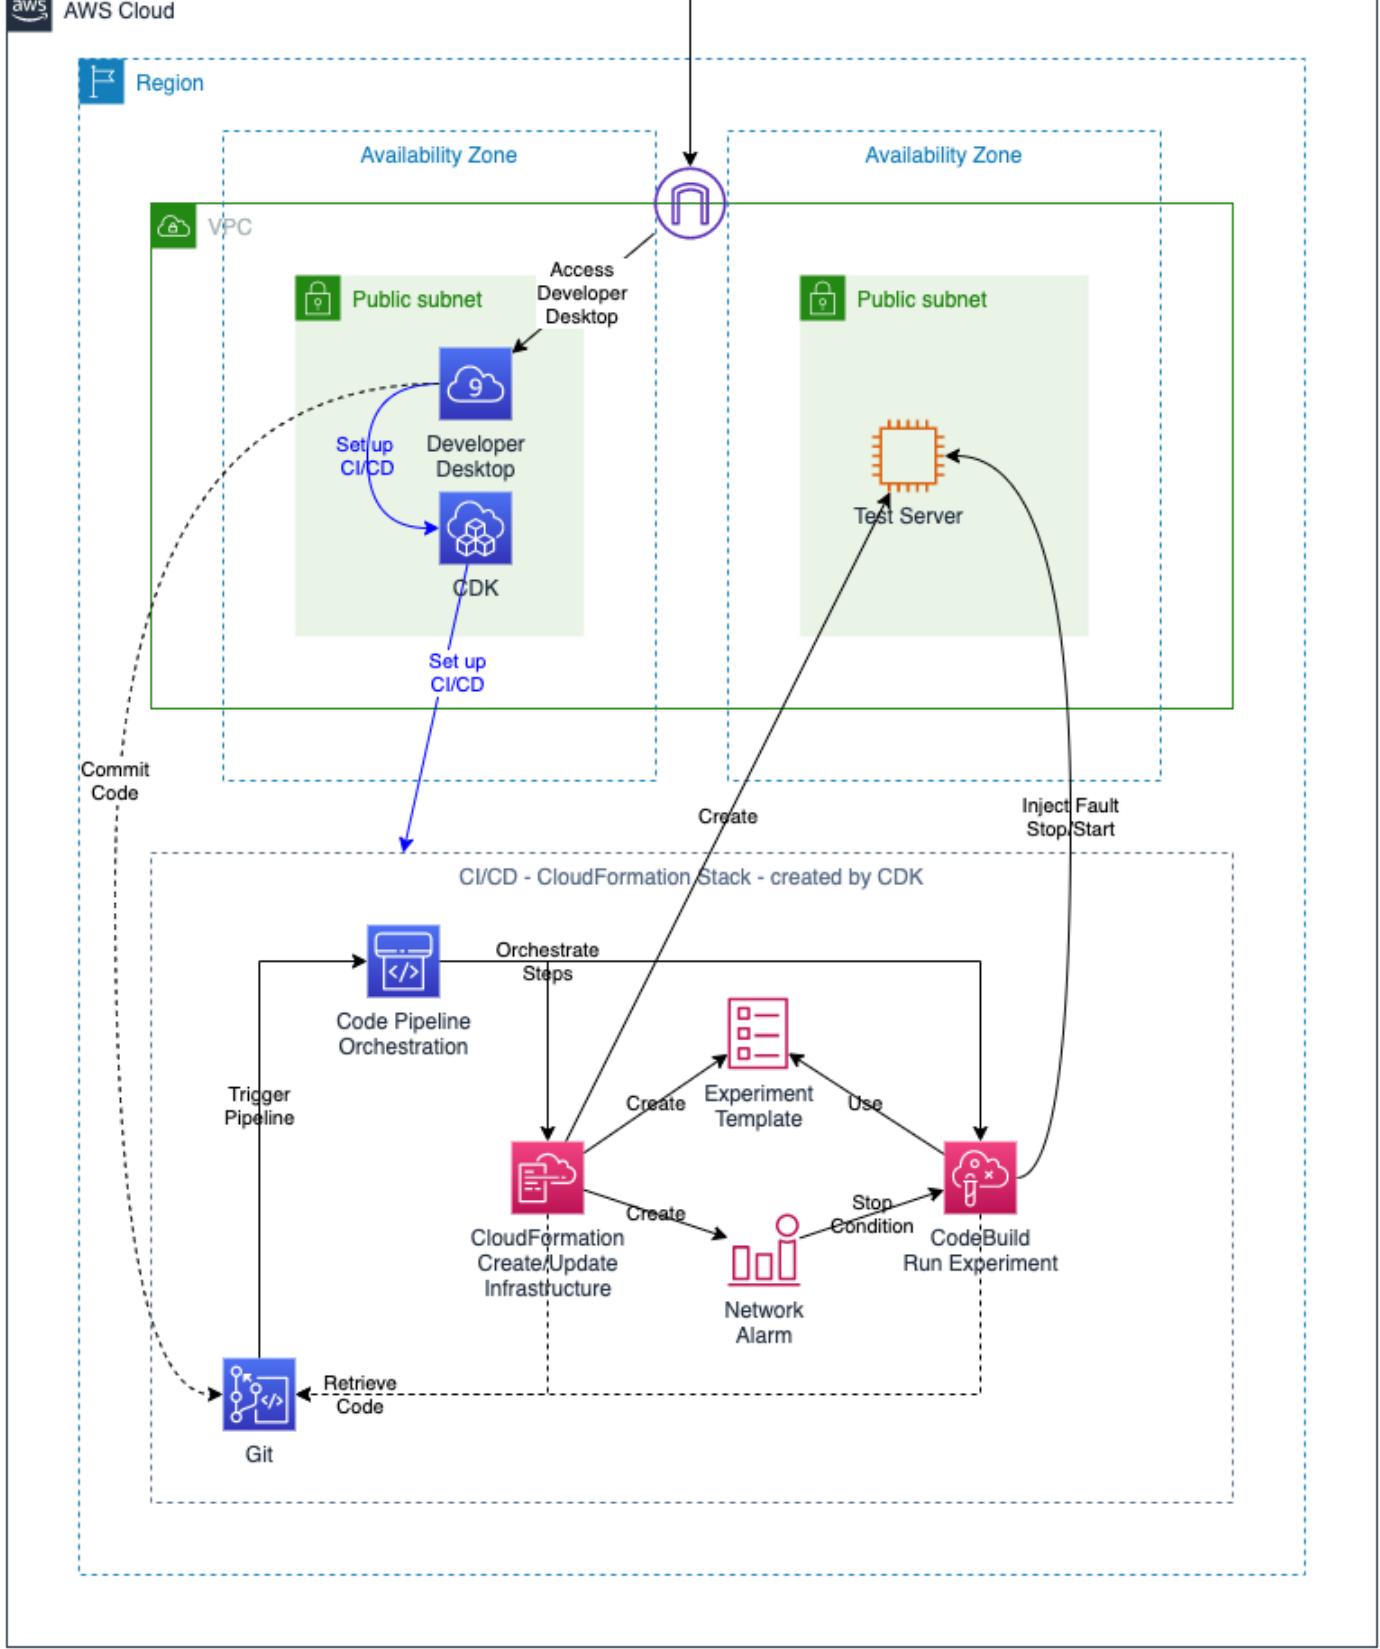
Just like for other tests, it is important to consider the scope and *duration* of recurrent fault injection experiments. Because fault injection experiments generally expose issues across a large number of linked systems they will typically require extended run times to ensure sufficient data collection. In order to not slow down developers they should be run in later stages of CI/CD pipelines.

Architecture

For demonstration purposes we have made the following choices but there are many other ways to build valuable automation:

- **CI/CD** - We focus on running experiments in a CI/CD pipeline with the argument that it's easy to slow down a pipeline to run only once a year but hard to speed up a manual process to run multiple times every day.
- **One repo** - We use a single repository to host the definition of the pipeline, the infrastructure, and fault injection template. We do this to show how one would co-version all components of a system but whether this is a good idea for you depends on your governance processes and each of the parts could easily be independent.

The setup looks like this:



In the next section we will:

- create a code repository and a pipeline using AWS CDK

- trigger the pipeline to instantiate sample infrastructure
- trigger the pipeline to update infrastructure and perform fault injection



SETUP

In this section, we will integrate an AWS Fault Injection Simulator experiment with a CI/CD pipeline.

Create The Pipeline

We will use the AWS CDK to provision our CI/CD pipeline.

If you have not done so yet, in your cloud9 terminal clone the repository for the workshop.

```
cd ~/environment
git clone https://github.com/aws-samples/aws-fault-injection-simulator-workshop.git
```

Next change directory into the CI/CD CDK project and restore the npm packages used for the pipeline.

```
cd aws-fault-injection-simulator-workshop/resources/code/cdk/cicd/
npm install
```

Finally lets deploy our stack.

```
cdk deploy --require-approval never
```

The stack will take a few minutes to complete. You can monitor the progress from the CloudFormation Console.

Once stack creation is complete, continue to the next section.

<

>



REVIEW THE PIPELINE

Lets review the components our previous section created.

CodeCommit

Open the [AWS CodeCommit Console](#). You should see the newly created `FIS_Workshop` repository.

The screenshot shows the AWS CodeCommit console interface. The top navigation bar includes 'Developer Tools' > 'CodeCommit' > 'Repositories'. Below the navigation is a toolbar with 'Repositories' (Info), a refresh icon, 'Notify' (dropdown), 'Clone URL' (dropdown), 'View repository', 'Delete repository', and a prominent orange 'Create repository' button. A search bar and pagination controls (1) are also present. The main table lists the repository details:

Name	Description	Last modified	Clone URL
FIS_Workshop	Sample Fault Injection Simulator Workshop Repository	3 minutes ago	HTTPS SSH HTTPS (GRC)

CodeBuild

Open the [AWS CodeBuild Console](#). **Note:** You may have to select a region at the top right. You should see the `FIS_Workshop` build project.

The screenshot shows the AWS CodeBuild console interface. The top navigation bar includes 'Developer Tools' > 'CodeBuild' > 'Build projects'. Below the navigation is a toolbar with 'Build projects' (Info), a refresh icon, 'Notify' (dropdown), 'Start build' (dropdown), 'View details', 'Edit' (dropdown), 'Delete build project', and a prominent orange 'Create build project' button. A search bar and a dropdown for 'Your projects' are also present. The main table lists the build project details:

Name	Source provider	Repository	Latest build status	Description
FIS_Workshop	AWS CodePipeline	-	-	-

CodePipeline

Open the [AWS CodePipeline Console](#). You should now see the `FIS_Workshop` pipeline.

The screenshot shows the AWS CodePipeline Pipelines list page. At the top, there are navigation links: Developer Tools > CodePipeline > Pipelines. Below this is a search bar and a header row with columns for Name, Most recent execution, and Latest source revisions. A single pipeline entry is listed: FIS_Workshop, which is currently Failed.

Name	Most recent execution	Latest source revisions
FIS_Workshop	Failed	-

The pipeline will start in a failed state, since we have not uploaded any files to our repository.

To review, click on the pipeline name.

This pipeline has 3 stages.

1. **Source:** This stage will trigger the pipeline when a commit occurs in our repository.
2. **Infrastructure_Provisioning:** This stage use an AWS Cloud Formation template from our repo to create our test infrastructure and create our experiment templates.
3. **FIS:** This stage will use the AWS Code Build project to make an API call to run our experiment and monitor the results.

The screenshot shows the AWS CodePipeline Pipeline details page for FIS_Workshop. The pipeline is currently Failed. The Source stage is expanded, showing a single step named CodeCommit_So... (AWS CodeCommit) which also failed 26 minutes ago. A red arrow points down to a button labeled "Disable transition".

Notify	Edit	Stop execution	Clone pipeline	Release change
--------	------	----------------	----------------	----------------

Source Failed
Pipeline execution ID: 6c53ef84-1783-477c-a79d-73b6f61efd87

CodeCommit_So... ⓘ
AWS CodeCommit
Failed - 26 minutes ago
Details

Disable transition

⊖ Infrastructure_Provisioning Didn't Run

Create_Infrastruc... ⓘ

AWS CloudFormation ↗

⊖ Didn't Run

No executions yet

Disable transition

⊖ FIS Didn't Run

Fault_Injection ⓘ

AWS CodeBuild

⊖ Didn't Run

No executions yet

Continue to the next section to start the pipeline.



START THE PIPELINE

The pipeline is configured to run every time new code is committed to our AWS CodeCommit repository. To start our pipeline we need to commit files to our repository.

Update repository

Note

The instructions below are designed for use with Cloud9. On Cloud9 git access is enabled via the IAM role associated with the Cloud9 instance. If you would like to access the AWS CodeCommit repository from your local machine, review the [getting started documentation](#).

Adding files to our repository is a 3-step process:

- create a local copy of our repository (`clone`)
- add or update files and save them (`add` / `commit`)
- upload them to our repo (`push`)

Clone

Open the [AWS Code Commit Console](#). Click the `HTTPS` link next to the `FIS_Workshop` repository name to copy the URL to the clipboard.

In your Cloud9 terminal clone the repository (replace the URL in the example by pasting from the clipboard):

```
cd ~/environment
git clone https://git-codecommit.us-east-1.amazonaws.com/v1/repos/FIS_Workshop
cd FIS_Workshop
```

Add/Update

Copy the sample files from the resources section into the newly cloned repository.

```
cp ~/environment/aws-fault-injection-simulator-workshop/resources/code/cdk/cicd/resources/* ~/environment/FIS_Workshop/
```

Since this is the first time working with code commit, we should setup our username and email for the commit history. Run the below commands, be sure to replace the details with your information.

```
git config --global user.name "Your Name"  
git config --global user.email you@example.com
```

Finally **add** all the files in the directory, and **commit** them as a new version with a label of **Uploading Workshop files**.

```
git add .  
git commit -am "Uploading Workshop files"
```

Push

Finally **push** the files to copy them to our repository and to trigger the pipeline:

```
git push -u
```

View Progress

After you **push** the files, the pipeline will start. Open the **AWS CodePipeline Console**. You should now see the **FIS_Workshop** pipeline is in progress. Click on the pipeline name to view the step details.

Pipelines		Info		Notify	View history	Release change	Delete pipeline	Create pipeline
<input type="text"/> < 1 > 								
Name	Most recent execution		Latest source revisions			Last executed		
FIS_Workshop	In progress		CodeCommit_Source - 7c66f758: Uploading workshop files			Just now		

The pipeline runs in sequence, first running the Wait for the "Infrastructure_Provisioning" step, and on success starting the "FIS" step.

You can monitor the progress of our experiment either from the CodePipeline details page or from the AWS FIS console.

Navigate to the **FIS console**. Click on the "Experiment ID" of the running experiment. You should see the experiment in a running status:

AWS FIS > Experiments > EXPET3j5TfTmoSKkSm

EXPET3j5TfTmoSKkSm [Info](#) Refresh [Actions ▾](#)

Details			
Experiment ID EXPET3j5TfTmoSKkSm	Start time June 28, 2021, 15:12:35 (UTC-07:00)	State Running	Experiment template ID EXT2T4cJuGYDgXX6
Creation time June 28, 2021, 15:12:34 (UTC-07:00)	End time -	IAM role fisWorkshopDemo-fisrole33E76559-1B2AWY0AS3CZ3	Stop conditions -

[Actions](#) [Targets](#) [Tags](#)

Actions (1)
View your experiment template actions, action duration, and action sequences.

▼ instanceActions / aws:ec2:stop-instances Running Start: At beginning of experiment / Target: instanceTargets			
Name InstanceActions	State Running	Description -	Start after -
Targets Instances / instanceTargets	startInstancesAfterDuration PT1M		

If you expand the "instanceActions / aws:ec2:stop-instance" card (as shown above) you can see that the experiment stops the test instance, waits for 1 minute, then restarts the instance.

Wait a couple minutes for the instance to restart and the experiment to finish and refresh the page. You should see the experiment is completed successfully.

Details			
Experiment ID EXPET3j5TfTmoSKkSm	Start time June 28, 2021, 15:12:35 (UTC-07:00)	State ✔ Completed	Experiment template ID EXT2T4cJuGYDgXX6
Creation time June 28, 2021, 15:12:34 (UTC-07:00)	End time June 28, 2021, 15:13:49 (UTC-07:00)	IAM role fisWorkshopDemo-fisrole33E76559-1B2AWY0A53CZ3 ↗	Stop conditions –

Finally navigate back to the [AWS CodePipeline Console](#). You should also see that your pipeline has completed successfully.

✓ Infrastructure_Provisioning Succeeded

Pipeline execution ID: [45b360c4-c02a-43b8-83c8-9f709ecfcdba](#)

Create_Infrastructure



[AWS CloudFormation](#)

✓ Succeeded - 10 minutes ago

[Details](#)

[1c211c6d](#) CodeCommit_Source: updated template



Disable transition

✓ FIS Succeeded

Pipeline execution ID: [45b360c4-c02a-43b8-83c8-9f709ecfcdba](#)

Fault_Injection



[AWS CodeBuild](#)

✓ Succeeded - 7 minutes ago

[Details](#)

[1c211c6d](#) CodeCommit_Source: updated template

Congratulations! You have successfully integrated a Fault Injection Simulator Experiment into a CI/CD pipeline. In this scenario, we completed a happy path to ensure that our infrastructure and experiment completed without error. Continue on to the next section, where we will deploy a new version of our CloudFormation template and force our experiment (and pipeline) to fail.



FORCE A PIPELINE ERROR

In this section we will update the experiment template defined in our repository to contain a stop condition that will prevent or abort experiment execution if our cloudwatch alarm is in ALARM state.

We will then **push** the new revision to our repository which will trigger a new pipeline run, update our experiment template and execute our experiment template. Then while our pipeline is running, we will force an ALARM state. This will lead to a failure of the AWS FIS experiment and in turn to a failure of the pipeline.

Change the Infrastructure Template

Note

In this section we are directly updating the file in AWS CodeCommit. This is equivalent to the **add** / **commit** / **push** process that we performed in the previous section and creates a new revision. To subsequently synchronize the copy on your Cloud9 instance you would **git pull**

We will be making a change to our CloudFormation template that creates our EC2 Instance and defines our experiment.

Open the [AWS CodeCommit Console](#) and select the **FIS_Workshop** repository. Click on **cfn_fis_demos.yaml** and select “**Edit**” in the upper right hand corner. Edit the file as shown below to enable am AWS CloudWatch alarm as a Stop Condition.

Before:

```
121     StopConditions:  
122         - Source: none  
123         #   - Source: aws:cloudwatch:alarm  
124         #     Value:  
125         #       Fn::GetAtt:  
126         #           - cwalarm8A77F56F  
127         #           - Arn  
128         -
```

After:

```
121     StopConditions:  
122         #   - Source: none  
123             - Source: aws:cloudwatch:alarm  
124                 Value:  
125                     Fn::GetAtt:  
126                         - cwalarm8A77F56F  
127                         - Arn
```

Finally, enter your name and email at the bottom of the page and select “Commit changes”. Just like our prior `git push` this will trigger the pipeline to start.

Forcing an Error

To trigger the stop condition and simulate a failed experiment, we will manually set our CloudWatch alarm to an error state.

Navigate back to the [AWS CodePipeline Console](#) and watch the pipeline status. Once the FIS section changes to in progress, run the below command from your Cloud9 instance to force an error.

```
aws cloudwatch set-alarm-state --alarm-name "NetworkInAbnormal" --state-value "ALARM" --state-reason "testing FIS"
```

By setting this CloudWatch alarm to an error state, this will stop a running experiment or prevent the experiment from starting.

Note

We are artificially changing the alarm state. The alarm will reset to OK state after a short period of time. If you want to persist the ALARM state for longer try running the command in a loop.

To verify the Experiment was stopped, navigate to the **FIS console**. You should see that your latest experiment has failed due to the stop condition.

AWS FIS > Experiments > EXPi6qKDit3RWUm8ZP

EXPi6qKDit3RWUm8ZP [Info](#)

[Refresh](#) [Actions ▾](#)

Details			
Experiment ID EXPi6qKDit3RWUm8ZP	Start time June 28, 2021, 15:52:11 (UTC-07:00)	State ✖ Stopped	Experiment template ID EXT2T4cJuGYDgXX6
Creation time June 28, 2021, 15:52:10 (UTC-07:00)	End time June 28, 2021, 15:54:12 (UTC-07:00)	IAM role fisWorkshopDemo-fisrole33E76559-1B2AWY0AS5CZ3	Stop conditions NetworkInAbnormal

[Actions](#) [Targets](#) [Tags](#)

Actions (1)
View your experiment template actions, action duration, and action sequences.

▼ instanceActions / aws:ec2:stop-instances			
Start: At beginning of experiment / Target: instanceTargets			
Name instanceActions	State ✖ Failed	Description	Start after
Targets Instances / instanceTargets	startInstancesAfterDuration PT1M	-	-

To verify that this resulted in a failed pipeline execution navigate back to the **AWS CodePipeline Console**. You should see that your pipeline has also failed do to the experiment stopping.

✖ **FIS** Failed

Pipeline execution ID: [739f5736-0d75-4c3c-94f3-158b96441e11](#)

Fault_Injection (i)

AWS CodeBuild

✖ Failed - 5 minutes ago

[Details](#)

Congratulations! You have built a CI/CD pipeline, instrumented it with an AWS FIS experiment, and demonstrated both successful and failed experiment outcomes.

Next steps

From this starting point you can explore improvements like:

- **add more pipeline stages** - In our pipeline the experiment is the last step and does not gate progress. In a production scenario there might be additional steps that would only run if the AWS FIS experiment succeeds. Try adding a pipeline stage and verify that it only runs if the experiment succeeds.
- **explore alternative ways to change the template** - in this example we are using an AWS CloudFormation template to define the experiment template as shown in the **Experiment (CloudFormation)** section. Could you store the experiment template as a separate file and update it using the CLI as shown in **Experiment (CLI)** or expand the `runExperiment.py` script (see [code in GitHub](#))?
- **trigger AWS CloudWatch alarm from experiment template** - AWS FIS templates allow you to run a sequence of steps, try triggering the alarm from a step in the template. Hint: you could do this via the [AWS Systems Manager integration](#)/en/030_basic_content/040_ssm.html.
- **set up a real alarm**



COMMON SCENARIOS

This section covers common scenarios customers ask about.

- Simulating AZ Issues



SIMULATING AZ ISSUES

A common ask we hear is for “AZ outage” simulation. Because AWS has spent more than a decade working to *prevent* exactly those scenarios and to self-heal any disruption, there is currently no easy button solution to simulate this.

In this section we will present failure paths you can group together to build an experiment that approximates an AZ failure for your particular workload.





BACKGROUND

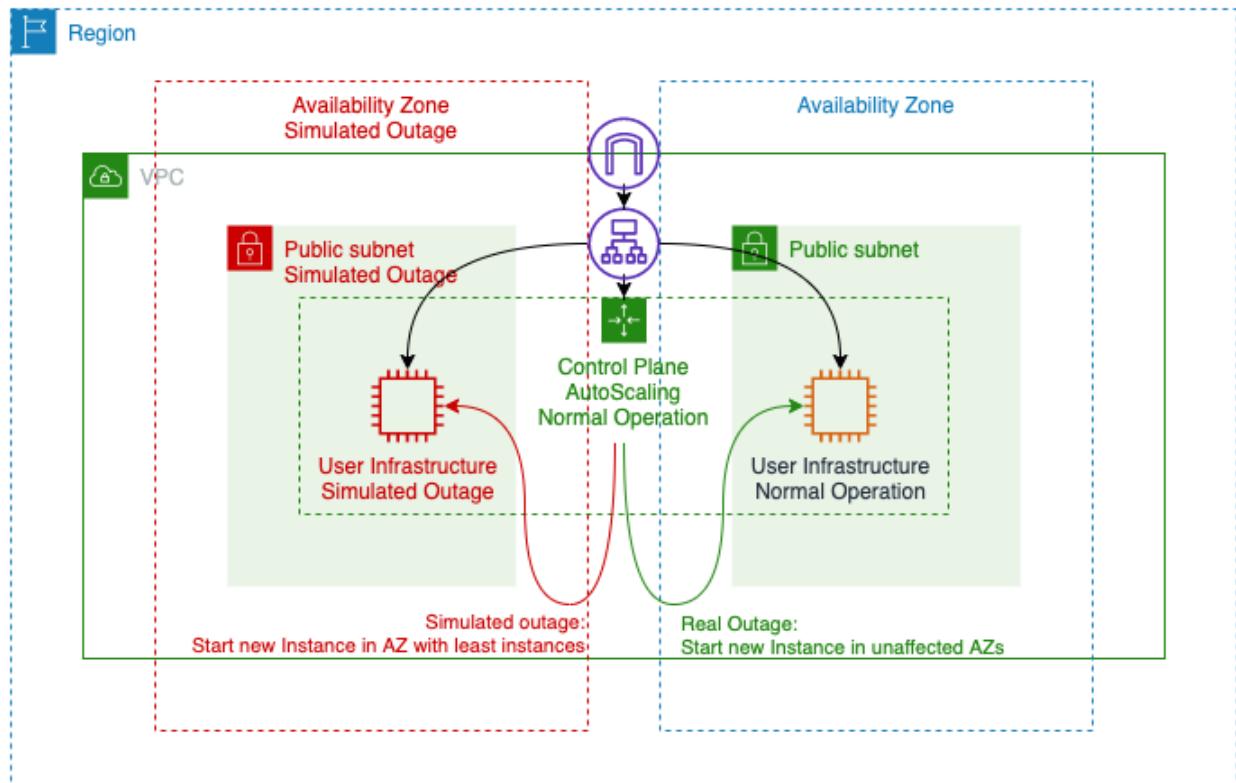
Before attempting to simulate an AZ failure it's worth considering what we mean by "AZ failure".

AZ vs. datacenter

Many of our customers phrase their idea of an AZ failure as "the whole datacenter goes away" but [AWS Availability Zones](#) are "one or more discrete data centers with redundant power, networking, and connectivity in an AWS Region" so even a full "datacenter" outage at AWS may not have the level of impact you would expect on-prem.

Control plane vs. active infrastructure

An important thing to consider is the difference between the AWS backplane and the provisioned customer infrastructure. A full datacenter outage would create awareness in the AWS control plane that resources are unavailable. In contrast a *simulated* outage to just the provisioned customer resources has to ensure that only the resources for *a single* customer are affected.



For example in the auto scaling setup we built for the **First Experiment** section, we can target EC2 instances in a given AZ for termination by filtering on `Placement.AvailabilityZone`. We expect that the “control plane”, in this case the associated Auto Scaling group, will start new instances to replace those terminated. However, since there is no actual AZ failure and the Auto Scaling group thus has no awareness of our experiment, the new instances will most likely be re-created in the AZ for which we wanted to simulate a failure.

Simulating AZ outage options

In the following sections we will cover how to approximate AZ outages for different configurations and how to build that into a bigger experiment.



OBSERVABILITY

A core aspect of chaos engineering is observability. In this section we will cover AWS tooling that supports FIS in providing observability for experiments and help in gaining the understanding needed to improve your system.



DEVOPS GURU

Warning

This section requires that you followed the [setup instructions](#) at the beginning of the workshop and allowed enough time for DevOps Guru to establish a baseline. This section also presumes that you followed the load generating steps in the **Synthetic user experience** section.

Dashboard overview

Navigate to the [DevOps Guru console](#). Once enough time has passed for DevOps Guru to generate insights you should see a dashboard similar to this:

The screenshot shows the Amazon DevOps Guru dashboard. On the left, there's a sidebar with links for Dashboard, Insights, Settings, and Cost estimator. The main area is titled "Dashboard". It features a "System health summary" card with four metrics: Total resources analyzed last hour (7), Impacted stacks (0), Ongoing reactive insights (0), and Ongoing proactive insights (0). Below this is a "System health overview" card for 16 stacks. A search bar shows "FisStack" with 4 matches. A dropdown menu shows "All stacks". The overview lists three stacks: FisStackVpc (Healthy, 0 reactive, 0 proactive insights, MTTR 38 Hours), FisStackRdsAurora (Healthy, 0 reactive, 0 proactive insights, MTTR 12 Minutes), and FisStackAsg (Healthy, 0 reactive, 0 proactive insights, MTTR 60 Minutes).

Reactive insights

Select “Insights” on the left and explore the reactive insights generated from our fault injection activities. You should see an event relating to “Application ELB” (depending on the exact order of events your dashboard may vary slightly):

The screenshot shows the Amazon DevOps Guru interface. On the left, there's a sidebar with options: Dashboard, Insights (which is selected and highlighted in orange), and Settings. Below that is a Cost estimator. The main area is titled "Insights" and has tabs for "Reactive" (which is selected) and "Proactive". Under the "Reactive" tab, there's a section titled "Reactive insights (13) Info". It says "A reactive insight lets you know about recommendations to improve the performance of your application now." Below this is a search bar with placeholder text "Filter insights by status, severity or affected resource" and a time range selector showing "12h 1d 1w 1M Custom". To the right of the search bar are navigation arrows and a refresh icon. A table below the search bar lists one insight: "ApplicationELB TargetResponseTime Anomalous In Stack FisStackAsg". The table columns are Name, Status, Severity, and Affected resources. The status is "Closed" (green checkmark), the severity is "High" (red circle), and there is 1 affected resource.

Visualizing anomalies

Selecting the event exposes more detailed information. The “Aggregated metrics” view will show timelines of different anomalous events that happened during the overall anomaly window:

Aggregated metrics**Graphed anomalies****Aggregated metrics (7)** July 21, 20:48–July 21, 21:53 UTC [Info](#)

Group by

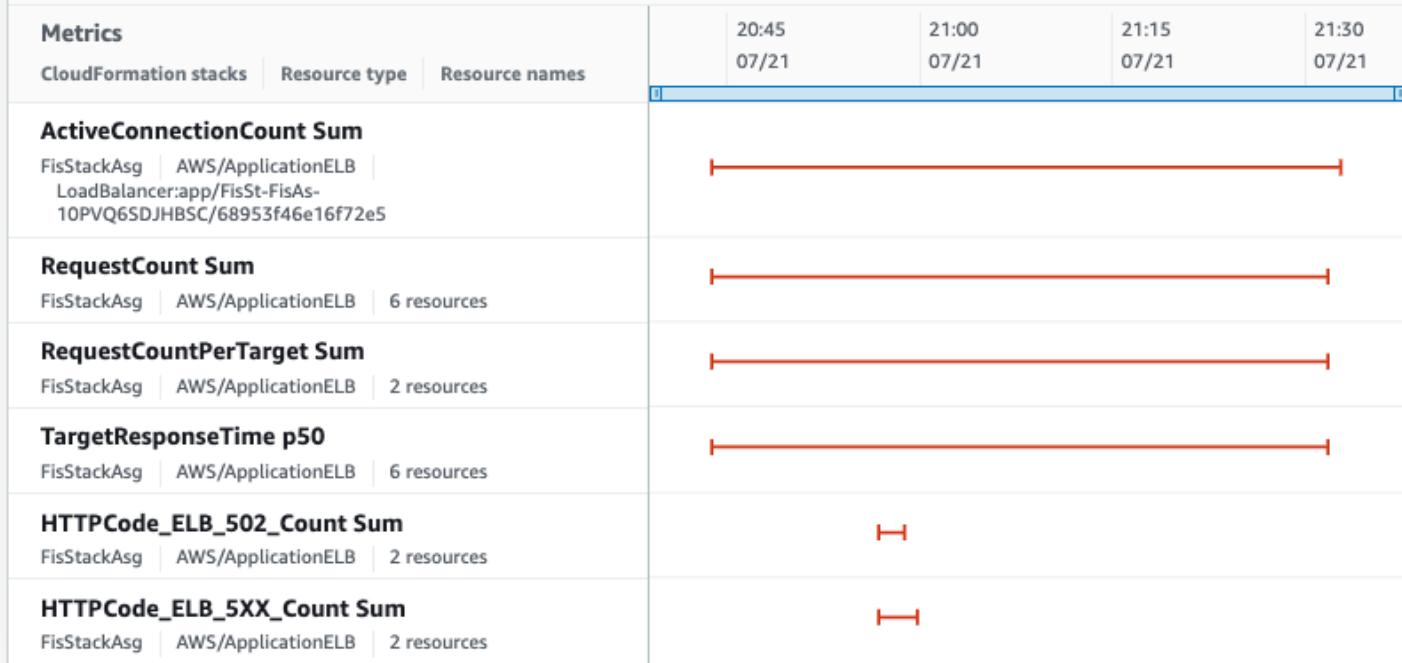
No grouping ▾

Metrics in your AWS account are analyzed to find anomalies in an insight. The timeline shows the start time of the anomaly to current time.

 Find metric by metric name, stack, resource type

< 1 2 >

🔍 🔍



Note that there may be multiple additional pages for additional events:

Aggregated metrics**Graphed anomalies****Aggregated metrics (7)** July 21, 20:48–July 21, 21:53 UTC [Info](#)

Group by

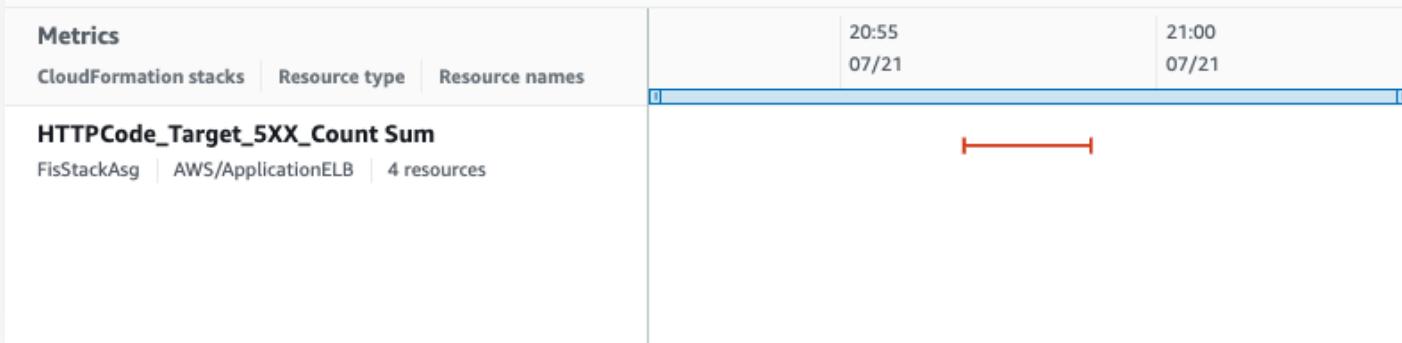
No grouping ▾

Metrics in your AWS account are analyzed to find anomalies in an insight. The timeline shows the start time of the anomaly to current time.

 Find metric by metric name, stack, resource type

< 1 2 >

🔍 🔍



Examining the example above we see that during the event

- an unusually high number of connections were made - by our external load testing tool,

- the high number of connections led to a high number of overall requests on the load balancer,
- the high number of connections led to a high number of connections to each target,
- the response time for the servers associated with the target increased substantially.

In addition to the expected direct impact of more connections, we also see unusual responses being sent:

- the number of HTTP 5xx errors increased at the load balancer,
- specifically the number of [HTTP 502](#) error increased at the load balancer,
- the number of HTTP 5xx errors originated at the load balancer target, i.e. our web servers.

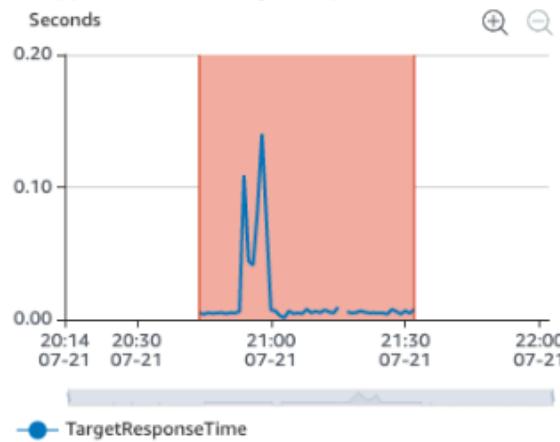
Switching to the "Graphed Anomalies" view shows the more detailed time data for each anomalous metric:

Graphed anomalies (7) July 21, 20:48–July 21, 21:53 UTC [Info](#)

Amazon DevOps Guru captures and can display the occurrence of anomalies over time.

 Find metric by metric name, stack, resource type

< 1 2 >

[1H](#) [3H](#) [12H](#) [1D](#) [3D](#) [1W](#) [2W](#) [C](#)
**AWS/ApplicationELB:TargetResponseTime**

Resource type

AWS/ApplicationELB

Resource names

TargetGroup:targetgroup/FisSt-FisAs-

L513EEGYI2M8/71ff6b943

8567639,

LoadBalancer:app/FisSt-FisAs-

10PVQ6SDJHBSC/68953f4

6e16f72e5

Stack

FisStackAsg

[View all statistics and dimensions](#)

Dimensions

TargetGroup:targetgroup/FisSt-FisAs-

L513EEGYI2M8/71ff6b943

8567639,

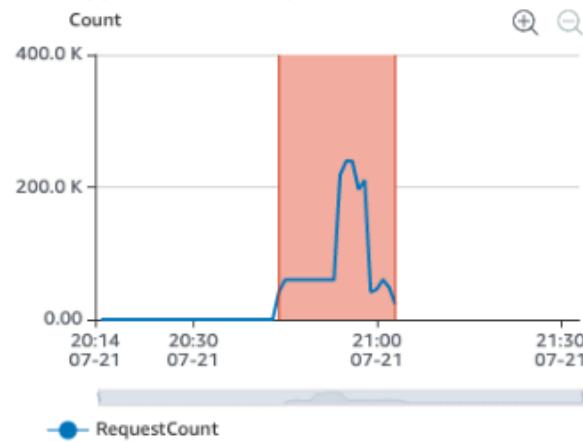
LoadBalancer:app/FisSt-FisAs-

10PVQ6SDJHBSC/68953f4

6e16f72e5

Statistics

p50

[1H](#) [3H](#) [12H](#) [1D](#) [3D](#) [1W](#) [2W](#) [C](#)
**AWS/ApplicationELB:RequestCount**

Resource type

AWS/ApplicationELB

Resource names

TargetGroup:targetgroup/FisSt-FisAs-

L513EEGYI2M8/71ff6b943

8567639,

AvailabilityZone:us-east-2a,

LoadBalancer:app/FisSt-FisAs-

10PVQ6SDJHBSC/68953f4

6e16f72e5

Stack

FisStackAsg

Dimensions

Dimensions

TargetGroup:targetgroup/FisSt-FisAs-

L513EEGYI2M8/71ff6b943

8567639,

AvailabilityZone:us-east-2a,

LoadBalancer:app/FisSt-FisAs-

10PVQ6SDJHBSC/68953f4

Statistics

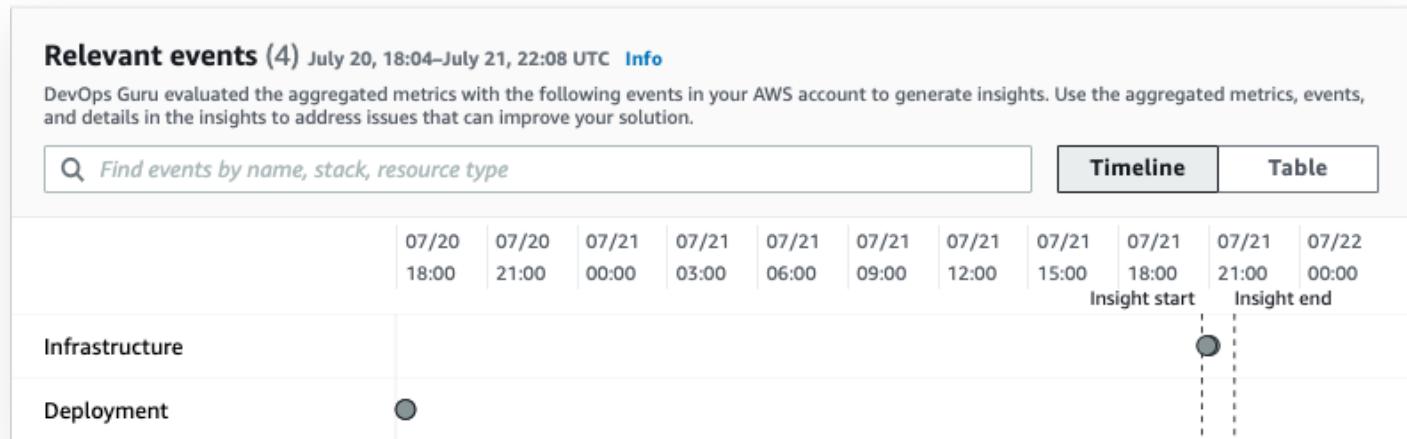
Sum

[View all statistics and dimensions](#)

Note that in this view data outside the anomaly window are set to zero to allow focusing on the relevant details during the outage.

Contextualizing with infrastructure events

In our case the anomalies arose from external load but frequently anomalies are caused by changes to code or infrastructure configuration. To help you diagnose this, DevOps Guru provides visibility into deployment and infrastructure changes associated with the anomaly. These events can be visualized in a timeline view (you can get details by clicking on the dots):



or in table format:

The figure shows a table view of the same event data. At the top, it displays "Relevant events (4) July 20, 18:04–July 21, 22:08 UTC" with a "Info" link. Below this, a message states: "DevOps Guru evaluated the aggregated metrics with the following events in your AWS account to generate insights. Use the aggregated metrics, events, and details in the insights to address issues that can improve your solution." A search bar at the top left contains the placeholder "Find events by name, stack, resource type". On the right, there are two tabs: "Timeline" (which is selected) and "Table". The table has columns: "Event name", "Resource type", "Resource name", "Time", and "AWS service". There are five rows of data:

Event name	Resource type	Resource name	Time	AWS service
CreateChangeSet	AWS::CloudFormation::Stack	FisStackAsg	Jul 20, 2021 18:19 UTC	cloudform
ExecuteChangeSet	AWS::CloudFormation::Stack	FisStackAsg	Jul 20, 2021 18:19 UTC	cloudform
RegisterTargets	AWS::ElasticLoadBalancingV2::TargetGroup		Jul 21, 2021 20:57 UTC	elasticload
DeregisterTargets	AWS::ElasticLoadBalancingV2::TargetGroup		Jul 21, 2021 21:03 UTC	elasticload

From the table format we can see that about 2h before the anomaly some changes were made to the stack configuration and deployed code. We can also see that around the time of the event instances were added to the load balancer in response to the increased load, and subsequently removed from the load balancer due to the external event subsiding.

Recommendations for improvement

Finally DevOps Guru provides “Recommendations”, links to relevant articles to help troubleshoot issues and improve overall system performance:

Recommendations (2) [Info](#)
View updates we recommend you implement to address the anomalies in this insight.

Troubleshoot errors in AWS Application Elastic Load Balancer (ELB) [🔗](#)

Your load balancer is throwing errors. To learn more and troubleshoot load balancer errors, see [Troubleshoot errors in AWS Application Elastic Load Balancer \(ELB\)](#) [🔗](#).

Why is DevOps Guru recommending this?

The **HTTPCode_Target_5XX_Count** metric in ApplicationELB breached a high threshold.

Related metric (2)

HTTPCode_ELB_5XX_Count
ApplicationELB | us-east-2b, app/FisSt-FisAs-10PVQ6SDJHBSC/68953f46e16f72e5

HTTPCode_Target_5XX_Count
ApplicationELB | targetgroup/FisSt-FisAs-L513EEGYI2M8/71ff6b9438567639, us-east-2a, app/FisSt-FisAs-10PVQ6SDJHBSC/68953f46e16f72e5

Further reading

To learn more about DevOps Guru, see the [documentation](#), and explore using [DevOps guru on serverless infrastructure](#) as well as [larger deployment strategies](#).

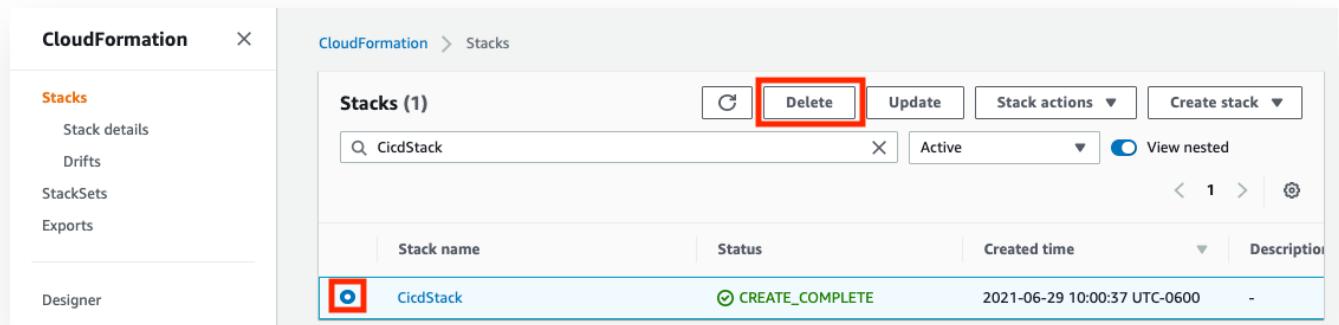


CLEANUP

To ensure you don't incur any further costs after the workshop, please follow these instructions to delete the resources you created.

Manually

- If you created the CI/CD stack and ran the pipeline, first start by deleting the infrastructure provisioned by the pipeline:
 - Navigate to the [AWS CloudFormation console](#) and find the stack named `fisworkshopDemo`
 - Select the stack
 - Select "Delete"
- Once, the `fisworkshopDemo` is deleted, following the same procedure as above, delete the `CicdStack` stack



- Following the same procedure as above, delete the following stacks
 - `FisStackEks`
 - `FisStackEcs`
 - `FisStackRdsAurora`
 - `FisStackLoadGen`
 - `FisStackAsg`
 - `FisStackVpc`
- Delete the CloudWatch log groups:

- Navigate to the [AWS CloudWatch console](#)
- Search for `fis-workshop`
- Select the checkboxes
- Under “Actions” select “Delete log group(s)”
- Delete Cloud9 Environments
 - Navigate to the [AWS Cloud9 console](#)
 - Delete the Cloud9 environment that you use during the workshop

Using a script

In your cloud9 terminal where you performed the **Provision AWS resources** step run the following commands:

```
cd ~/environment
```

```
cd aws-fault-injection-simulator-workshop  
cd resources/templates  
../cleanup.sh
```

Retained resources

CloudWatch metrics and FIS experiments will be retained until the end of their respective expiration periods.

